

332

Advanced Computer Architecture

Chapter 3

Dynamic scheduling, out-of-order execution, register renaming and speculative execution

January 2018
Paul H J Kelly

These lecture notes are partly based on the course text, Hennessy and Patterson's *Computer Architecture, a quantitative approach* (4-6th ed), and on the lecture slides of David Patterson's Berkeley course (CS252)

Course materials online at
<http://www.doc.ic.ac.uk/~phjk/AdvancedCompArchitecture.html>

Advantages of Dynamic Scheduling

- Handles cases when dependences unknown at compile time
 - (e.g., because they may involve a memory reference)
- It simplifies the compiler
- Allows code that compiled for one pipeline to run efficiently on a different pipeline
- Hardware speculation, a technique with significant performance advantages, builds on dynamic scheduling

HW Schemes: Instruction Parallelism

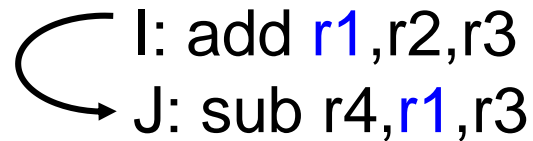
- Key idea: Allow instructions behind stall to proceed

```
DIVD    F0,F2,F4
ADDD    F10,F0,F8
SUBD    F12,F8,F14
```

- Enables **out-of-order execution** and allows **out-of-order completion**
- We will distinguish when an instruction is **issued**, *begins execution* and when it *completes execution*; between these two times, the instruction is *in execution*
- In a dynamically scheduled pipeline, all instructions pass through issue stage in order (**in-order issue**)

Data Dependence and Hazards

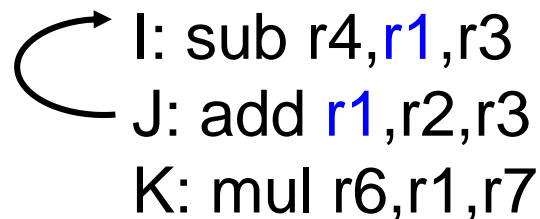
- Instr_j is **data dependent** on Instr_i
 Instr_j tries to read operand before Instr_i writes it



- or Instr_j is data dependent on Instr_k which is dependent on Instr_i
- Caused by a “**True Dependence**” (compiler term)
- If true dependence caused a hazard in the pipeline, called a **Read After Write (RAW) hazard**

Name Dependence #1: Anti-dependence

- **Name dependence**: when two instructions use same register or memory location, called a **name**, but no flow of data between the instructions associated with that name
- There are two kinds:
- Name dependence #1: anti-dependence/WAR
Instr_j writes operand before Instr_i reads it:



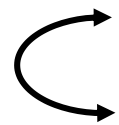
```
I: sub r4, r1, r3
J: add r1, r2, r3
K: mul r6, r1, r7
```

Called an “**anti-dependence**” by compiler writers.
This results from reuse of the name “**r1**”

- If anti-dependence caused a hazard in the pipeline, called a **Write After Read (WAR) hazard**

Name Dependence #2: Output dependence

- Instr_j writes operand before Instr_i writes it.

 I: sub **r1**,r4,r3
J: add **r1**,r2,r3
K: mul r6,r1,r7

- Called an “**output dependence**” by compiler writers
This also results from the reuse of name “**r1**”
- If anti-dependence caused a hazard in the pipeline,
called a **Write After Write (WAW) hazard**

Dynamic Scheduling Step 1

- Simple pipeline had one stage to check both structural and data hazards: Instruction Decode (ID), also called Instruction Issue
- Split the ID pipe stage of simple 5-stage pipeline into 2 stages:
- *Issue*—Decode instructions, check for structural hazards
- *Read operands*—Wait until no data hazards, then read operands

Instructions are *issued* in-order


But may stall at the Read Operands stage while others execute

Tomasulo's Algorithm

- For IBM 360/91 (before caches!)
- Goal: High Performance without special compilers
- Small number of floating point registers (4 in 360) prevented interesting compiler scheduling of operations
 - This led Tomasulo to try to figure out how to increase the effective number of registers — **renaming in hardware!**
- Why study a 1966 Computer?
- The descendents of this have flourished!
 - **Alpha 21264, HP 8000, MIPS 10000/R12000, Pentium II/III/4, Core, Core2, Nehalem, Sandy Bridge, Ivy Bridge, Haswell, AMD K5,K6,Athlon, Opteron, Phenom, PowerPC 603/604/G3/G4/G5, Power 3,4,5,6, ARM A15, ...**



- CPU cycle time: 60 nanoseconds
- memory cycle time (to fetch and store eight bytes in parallel): 780ns
- Standard memory capacity: 2,097,152B interleaved 16 ways (magnetic cores)
- Up to 6,291,496 bytes of main storage
- Up to 16.6-million additions/second
- Ca.120K gates, ECL
- Solid Logic Technology (SLT), an IBM invention which encapsulated 5-6 transistors into a small module--a transition technology between discrete transistors and the IC
- About 12 were made



NASA's Space Flight Center in Greenbelt, Md, January 1968

NASA Center for Computational Sciences


See:

Some Reflections on Computer Engineering: 30 Years after the IBM System 360 Model 91

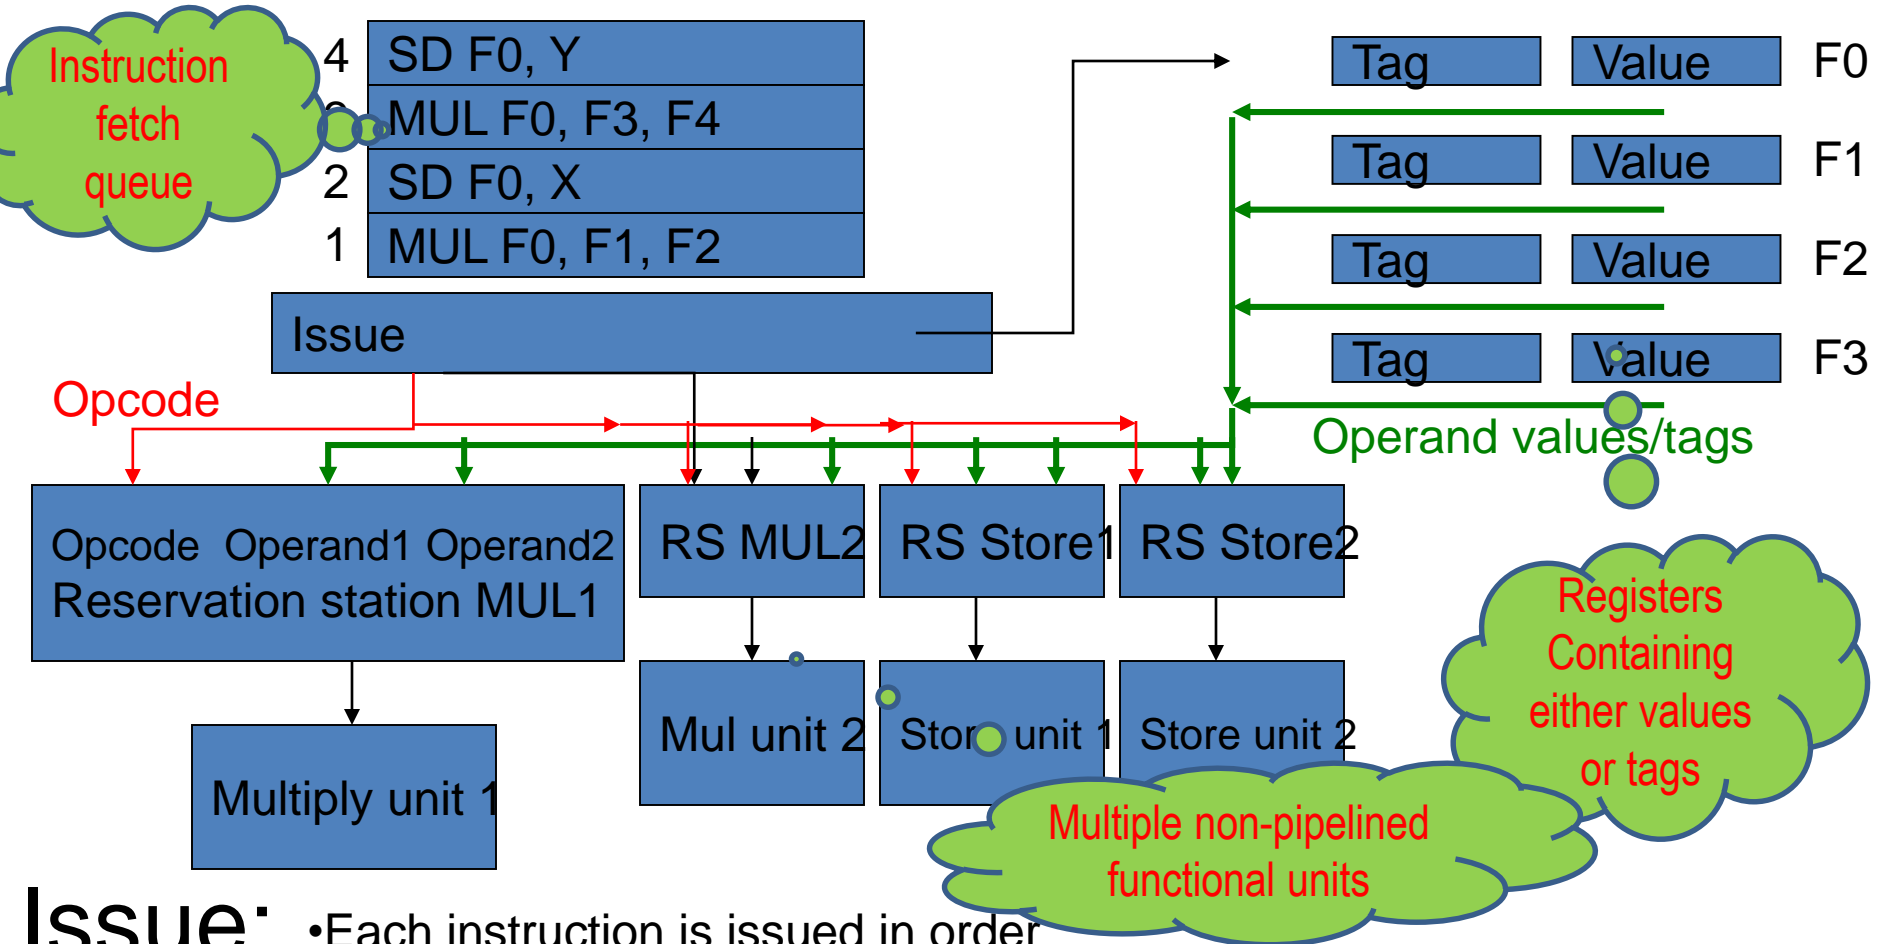
Michael J. Flynn

<ftp://arith.stanford.edu/tr/micro30.ps.Z>

Source: <http://www.columbia.edu/acis/history/36091.html>



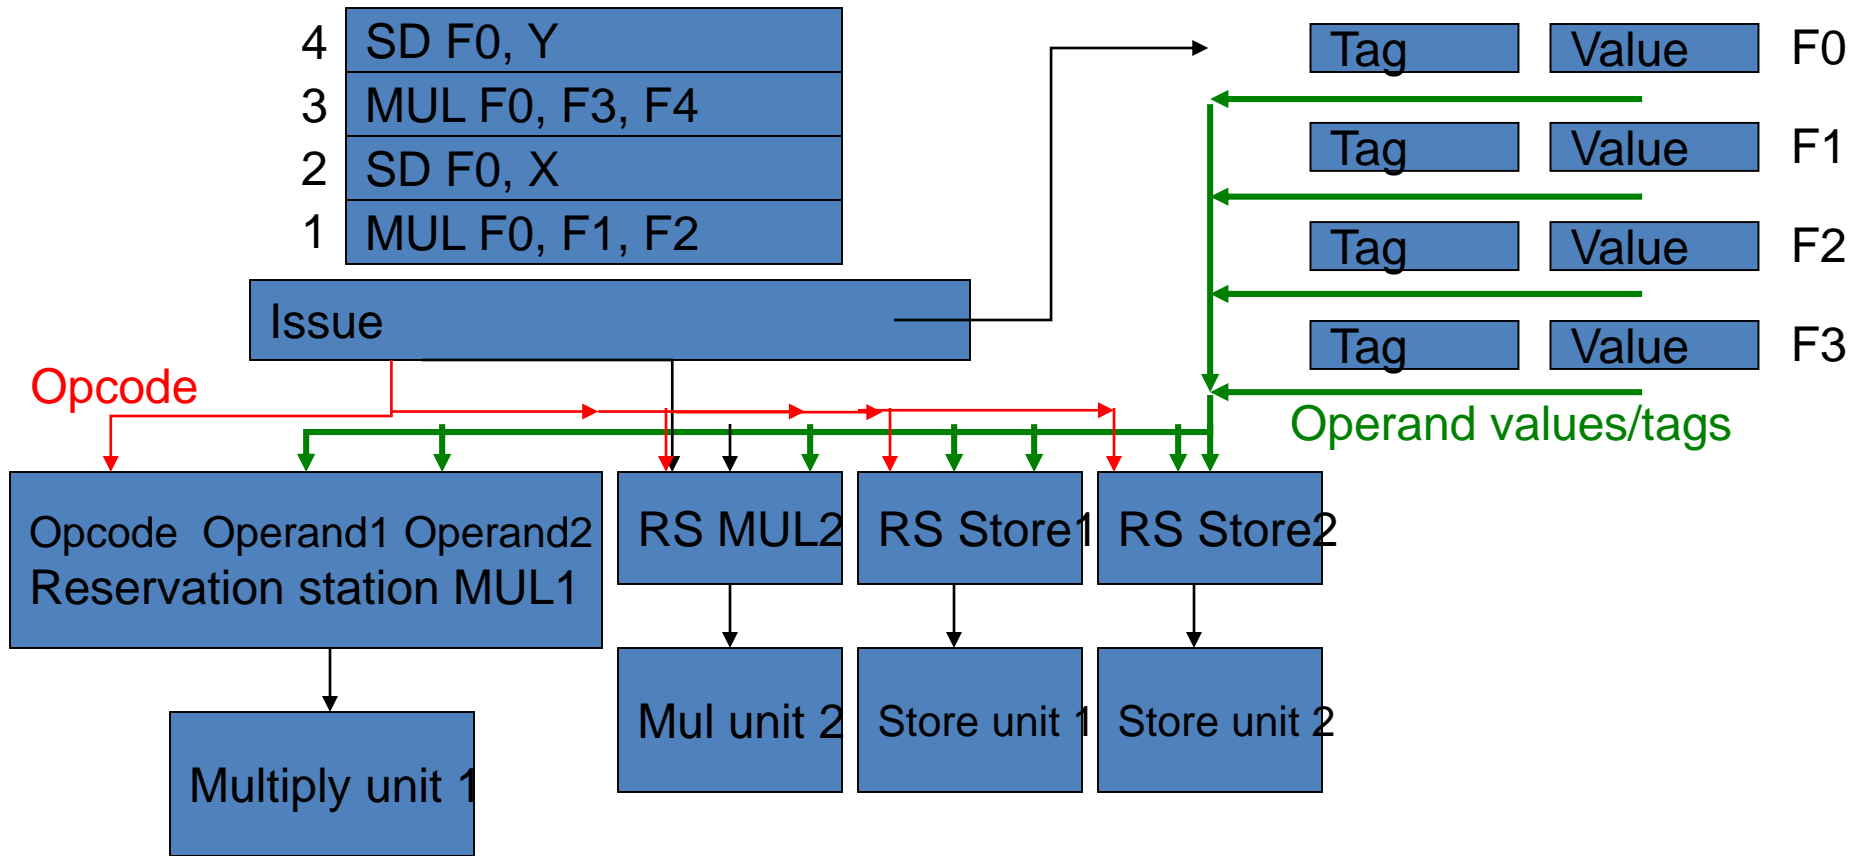
Tomasulo – closer look at instruction processing



Issue:


- Each instruction is issued in order
- Issue unit collects operands from the two instruction's source registers
- Result may be a value, or, if value will be computed by an uncompleted instruction, the tag of the RS to which it was issued.
- When instruction 1 is issued, F0 is updated to get result from MUL1
- When instruction 3 is issued, F0 is updated to get result from MUL2

Tomasulo – closer look at instruction processing



- Issue:**
- Each instruction is issued in order
 - Issue unit collects operands from the two instruction's source registers
 - Result may be a value, or, if value will be computed by an uncompleted instruction, the tag of the RS to which it was issued.
 - When instruction 1 is issued, F0 is updated to get result from MUL1
 - When instruction 3 is issued, F0 is updated to get result from MUL2

Tomasulo – closer look at instruction processing



- Write-back:**
- Instructions may complete out of order
 - Result is broadcast on CDB
 - Carrying tag of RS to which instruction was originally issued
 - All RSs and registers monitor CDB and collect value if tag matches
 - Any RS which has both operands and whose FU is free fires.
 - When MUL1 completes result goes to store unit but not F0

Three Stages of Tomasulo Algorithm

1. **Issue**—get instruction from FP Op Queue

If reservation station free (no structural hazard),
control issues instr & sends operands (renames registers).

2. **Execute**—operate on operands (EX)

When both operands ready then execute;
if not ready, watch Common Data Bus for result

3. **Write result**—finish execution (WB)

Write on Common Data Bus to all awaiting units;
mark reservation station available

Two buses:

- Normal data bus: data+destination (“go to” bus)
 - Used at Issue
- Common data bus: data+source (“come from” bus)
 - Used at WB
 - 64 bits of data + 4 bits of Functional Unit source address
 - Write if matches expected Functional Unit (produces result)
 - Does the broadcast

360/91 pipeline

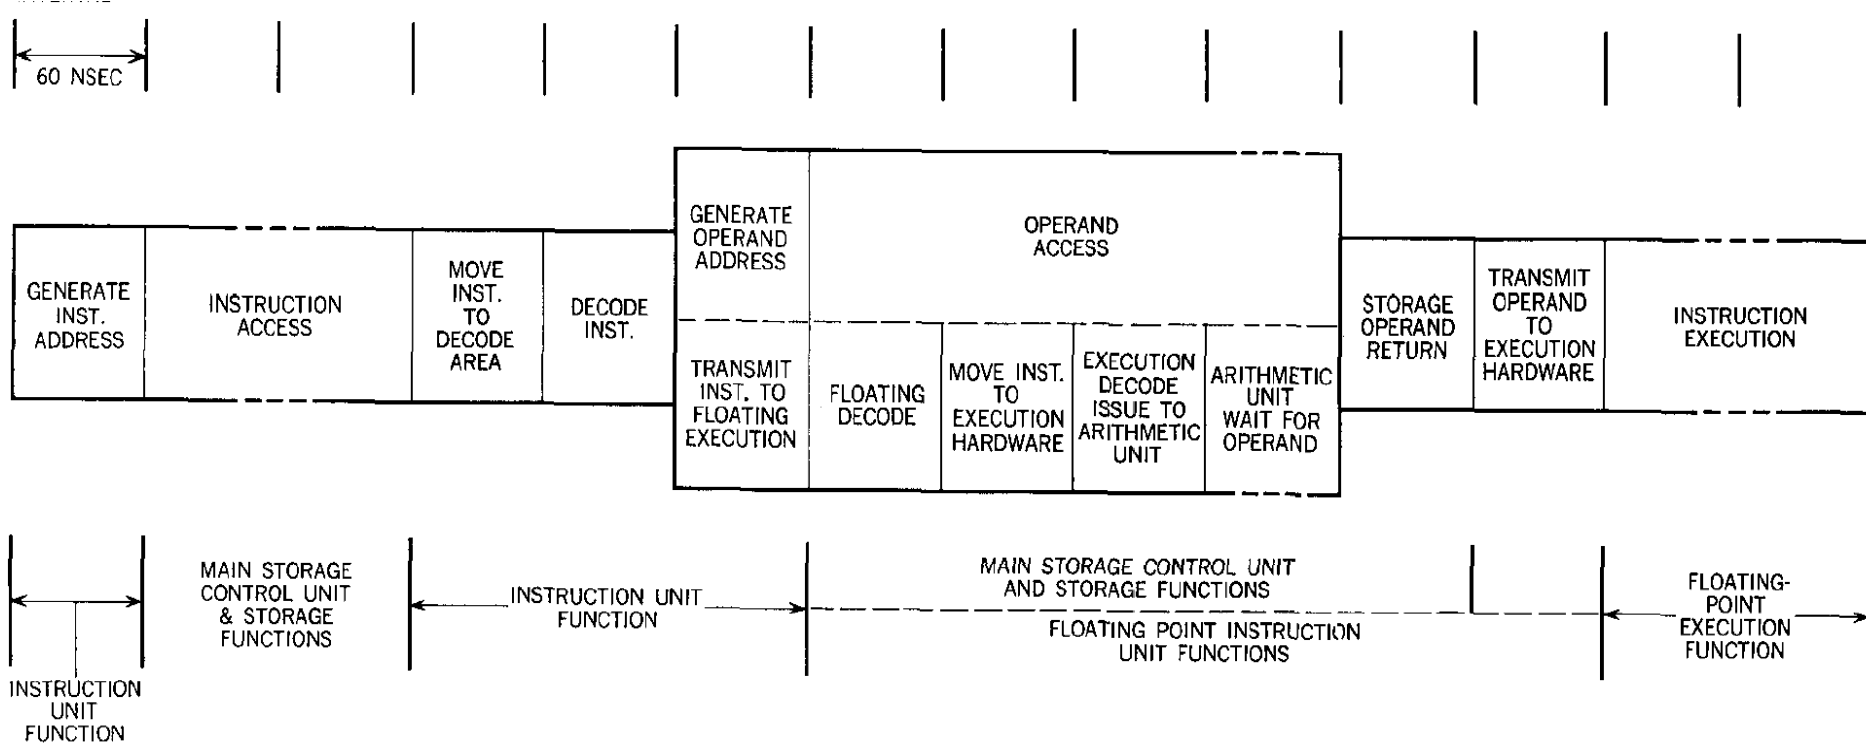


Figure 3 CPU “assembly-line stations required to accommodate a typical floating-point storage-to-register instruction.

11-12 circuit levels per pipeline stage, of 5-6ns each

CPU consists of three physical frames, each having dimensions 66" L X 15" D X 78" H

See: The IBM System/360 Model 91: Machine Philosophy and Instruction-Handling, by D. W. Anderson, F. J. Sparacio, R. M. Tomasulo. IBM J. R&D (1967), <http://www.research.ibm.com/journal/rd/111/ibmrd1101C.pdf>

Tomasulo Drawbacks

- Complexity
 - delays of 360/91, MIPS 10000, Alpha 21264, IBM PPC 620
 - Many associative stores (CDB) at high speed
- Performance limited by Common Data Bus
 - Each CDB must go to multiple functional units
⇒ high capacitance, high wiring density
 - Number of functional units that can complete per cycle limited to one!
 - Multiple CDBs ⇒ more FU logic for parallel assoc stores
- Non-precise interrupts!
 - We will address this later

Tomasulo Loop Example

Loop: LD	F0	0	R1
MULTD	F4	F0	F2
SD	F4	0	R1
SUBI	R1	R1	#8
BNEZ	R1	Loop	

- Assume Multiply takes 4 clocks
- Assume 1st load takes 8 clocks (L1 cache miss), 2nd load takes 1 clock (hit)
- To be clear, will show clocks for SUBI, BNEZ
 - Reality: integer instructions ahead of Fl. Pt. Instructions
- Show 2 iterations

Loop Example

Instruction status:

						<i>Exec Write</i>			
<i>ITER</i>	<i>Instruction</i>	<i>j</i>	<i>k</i>	<i>Issue</i>	<i>CompResult</i>	<i>Busy</i>	<i>Addr</i>	<i>Fu</i>	
Iteration Count	1	LD	F0	0	R1		Load1	No	
	1	MULTD	F4	F0	F2		Load2	No	
	1	SD	F4	0	R1		Load3	No	
	2	LD	F0	0	R1		Store1	No	
	2	MULTD	F4	F0	F2		Store2	No	
	2	SD	F4	0	R1		Store3	No	

Reservation Stations:

<i>Time</i>	<i>Name</i>	<i>Busy</i>	<i>Op</i>	<i>Vj</i>	<i>Vk</i>	<i>Qj</i>	<i>Qk</i>	<i>Code:</i>	
	Add1	No						LD	F0 0 R1
	Add2	No						MULTD	F4 F0 F2
	Add3	No						SD	F4 0 R1
	Mult1	No						SUBI	R1 R1 #8
	Mult2	No						BNEZ	R1 Loop

Register result status

<i>Clock</i>	<i>R1</i>	<i>F0</i>	<i>F2</i>	<i>F4</i>	<i>F6</i>	<i>F8</i>	<i>F10</i>	<i>F12</i>	...	<i>F30</i>
0	80									

Value of Register used for address, iteration control

Reservation Station Components

Op: Operation to perform in the unit (e.g., + or −)

V_j, V_k: **Value** of Source operands

- Store buffers has V field, result to be stored

Q_j, Q_k: Reservation stations producing source registers (value to be written)

- Note: Q_j, Q_k=0 => ready
- Store buffers only have Q_i for RS producing result

Busy: Indicates reservation station or FU is busy

Register result status—Indicates which functional unit will write each register, if one exists. Blank when no pending instructions that will write that register.

Loop Example Cycle 1

Instruction status:

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1		Load1	Yes	80
							Load2	No	
							Load3	No	
							Store1	No	
							Store2	No	
							Store3	No	

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
	Mult1	No						SUBI R1 R1 #8
	Mult2	No						BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
1	80	Load1								

Instruction status:

					Exec Write				
ITER	Instruction	<i>j</i>	<i>k</i>	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1		Load1	Yes	80
1	MULTD	F4	F0	F2	2		Load2	No	
							Load3	No	
							Store1	No	
							Store2	No	
							Store3	No	

Reservation Stations:

					S1		S2	RS				
Time	Name	Busy	Op	<i>Vj</i>	<i>Vk</i>	<i>Qj</i>	<i>Qk</i>	Code:				
	Add1	No						LD	F0	0	R1	
	Add2	No						MULTD	F4	F0	F2	←
	Add3	No						SD	F4	0	R1	
	Mult1	Yes	Multd		R(F2)	Load1		SUBI	R1	R1	#8	
	Mult2	No						BNEZ	R1	Loop		

Register result status

Clock	R1		F0	F2	F4	F6	F8	F10	F12	...	F30
2	80	Fu	Load1		Mult1						

Loop Example Cycle 3

Instruction status:

					<i>Exec Write</i>				
<i>ITER</i>	<i>Instruction</i>		<i>j</i>	<i>k</i>	<i>Issue</i>	<i>CompResult</i>	<i>Busy</i>	<i>Addr</i>	<i>Fu</i>
1	LD	F0	0	R1	1		Load1	Yes	80
1	MULTD	F4	F0	F2	2		Load2	No	
1	SD	F4	0	R1	3		Load3	No	
							Store1	Yes	80
							Store2	No	
							Store3	No	
									Mult1

Reservation Stations:

<i>Time</i>	<i>Name</i>	<i>Busy</i>	<i>Op</i>	<i>Vj</i>	<i>Vk</i>	<i>Qj</i>	<i>Qk</i>	<i>Code:</i>
	Add1	No						LD
	Add2	No						F0
	Add3	No						0
	Mult1	Yes	Multd					R1
	Mult2	No						F2
								SD
								F4
								0
								R1
								#8
								BNEZ
								R1
								Loop

Register result status

<i>Clock</i>	<i>R1</i>	<i>F0</i>	<i>F2</i>	<i>F4</i>	<i>F6</i>	<i>F8</i>	<i>F10</i>	<i>F12</i>	<i>...</i>	<i>F30</i>
3	80	<i>Fu</i>	Load1	Mult1						

- Implicit renaming sets up data flow graph

Loop Example Cycle 5

Instruction status:

Exec Write

ITER	Instruction		<i>j</i>	<i>k</i>	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1			Load1	Yes	80
1	MULTD	F4	F0	F2	2			Load2	No	
1	SD	F4	0	R1	3			Load3	No	
								Store1	Yes	80
								Store2	No	
								Store3	No	
										Mult1

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:				
	Add1	No						LD	F0	0	R1	
	Add2	No						MULTD	F4	F0	F2	
	Add3	No						SD	F4	0	R1	
	Mult1	Yes	Multd		R(F2)	Load1		SUBI	R1	R1	#8	
	Mult2	No						BNEZ	R1	Loop		

Register result status

Clock	R1		F0	F2	F4	F6	F8	F10	F12	...	F30
5	72	Fu	Load1		Mult1						

- And, BNEZ instruction (not in FP queue)

Loop Example Cycle 6

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1		Load1	Yes 80	
1	MULTD	F4	F0	F2	2		Load2	Yes 72	
1	SD	F4	0	R1	3		Load3	No	
2	LD	F0	0	R1	6		Store1	Yes 80	Mult1
							Store2	No	
							Store3	No	

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load1		SUBI R1 R1 #8
	Mult2	No						BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
6	72	Fu	Load2			Mult1				

- Notice that F0 never sees Load from location 80

Loop Example Cycle 7


Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1		Load1	Yes	80
1	MULTD	F4	F0	F2	2		Load2	Yes	72
1	SD	F4	0	R1	3		Load3	No	
2	LD	F0	0	R1	6		Store1	Yes	80
2	MULTD	F4	F0	F2	7		Store2	No	
							Store3	No	
									Mult1

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2 
	Add3	No						SD F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load1		SUBI R1 R1 #8
	Mult2	Yes	Multd		R(F2)	Load2		BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
7	72	Fu	Load2	Mult2						

- Register file completely detached from computation
- First and Second iteration completely overlapped

Loop Example Cycle 8

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	CompResult	Busy	Addr	Fu
1	LD	F0	0	R1	1	Load1	Yes 80	
1	MULTD	F4	F0	F2	2	Load2	Yes 72	
1	SD	F4	0	R1	3	Load3	No	
2	LD	F0	0	R1	6	Store1	Yes 80	Mult1
2	MULTD	F4	F0	F2	7	Store2	Yes 72	Mult2
2	SD	F4	0	R1	8	Store3	No	

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load1		SUBI R1 R1 #8
	Mult2	Yes	Multd		R(F2)	Load2		BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
8	72	Fu	Load2	Mult2						

Loop Example Cycle 9

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	Load1	Yes	80
1	MULTD	F4	F0	F2	2		Load2	Yes	72
1	SD	F4	0	R1	3		Load3	No	
2	LD	F0	0	R1	6		Store1	Yes	80
2	MULTD	F4	F0	F2	7		Store2	Yes	72
2	SD	F4	0	R1	8		Store3	No	
									Mult1
									Mult2

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load1		SUBI R1 R1 #8
	Mult2	Yes	Multd		R(F2)	Load2		BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
9	72	Fu	Load2	Mult2						

- Load1 completing: who is waiting?
- Note: Dispatching SUBI

Loop Example Cycle 10

Instruction status:

					<i>Exec Write</i>				
<i>ITER</i>	<i>Instruction</i>	<i>j</i>	<i>k</i>	<i>Issue</i>	<i>Comp</i>	<i>Result</i>	<i>Busy</i>	<i>Addr</i>	<i>Fu</i>
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2		Load2	Yes	72
1	SD	F4	0	R1	3		Load3	No	
2	LD	F0	0	R1	6	10	Store1	Yes	80
2	MULTD	F4	F0	F2	7		Store2	Yes	72
2	SD	F4	0	R1	8		Store3	No	
									Mult1
									Mult2

Reservation Stations:

					<i>S1</i>	<i>S2</i>	<i>RS</i>				
<i>Time</i>	<i>Name</i>	<i>Busy</i>	<i>Op</i>	<i>Vj</i>	<i>Vk</i>	<i>Qj</i>	<i>Qk</i>	<i>Code:</i>			
	Add1	No						LD	F0	0	R1
	Add2	No						MULTD	F4	F0	F2
	Add3	No						SD	F4	0	R1
4	Mult1	Yes	Multd	M[80]	R(F2)			SUBI	R1	R1	#8
	Mult2	Yes	Multd		R(F2)	Load2		BNEZ	R1	Loop	

Register result status

<i>Clock</i>	<i>R1</i>	<i>F0</i>	<i>F2</i>	<i>F4</i>	<i>F6</i>	<i>F8</i>	<i>F10</i>	<i>F12</i>	...	<i>F30</i>
10	64	Load2	Mult2							

- Load2 completing: who is waiting?
- Note: Dispatching BNEZ

Loop Example Cycle 11

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2			Load2	No
1	SD	F4	0	R1	3			Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	Yes 80 Mult1
2	MULTD	F4	F0	F2	7			Store2	Yes 72 Mult2
2	SD	F4	0	R1	8			Store3	No

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
3	Mult1	Yes	Multd	M[80]	R(F2)			SUBI R1 R1 #8
4	Mult2	Yes	Multd	M[72]	R(F2)			BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
11	64	Fu	Load3			Mult2				

- Next load in sequence

Loop Example Cycle 12

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2			Load2	No
1	SD	F4	0	R1	3			Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	Yes 80 Mult1
2	MULTD	F4	F0	F2	7			Store2	Yes 72 Mult2
2	SD	F4	0	R1	8			Store3	No

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2 ←
	Add3	No						SD F4 0 R1
2	Mult1	Yes	Multd	M[80]	R(F2)			SUBI R1 R1 #8
3	Mult2	Yes	Multd	M[72]	R(F2)			BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
12	64	Fu	Load3	Mult2						

- Why not issue third multiply?

Loop Example Cycle 13

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2			Load2	No
1	SD	F4	0	R1	3			Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	Yes 80 Mult1
2	MULTD	F4	F0	F2	7			Store2	Yes 72 Mult2
2	SD	F4	0	R1	8			Store3	No

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2 ←
	Add3	No						SD F4 0 R1
1	Mult1	Yes	Multd	M[80]	R(F2)			SUBI R1 R1 #8
2	Mult2	Yes	Multd	M[72]	R(F2)			BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
13	64	Fu	Load3	Mult2						

- Why not issue third store?

Loop Example Cycle 14

Instruction status:

Exec Write

ITER	Instruction		j	k	Issue	Comp	Result		Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No		
1	MULTD	F4	F0	F2	2	14		Load2	No		
1	SD	F4	0	R1	3			Load3	Yes	64	
2	LD	F0	0	R1	6	10	11	Store1	Yes	80	Mult1
2	MULTD	F4	F0	F2	7			Store2	Yes	72	Mult2
2	SD	F4	0	R1	8			Store3	No		

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2 ←
	Add3	No						SD F4 0 R1
0	Mult1	Yes	Multd	M[80]	R(F2)			SUBI R1 R1 #8
1	Mult2	Yes	Multd	M[72]	R(F2)			BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
14	64	Fu	Load3	Mult2						

- Mult1 completing. Who is waiting?

Loop Example Cycle 15

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2	14	15	Load2	No
1	SD	F4	0	R1	3			Load3	Yes
2	LD	F0	0	R1	6	10	11	Store1	Yes
2	MULTD	F4	F0	F2	7	15		Store2	Yes
2	SD	F4	0	R1	8			Store3	No
									[80]*R2
									Mult2

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD
	Add2	No						F0
	Add3	No						0
	Mult1	No						R1
0	Mult2	Yes	Multd	M[72]	R(F2)			MULTD
								F4
								F0
								0
								R1
								SUBI
								R1
								R1
								#8
								BNEZ
								Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
15	64	Fu	Load3	Mult2						

- Mult2 completing. Who is waiting?

Loop Example Cycle 16

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2	14	15	Load2	No
1	SD	F4	0	R1	3			Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	Yes 80 [80]*R2
2	MULTD	F4	F0	F2	7	15	16	Store2	Yes 72 [72]*R2
2	SD	F4	0	R1	8			Store3	No

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2 ←
	Add3	No						SD F4 0 R1
4	Mult1	Yes	Multd		R(F2)	Load3		SUBI R1 R1 #8
	Mult2	No						BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
16	64	Fu	Load3	Mult1						

Loop Example Cycle 17

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2	14	15	Load2	No
1	SD	F4	0	R1	3			Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	Yes 80 [80]*R2
2	MULTD	F4	F0	F2	7	15	16	Store2	Yes 72 [72]*R2
2	SD	F4	0	R1	8			Store3	Yes 64 Mult1

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load3		SUBI R1 R1 #8
	Mult2	No						BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
17	64	Fu	Load3	Mult1						

Loop Example Cycle 18

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2	14	15	Load2	No
1	SD	F4	0	R1	3	18		Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	Yes 80 [80]*R2
2	MULTD	F4	F0	F2	7	15	16	Store2	Yes 72 [72]*R2
2	SD	F4	0	R1	8			Store3	Yes 64 Mult1

Reservation Stations:

S1 S2 RS

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
	Add1	No						LD F0 0 R1
	Add2	No						MULTD F4 F0 F2
	Add3	No						SD F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load3		SUBI R1 R1 #8
	Mult2	No						BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
18	64	Fu	Load3		Mult1					

Loop Example Cycle 19

Instruction status:

Exec Write

ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu
1	LD	F0	0	R1	1	9	10	Load1	No
1	MULTD	F4	F0	F2	2	14	15	Load2	No
1	SD	F4	0	R1	3	18	19	Load3	Yes 64
2	LD	F0	0	R1	6	10	11	Store1	No
2	MULTD	F4	F0	F2	7	15	16	Store2	Yes 72 [72]*R2
2	SD	F4	0	R1	8	19		Store3	Yes 64 Mult1

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:
Add1	No							LD F0 0 R1
Add2	No							MULTD F4 F0 F2
Add3	No							SD F4 0 R1
Mult1	Yes	Multd			R(F2)	Load3		SUBI R1 R1 #8
Mult2	No							BNEZ R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
19	56	Fu	Load3		Mult1					

Loop Example Cycle 20

Instruction status:

					Exec Write					
ITER	Instruction	j	k	Issue	Comp	Result	Busy	Addr	Fu	
1	LD	F0	0	R1	1	9	10	Load1	Yes	56
1	MULTD	F4	F0	F2	2	14	15	Load2	No	
1	SD	F4	0	R1	3	18	19	Load3	Yes	64
2	LD	F0	0	R1	6	10	11	Store1	No	
2	MULTD	F4	F0	F2	7	15	16	Store2	No	
2	SD	F4	0	R1	8	19	20	Store3	Yes	64
										Mult1

Reservation Stations:

Time	Name	Busy	Op	Vj	Vk	Qj	Qk	Code:	
	Add1	No						LD	F0 0 R1
	Add2	No						MULTD	F4 F0 F2
	Add3	No						SD	F4 0 R1
	Mult1	Yes	Multd		R(F2)	Load3		SUBI	R1 R1 #8
	Mult2	No						BNEZ	R1 Loop

Register result status

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
20	56	Fu	Load1		Mult1					

- Once again: In-order issue, out-of-order execution and out-of-order completion.

Why can Tomasulo overlap iterations of loops?

- Register renaming
 - Multiple iterations use different physical destinations for registers (dynamic loop unrolling).
- Reservation stations
 - Permit instruction issue to advance past integer control flow operations
 - Also buffer old values of registers - totally avoiding the WAR stall that we saw in the scoreboard.
- Other perspective:
 - The CDB is doing forwarding, bypassing the registers
 - Builds the data flow dependency graph on the fly

What about Precise Interrupts?

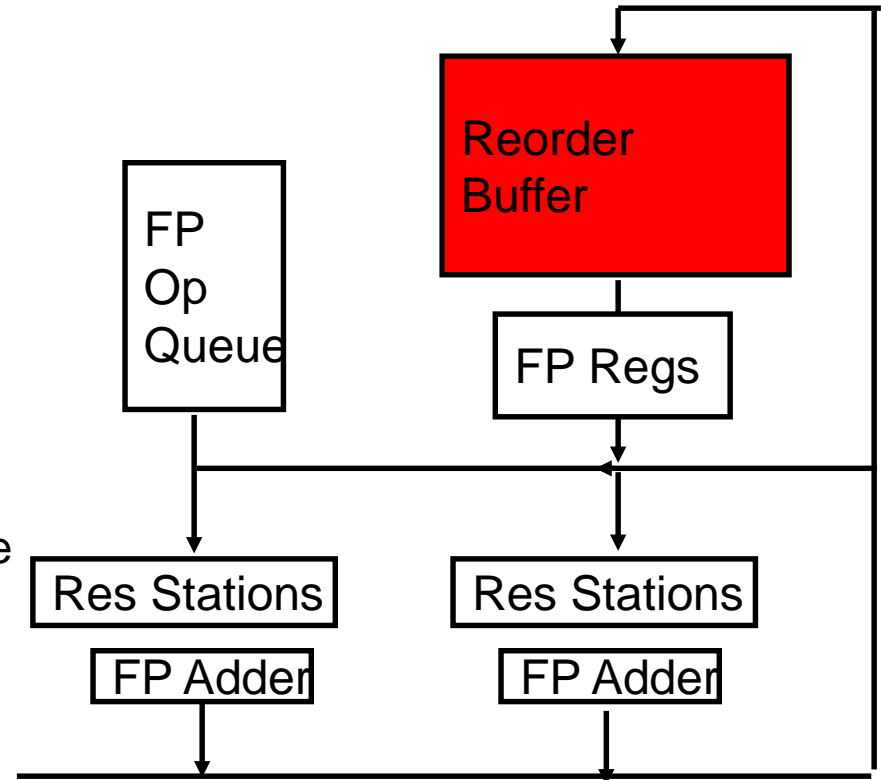
- Tomasulo had:

In-order issue, out-of-order execution, and out-of-order completion

- Need to “fix” the out-of-order completion aspect so that we can find precise breakpoint in instruction stream
- Actually we have the same issue with branch speculation...
- The answer: add a stage that “commits” the state
- In issue order

HW support for precise interrupts

- Need HW buffer for results of uncommitted instructions: *reorder buffer*
 - 3 fields: instr, destination, value
 - Use reorder buffer number instead of reservation station when execution completes
 - Supplies operands between execution complete & commit
 - (Reorder buffer can be operand source => more registers like RS)
 - Instructions commit
 - Once instruction commits, result is put into register
 - As a result, easy to undo speculated instructions on mispredicted branches or exceptions



Four Steps of the Speculative Tomasulo Algorithm

1. Issue—get instruction from FP Op Queue

If reservation station **and reorder buffer slot** free, issue instr & send operands **& reorder buffer no. for destination** (this stage sometimes called “dispatch”)

2. Execution—operate on operands (EX)

When both operands ready then execute; if not ready, watch CDB for result; when both in reservation station, execute; checks RAW (sometimes called “issue”)

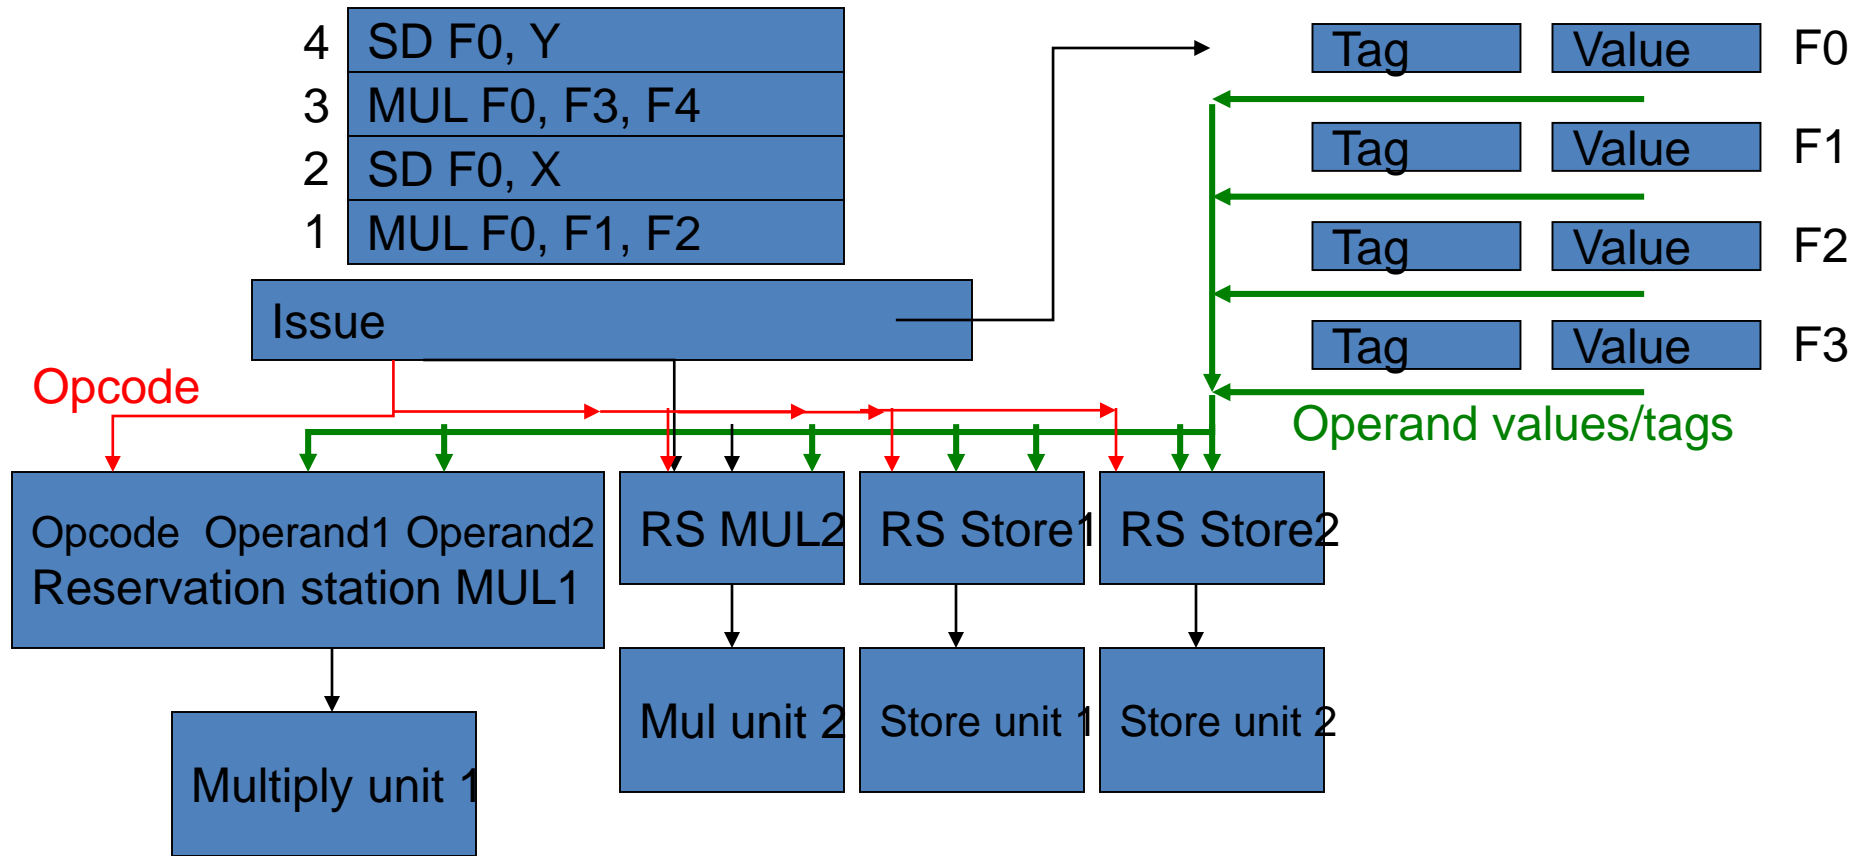
3. Write result—finish execution (WB)

Write on Common Data Bus to all awaiting FUs **& reorder buffer**; mark reservation station available.

4. Commit—**update register with reorder result**

When instr. at head of reorder buffer & result present, update register with result (or store to memory) and remove instr from reorder buffer. Mispredicted branch flushes reorder buffer (sometimes called “graduation”)

Tomasulo without Re-order Buffer

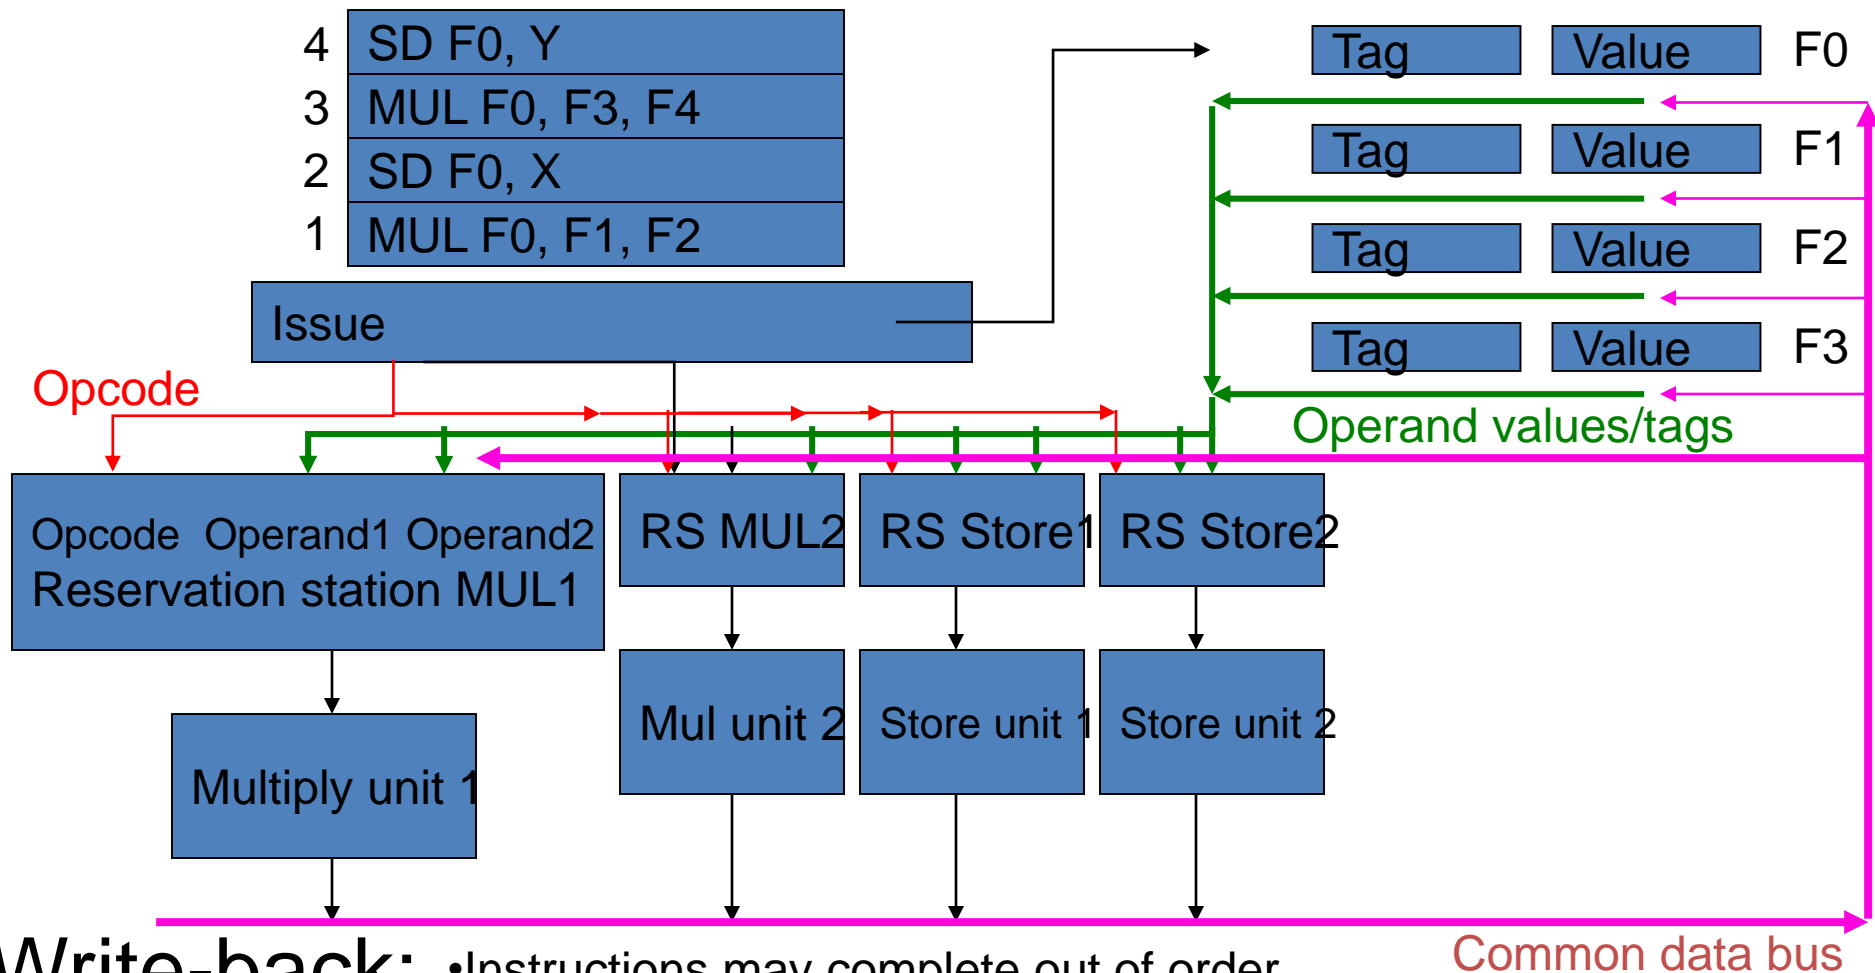


Issue:

- Each instruction is issued in order
- Issue unit collects operands from the two instruction's source registers
- Result may be a value, or, if value will be computed by an uncompleted instruction, the tag of the RS to which it was issued.
- When instruction 1 is issued, F0 is updated to get result from MUL1
- When instruction 3 is issued, F0 is updated to get result from MUL2

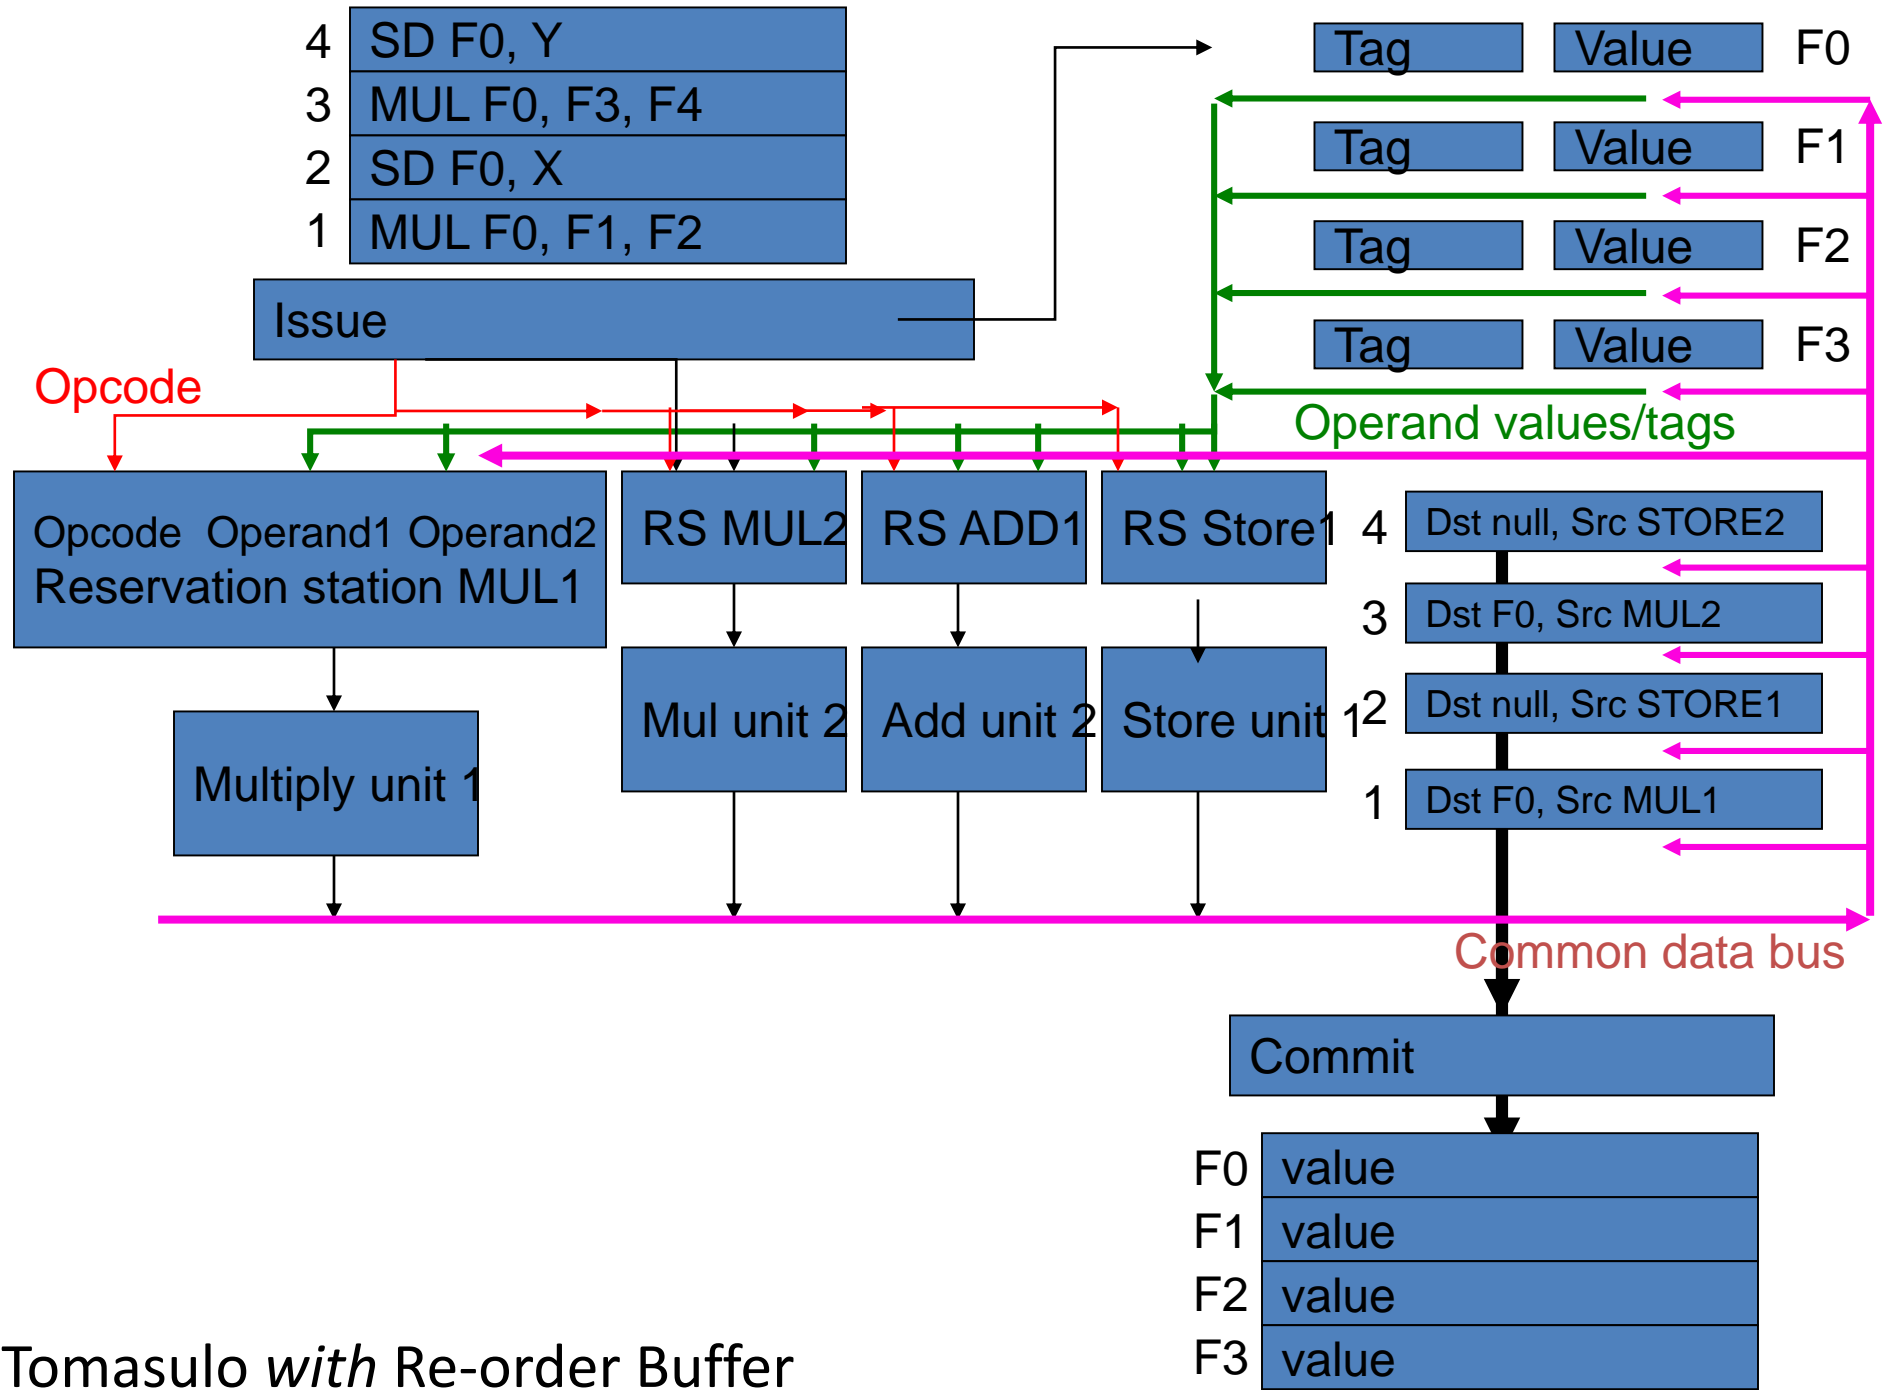
(Copy of
earlier
slide)

Tomasulo without Re-order Buffer

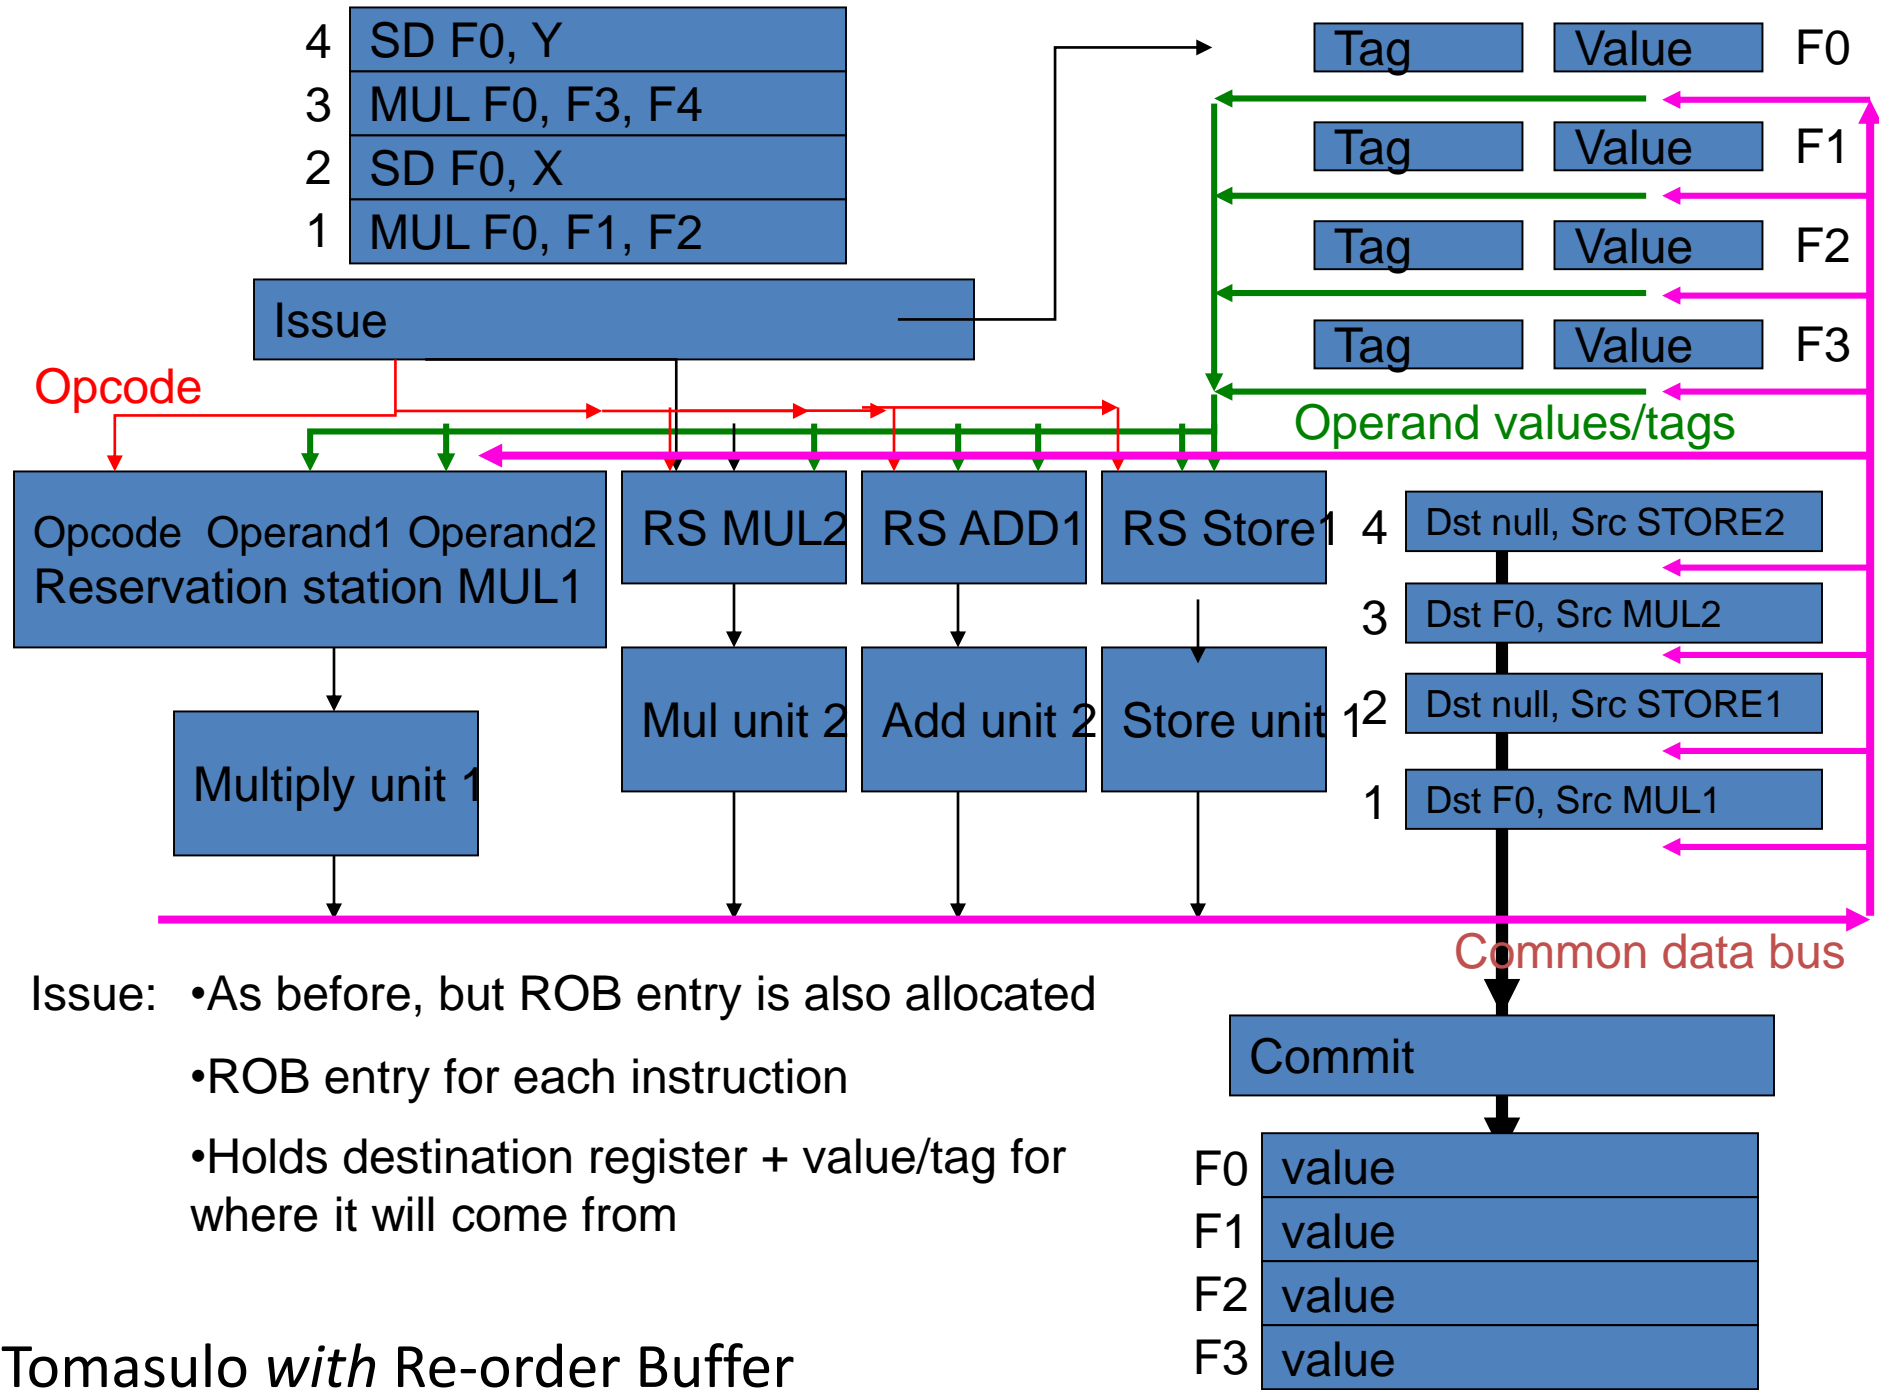


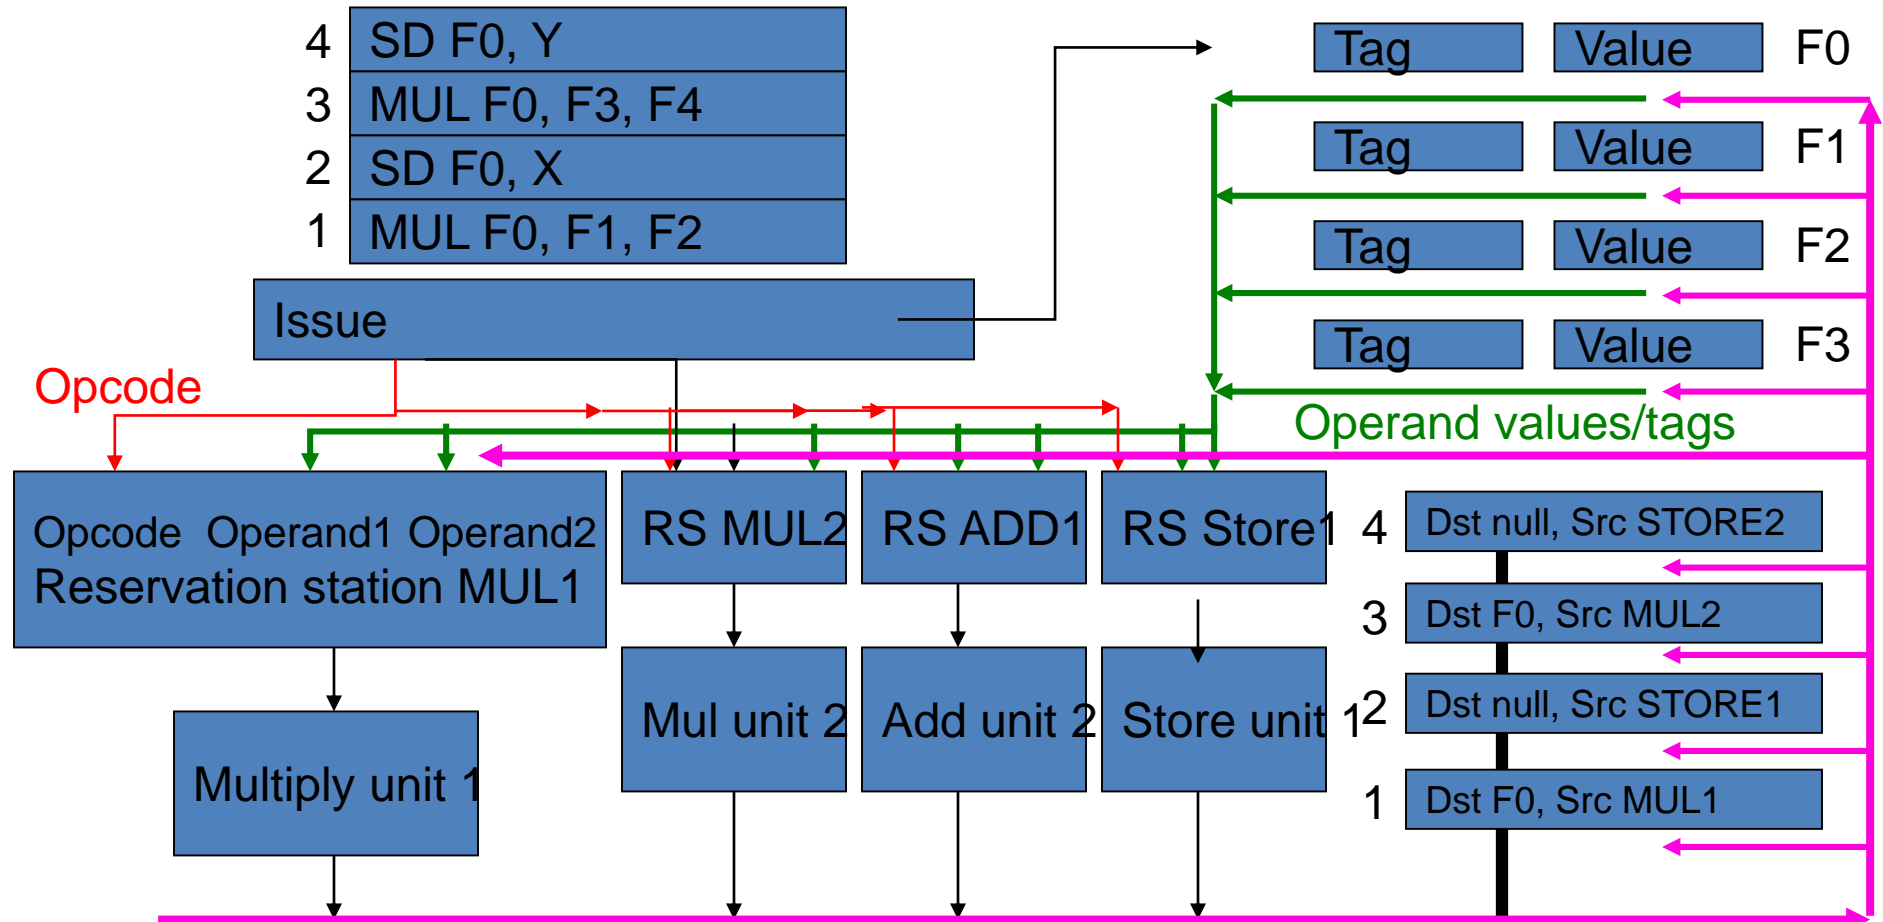
- Write-back:**
- Instructions may complete out of order
 - Result is broadcast on CDB
 - Carrying tag of RS to which instruction was originally issued
 - All RSs and registers monitor CDB and collect value if tag matches
 - Any RS which has both operands and whose FU is free fires.
 - When MUL1 completes result goes to store unit but not F0

(Copy of
earlier
slide)



Tomasulo *with* Re-order Buffer

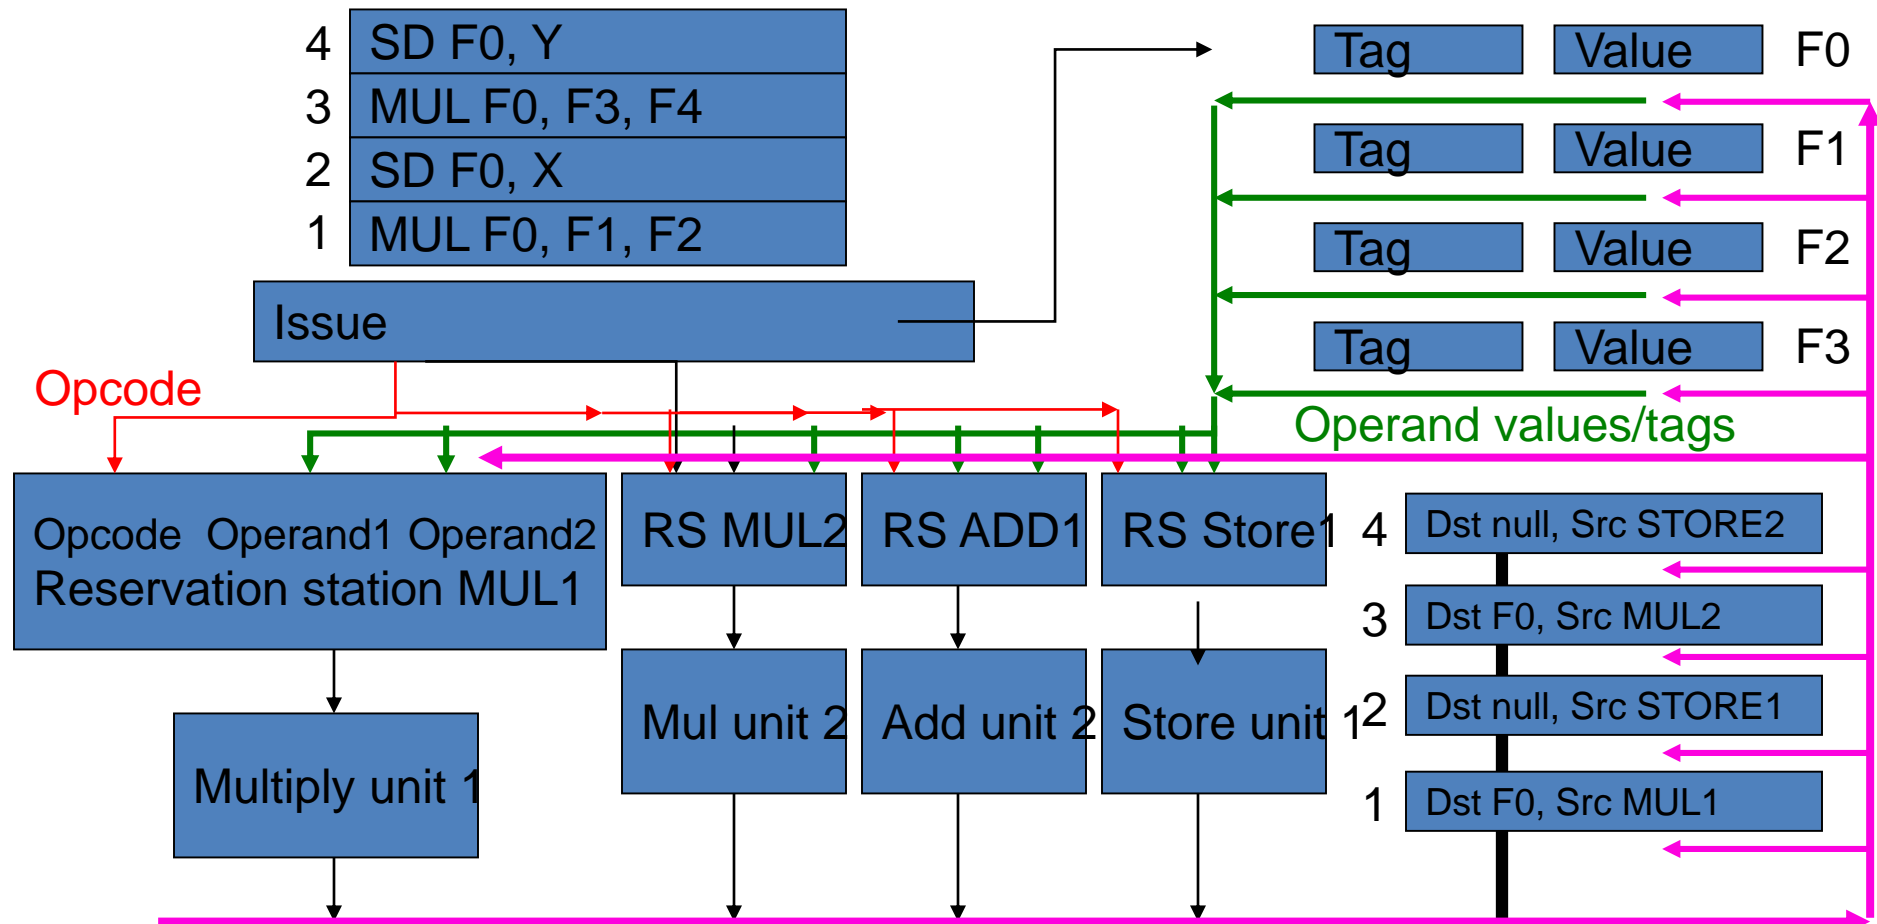




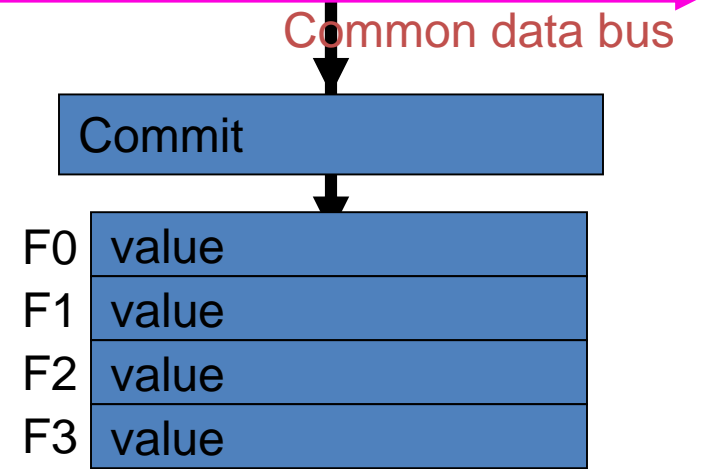
Write Back:

- As before, but ROB entry with matching tag also updated
- ROB entry for instruction 1 holds value for F0
- ROB entry for instruction 3 holds another value for F0

Tomasulo *with* Re-order Buffer



- Commit:
- Commit unit processes ROB entries in issue order
 - Each instruction waits in turn and commits when its operands are completed
 - Committed registers updated with values from ROB
 - F0 is updated first with result from MUL1 then result from MUL2



4	SD F0, Y
3	MUL F0, F3, F4
2	SD F0, X
1	MUL F0, F1, F2

Issue

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Operand values/tags

Opcode

Opcode	Operand1	Operand2
Reservation station	MUL1	

RS	MUL2
----	------

RS	ADD1
----	------

RS	Store1
----	--------

4	Dst null, Src STORE2
3	Dst F0, Src MUL2
12	Dst null, Src STORE1
1	Dst F0, Src MUL1

Multiply unit 1

Mul unit 2

Add unit 2

Store unit

Common data bus

Commit

F0	value
F1	value
F2	value
F3	value

Issue-side registers
(updated speculatively)

Commit-side registers
(updated when speculation resolved)

5	SD F0, Y
4	MUL F0, F3, F4
3	BEQ R10, Lab
2	SD F0, X
1	MUL F0, F1, F2

Issue

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Opcode

Operand values/tags

Opcode Operand1 Operand2
Reservation station MUL1

RS MUL2

RS ADD1

RS Store1

5 Dst null, Src STORE2

4 Dst F0, Src MUL2

3 BEQ R10, Lab (predNT)

2 Dst null, Src STORE1

1 Dst F0, Src MUL1

Multiply unit 1

Mul unit 2

Add unit 2

Store unit 1

Commit

F0	value
F1	value
F2	value
F3	value

- Now extend example with conditional branch
- Assume predicted Not Taken
- When BEQ reaches head of commit queue, all instructions which have been issued but have not yet committed are erroneous

5	SD F0, Y
4	MUL F0, F3, F4
3	BEQ R10, Lab
2	SD F0, X
1	MUL F0, F1, F2

Issue

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Operand values/tags

Opcode

Opcode Operand1 Operand2
Reservation station MUL1

RS MUL2

RS ADD1

RS Store1

Multiply unit 1

Mul unit 2

Add unit 2

Store unit 1

~~5 Dst null, Src STORE2~~

~~4 Dst F0, Src MUL2~~

3 BEQ R10, Lab (predNT)

Commit

F0	Value from MUL1
F1	value
F2	value
F3	value

- Misprediction: all ROB entries are trashed
- Issue-side registers reset from commit-side registers
- Correct branch target instruction fetched and issued

5	SD F0, Y
4	MUL F0, F3, F4
3	BEQ R10, Lab
2	SD F0, X
1	MUL F0, F1, F2

Issue

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Operand values/tags

Opcode

Opcode Operand1 Operand2
Reservation station MUL1

RS MUL2

RS ADD1

RS Store1

Multiply unit 1

Mul unit 2

Add unit 2

Store unit 1

5	Dst null, Src STORE2
4	Dst F0, Src MUL2
3	BEQ R10, Lab (predNT)

Commit

F0	Value from MUL1
F1	value
F2	value
F3	value

•Committed F0 holds value from first MUL

•RS of uncompleted speculatively-executed instruction cannot be re-used until its FU (eg MUL2) completes

Some subtleties...

- It's vital to reduce the branch misprediction penalty. Does the Tomasulo+ROB scheme described here roll-back as soon as the branch is found to be mispredicted?
- Stores are buffered in the ROB, and committed only when the instruction is committed. A load can be issued while several stores (perhaps to the same address) are uncommitted. We need to make sure the load gets the right data. See:
http://home.eng.iastate.edu/~zzhang/courses/cpre585_f03/slides/lecture11.pdf
- What if a second conditional branch is encountered, before the outcome of the first is resolved?
- This discussion has assumed a single-issue machine. How can these ideas be extended to allow multiple instructions to be issued per cycle?
 - Issue
 - Monitoring CDBs for completion
 - Handling multiple commits per cycle

8	SD F0, Y
7	MUL F0, F3, F4
6	BEQ R11, Lab
5	SD F0, Y
4	MUL F0, F3, F4
3	BEQ R10, Lab
2	SD F0, X
1	MUL F0, F1, F2

Issue

Opcode

Opcode Operand1 Operand2
Reservation station MUL1

Multiply unit 1

RS MUL2

Mul unit 2

RS ADD1

Add unit 2

RS Store1

Store unit 1

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Operand values/tags

8	Dst null, Src STORE2
7	Dst F0, Src MUL2
6	BEQ R11, Lab (predNT)
5	Dst null, Src STORE2
4	Dst F0, Src MUL2
3	BEQ R10, Lab (predNT)
2	Dst null, Src STORE1
1	Dst F0, Src MUL1

Commit

F0	value
F1	value
F2	value
F3	value

- Two conditional branches
- We speculate on *both* branches

8	SD F0, Y
7	MUL F0, F3, F4
6	BEQ R11, Lab
5	SD F0, Y
4	MUL F0, F3, F4
3	BEQ R10, Lab
2	SD F0, X
1	MUL F0, F1, F2

Issue

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Opcode

Operand values/tags

Opcode Operand1 Operand2
Reservation station MUL1

RS MUL2

RS ADD1

RS Store1

Multiply unit 1

Mul unit 2

Add unit 2

Store unit 1

8	Dst null, Src STORE2
7	Dst F0, Src MUL2
6	BEQ R11, Lab (predNT)
5	Dst null, Src STORE2
4	Dst F0, Src MUL2
3	BEQ R10, Lab (predNT)

Commit

F0	value
F1	value
F2	value
F3	value

- Two conditional branches
- When we come to commit the first branch we discover it was mispredicted

8	SD F0, Y
7	MUL F0, F3, F4
6	BEQ R11, Lab
5	SD F0, Y
4	MUL F0, F3, F4
3	BEQ R10, Lab
2	SD F0, X
1	MUL F0, F1, F2

Issue

Tag	Value	F0
Tag	Value	F1
Tag	Value	F2
Tag	Value	F3

Opcode

Operand values/tags

Opcode Operand1 Operand2
Reservation station MUL1

RS MUL2

RS ADD1

RS Store1

Multiply unit 1

Mul unit 2

Add unit 2

Store unit 1

8	Dst null, Src STORE2
7	Dst F0, Src MUL2
6	BEQ R11, Lab (predNT)
5	Dst null, Src STORE2
4	Dst F0, Src MUL2
3	BEQ R10, Lab (predNT)

Commit

F0	value
F1	value
F2	value
F3	value

- When we come to commit the first branch we discover it was mispredicted
- We squash all the issued instructions including the second branch

Tomasulo + ROB: Summary

- Reservations stations: *implicit register renaming* to larger set of registers + buffering source operands
 - Prevents registers as bottleneck
 - Avoids WAR, WAW hazards of Scoreboard (see textbook)
 - Allows loop unrolling in HW
- Not limited to basic blocks
(integer units gets ahead, beyond branches)
- Today, helps cache misses as well
 - Don't stall for L1 Data cache miss (insufficient ILP for L2 miss?)
- Lasting Contributions
 - Dynamic scheduling
 - Register renaming
 - Load/store disambiguation

Design alternatives for o-o-o processor architectures

- See:
 - The Microarchitecture of the Pentium 4 Processor (Hinton et al, Intel Tech Jnl Q1 2001)
 - The SimpleScalar Tool Set, Version 2.0 (Burger and Austin, http://www.simplescalar.com/docs/users_guide_v2.pdf)
 - Wattch: a framework for architectural-level power analysis and optimizations (Brooks et al, ISCA 2000) www.tortolaproject.com/papers/brooks00wattch.pdf
- *Specifically:*
 - *Register Update Unit (RUU, as in SimpleScalar) versus Re-Order Buffer*
 - *Realisation in Pentium III and Pentium 4 (“Netburst”)*
 - *Frontend and Retirement Register Alias Tables (RATs)*

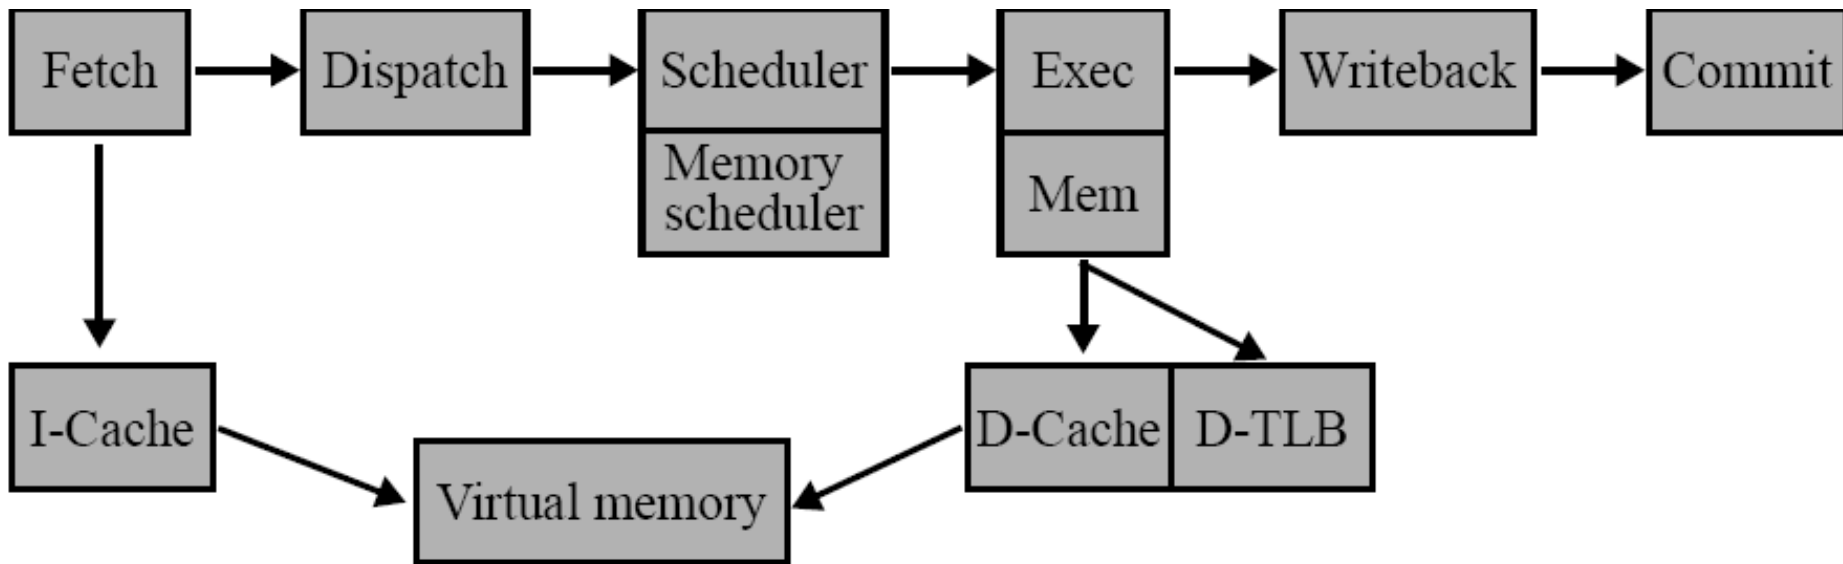
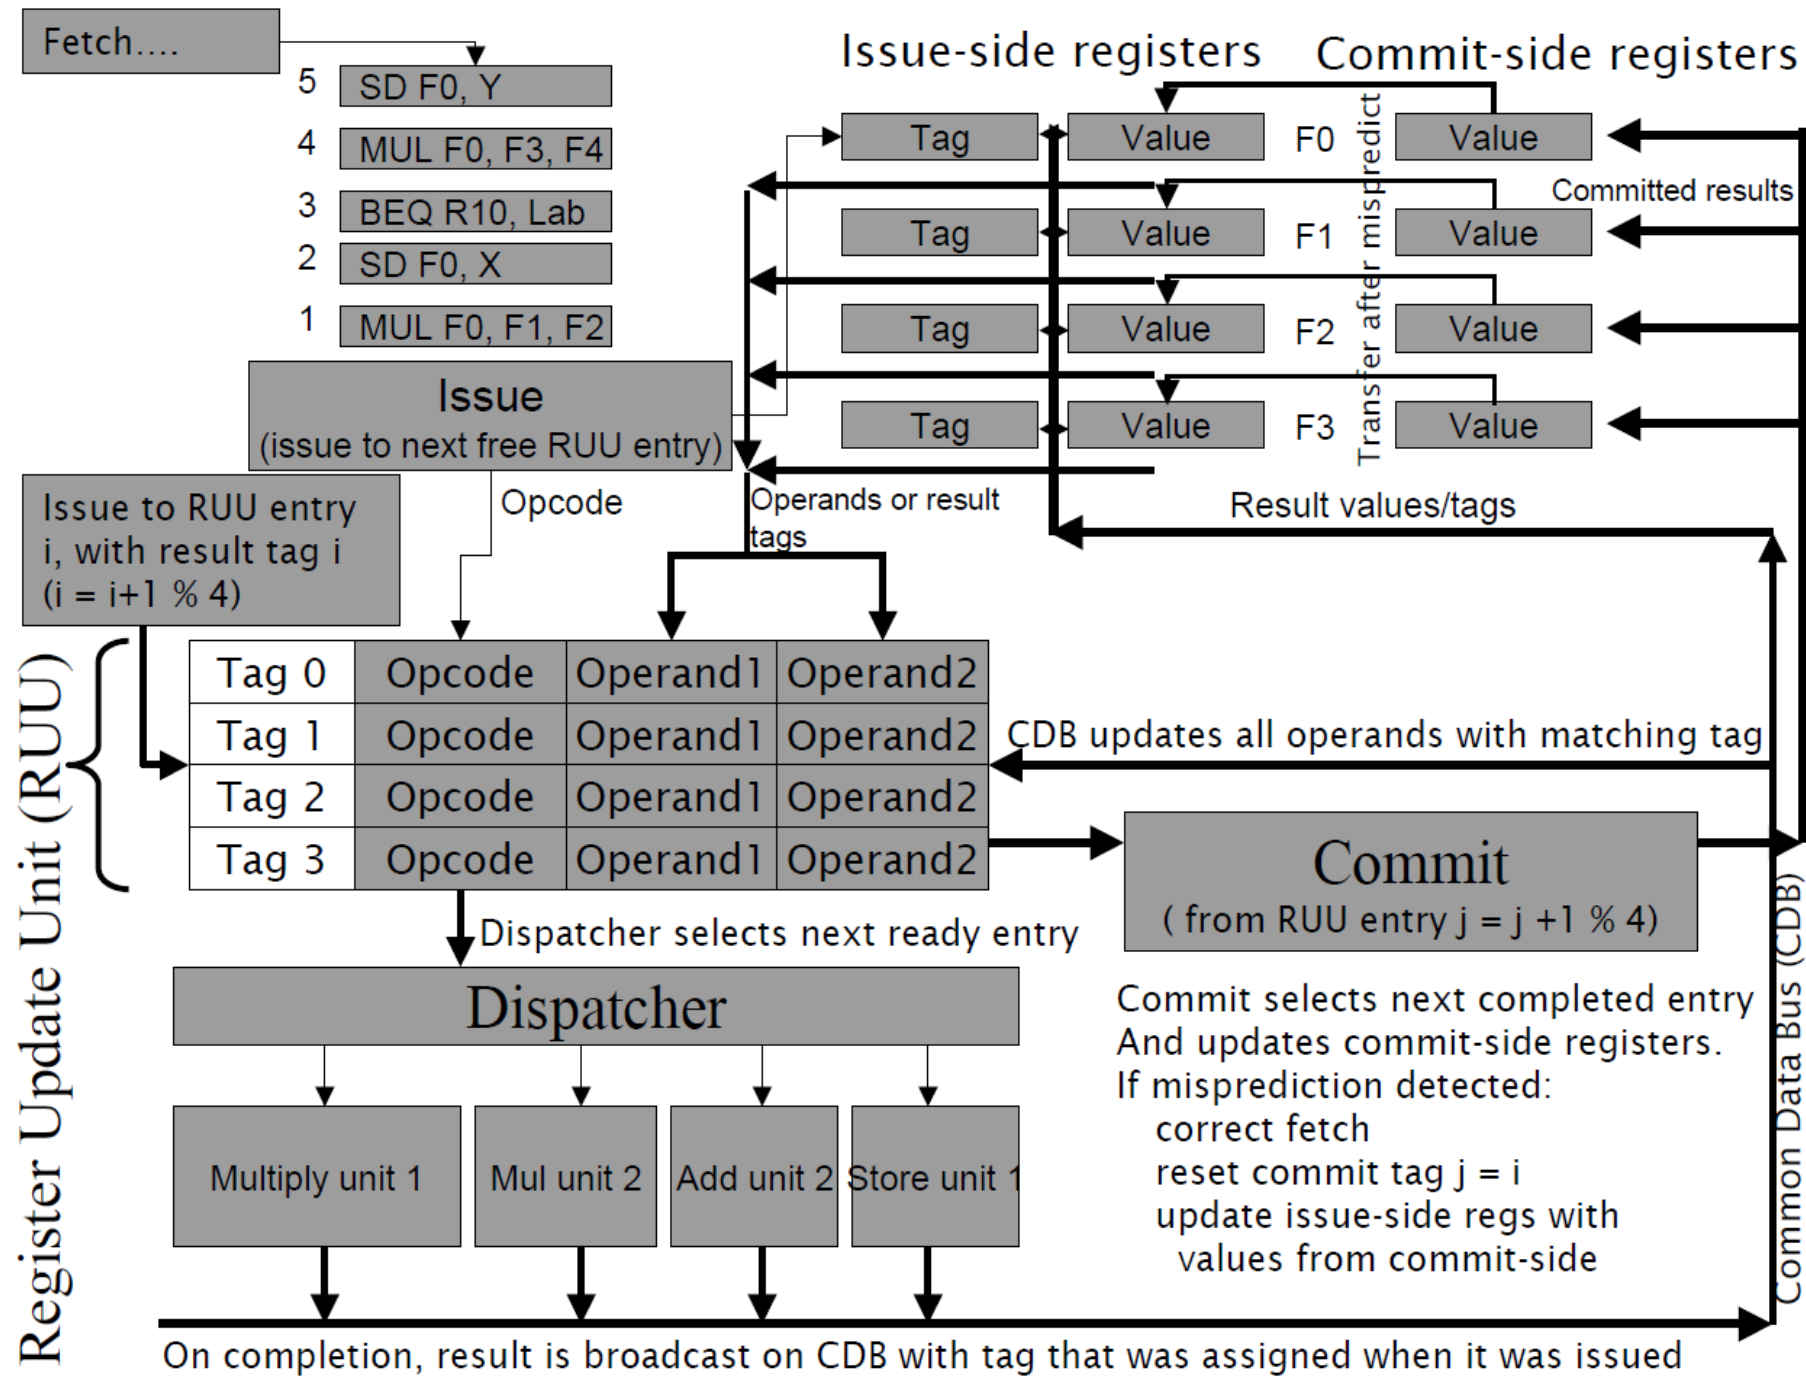
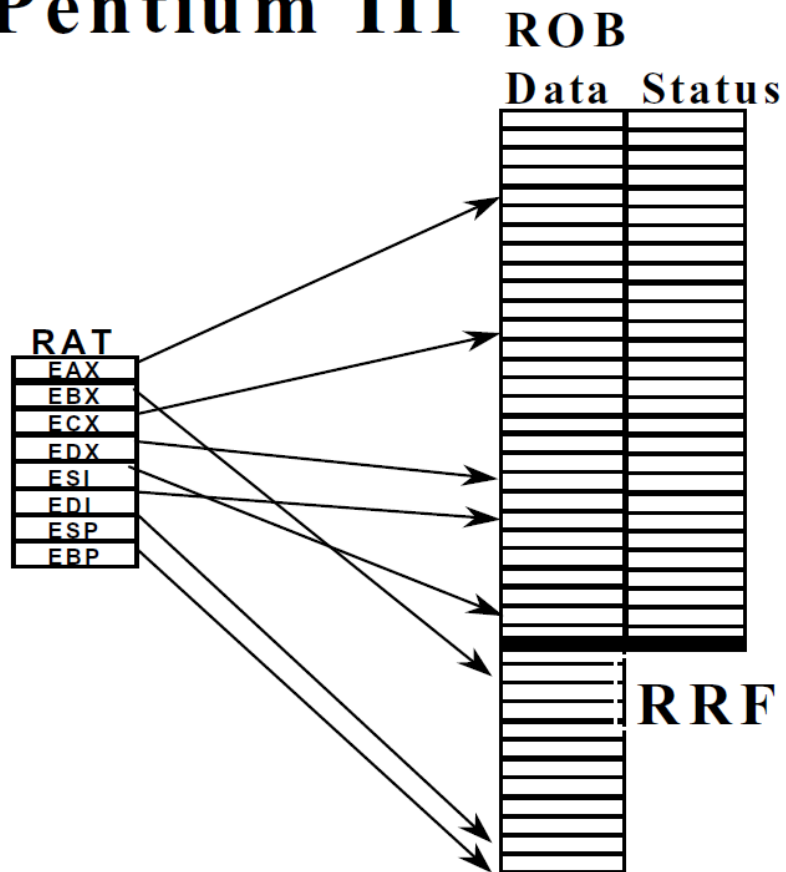


Figure 5. Pipeline for sim-outorder

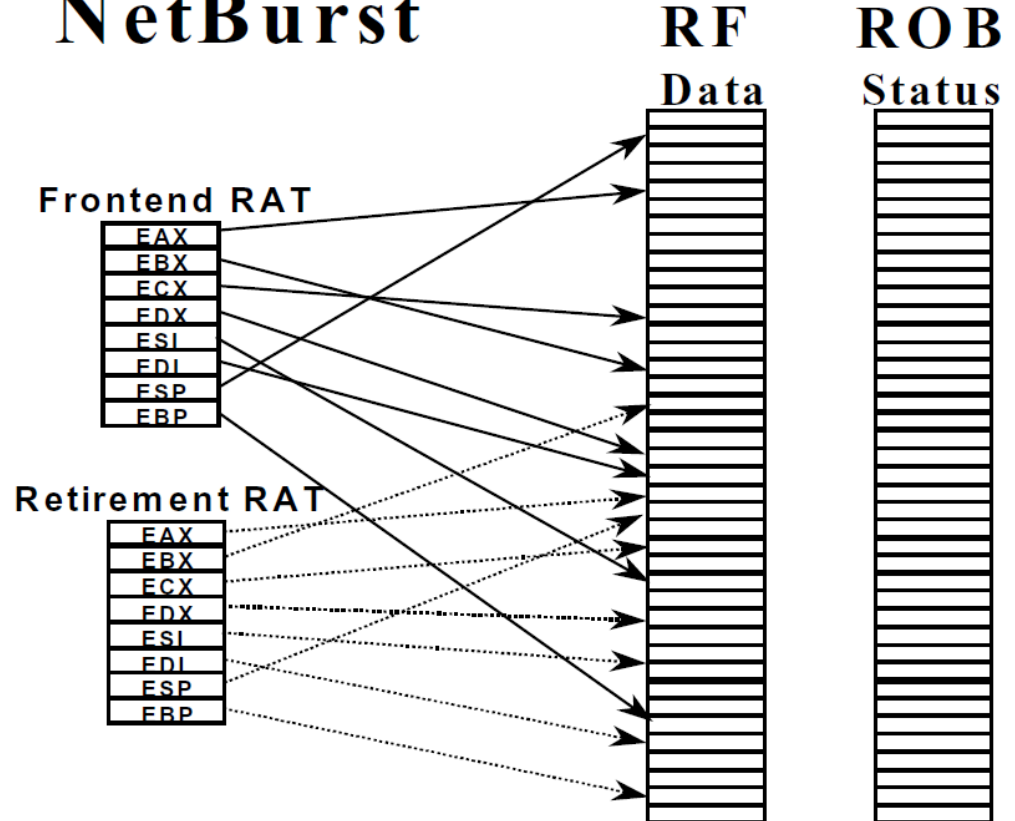
- Simplescalar is a software simulation of a processor microarchitecture
- It simulates a multi-issue out-of-order design with speculative execution
- Many aspects of the design can be controlled by parameters
- Simplescalar uses a Register Update Unit, which combines ROB and reservation stations in a single pool



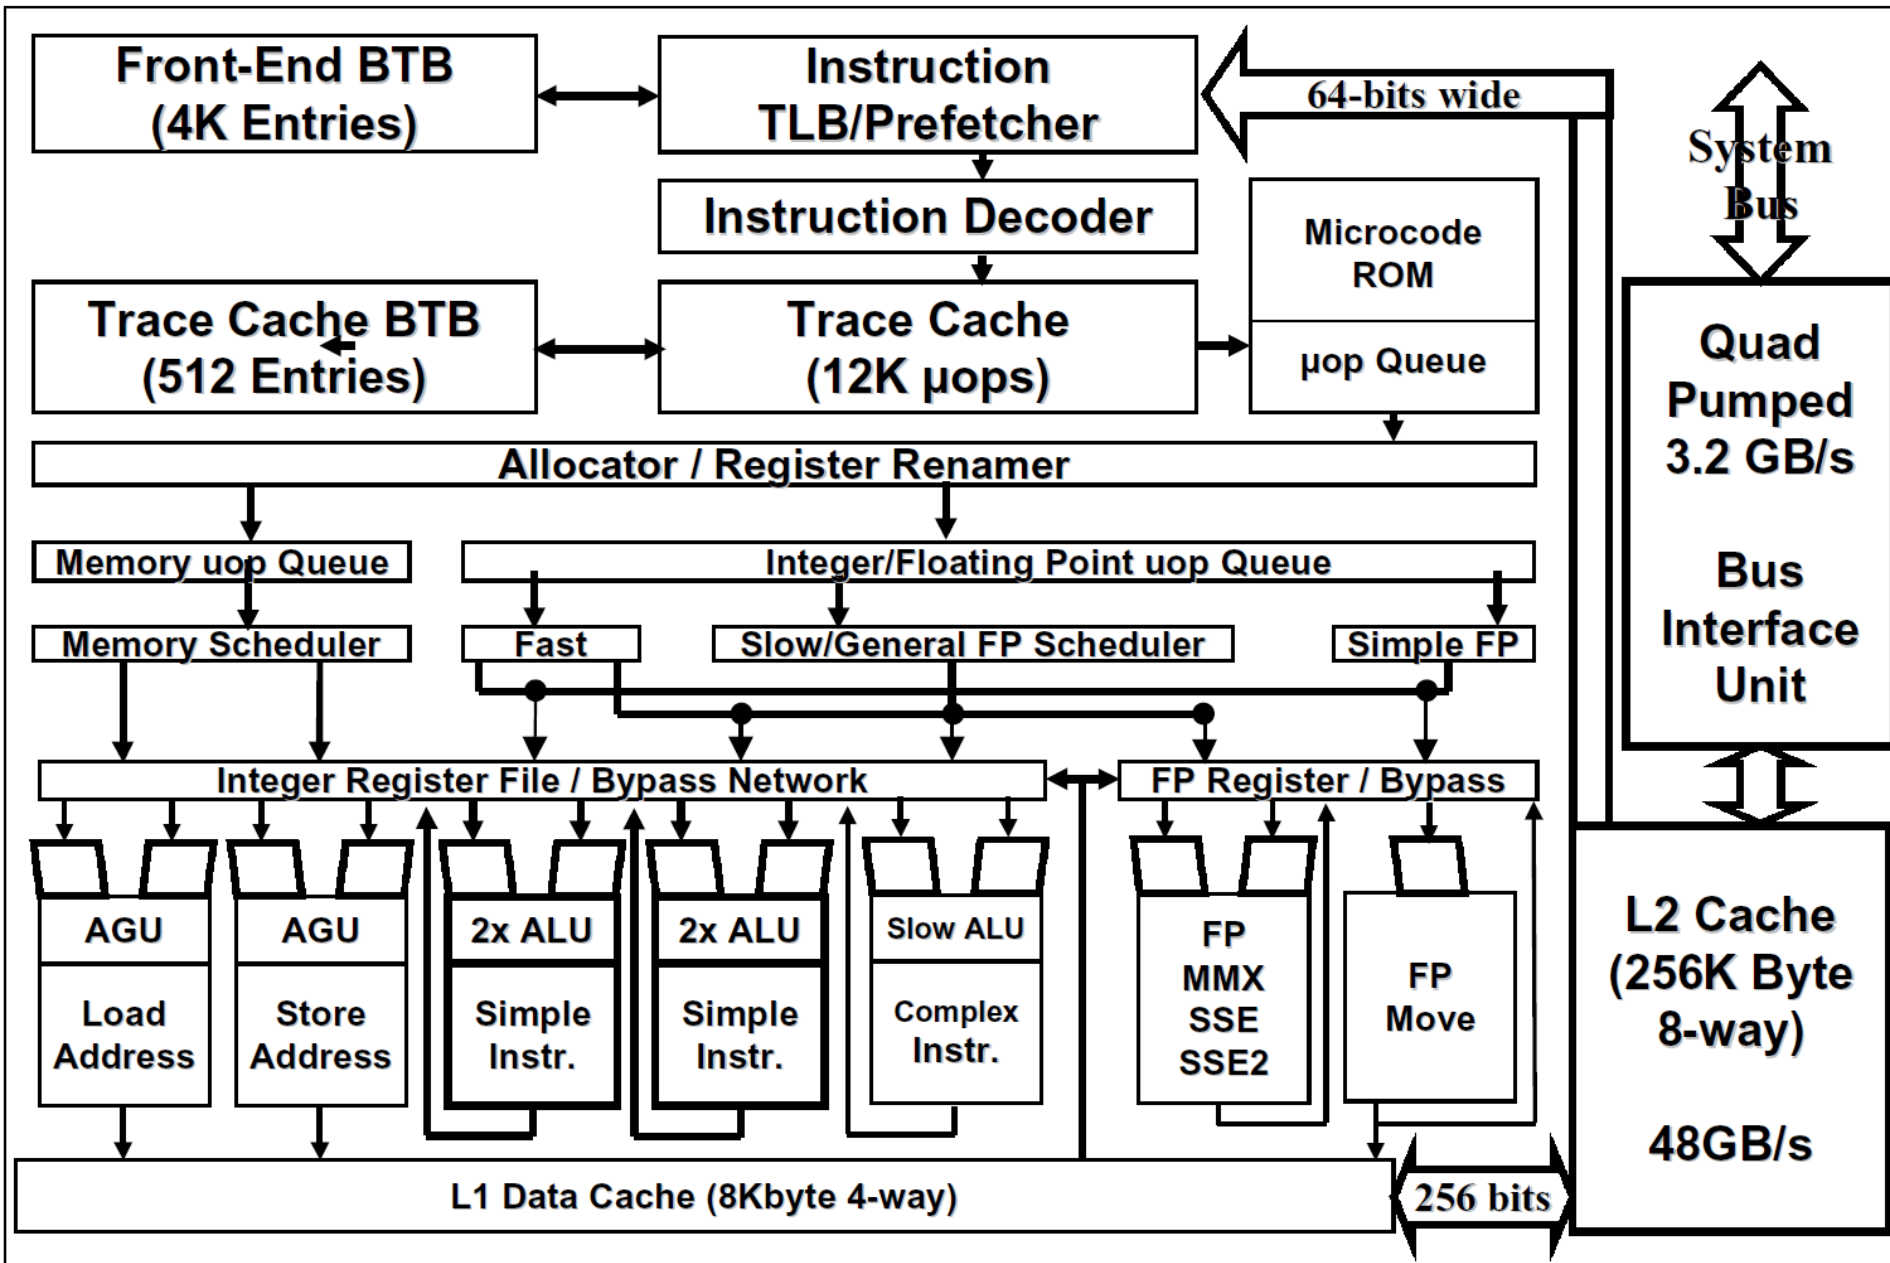
Pentium III



NetBurst



Pentium 4 ("Netburst") microarchitecture



Basic Pentium III Processor Misprediction Pipeline

1	2	3	4	5	6	7	8	9	10
Fetch	Fetch	Decode	Decode	Decode	Rename	ROB Rd	Rdy/Sch	Dispatch	Exec

Basic Pentium 4 Processor Misprediction Pipeline

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
TC Nxt IP		TC Fetch		Drive	Alloc	Rename		Que	Sch	Sch	Sch	Disp	Disp	RF	RF	Ex	Flgs	Br Ck	Drive


Out-of-order processing – Four instructions per cycle

Example: **Naive implementation (roughly from `cc -S`):**

```
void f() {  
    int i, a;  
    for (i=1;  
        i<=1000000000000;  
        i++)  
        a = a+i;  
}
```

**Real
example**

X86 code (slightly
tidied but without
register allocation)



```
movl $1,-4(%ebp)  
jmp .L4  
.L5  
    movl -4(%ebp),%eax  
    addl %eax,-8(%ebp)  
    incl -4(%ebp)  
.L4:  
    cmpl $1000000000000,-4(%ebp)  
    jle .L5
```

Unoptimised:

```
movl $1,-4(%ebp)
```

```
jmp .L4
```

```
.L5
```

```
movl -4(%ebp),%eax
```

```
addl %eax,-8(%ebp)
```

```
incl -4(%ebp)
```

```
.L4:
```

```
cmpl $1000000000,-4(%ebp)
```

```
jle .L5
```

5 instructions in the loop

Execution time on 2.13GHz Intel
Core2Duo: 3.87 seconds (3.87
nanoseconds/iteration, 8.24 cycles)

Optimised:

```
movl $1,%edx
```

```
.L6:
```

```
addl %edx,%eax
```

```
incl %edx
```

```
cmpl $1000000000,%edx
```

```
jle .L6
```

4 instructions in the loop, no
references to main memory

Execution time on 2.13GHz
Intel Core2Duo: 0.48 seconds
(0.48 nanoseconds/iteration,
1.02 cycles)

Time per instruction fell: 0.77 nanoseconds to 0.12

Optimised code runs at four instructions per cycle

Resources

- Wikipedia (!):

- http://en.wikipedia.org/wiki/Register_renaming

- Papers:

- Instruction issue logic for high-performance, interruptable pipelined processors. G. S. Sohi, S. Vajapeyam. International Conference on Computer Architecture, 1987 (<http://doi.acm.org/10.1145/30350.30354>)

- Towards Kilo-instruction processors. Cristal, Santana, Valero, Martinez ACM Trans. Architecture and Code Optimization (<http://doi.acm.org/10.1145/1044823.1044825>)

- Other simulators:

- SimpleScalar: www.simplescalar.com/

- Gem5: <http://www.gem5.org>

- Liberty: <http://liberty.cs.princeton.edu/>

- SimFlex: <http://parsa.epfl.ch/simflex/>

- SIMICS: <http://www.windriver.com/products/simics/>

Dynamic scheduling - summary

- Dynamic instruction scheduling is attractive:
 - Reduced dependence on compile-time instruction scheduling (and compiler knowledge of hardware)
 - Handles dynamic stalls due to cache misses
 - Register renaming frees architecture from constraints of the instruction set
- Comes with costs
 - Increases pipeline depth, and misprediction latency
 - Increased power consumption and area (but not by all that much if you are careful and clever)
 - Increased complexity and risk of design error
 - Hard to predict performance, hard to optimise code