

Registers are just arrows

Discussion exercise

How can we design our ISA without registers?

Can we simplify the microarchitecture as a result?

Perhaps encoding dependence directly (so it doesn't have to be discovered)

Perhaps avoiding register renaming?

Code is just an encoding for a graph

Load r1 A
Load r2 C
Add r3 r2 #1 // $r3 = C + 1$
Mul r4 r1 r3 // $r4 = r3 * A$
Store r4 D // $D = r4$

Register machine

Push A // load from location A, push onto stack
Push B
Add #1 // add 1 to the value on the top of the stack
Mul // multiply the top two items on the stack
Store // store the value on the top of the stack to memory

Stack machine

Load +3 A // load from location A, send to Mul instruction below
Load +1 C // load from location C, send to Add instruction below
Add +1 #1 // add 1 to the value received, send result to next
Mul +1 // multiply value received by from Add by C, send to store
Store // store the value received from instruction above

dataflow machine

Load A // Load A, drop the value on the conveyor “belt”
Load C // Load C, drop the value on the conveyor “belt”
Add -1 #1 // add 1 to the result of the instruction one instruction earlier
Mul -1 -3 // multiply the results from positions -1 and -3 on the belt
Store -1 // store the result of the multiply to memory


“belt” machine

Belt addressing

Belt operands are addressed by relative position

add b3, b5

No result address!



“b3” is the fourth most recent value to drop to the belt

“b5” is the sixth most recent value to drop to the belt



This is temporal addressing



Wavescalar's ISA encodes graph's forward arrows


```
int *V;  
int a, b;  
int c, d, r;
```

```
r = a*c + b*d;  
V[a] = 2*r + d << 2;
```



- **Three views of code in WaveScalar:** At left is the C code for a simple computation. The WaveScalar dataflow graph is shown at center, and the same graph is mapped onto 2 8-PE domains in the WaveCache substrate at right.

Wavescalar's ISA encodes graph's forward arrows



- Graph fragments are mapped onto clustered on-chip array of processing elements
- Each processing element waits for its operands
- Operands are tagged – processing element fires when it receives operands with matching tags

References:

- Sinha, Steve & Chatterjee, Satrajit & Ravindran, Kaushik. (2019). **BOOST: Berkeley's Out-of-Order Stack Thingy**.
https://www.researchgate.net/publication/228556746_BOOST_Berkeley's_Out-of-Order_Stack_Thingy
- Mill Computing: <https://millcomputing.com/docs/>
- Schmit, Herman; Benjamin Levine; Benjamin Ylvisaker (2002). **Queue Machines: Hardware Compilation in Hardware**. *10th Annual IEEE Symposium on Field-Programmable Custom Computing Machines (FCCM'02)*: 152. [doi:10.1109/FPGA.2002.1106670](https://doi.org/10.1109/FPGA.2002.1106670)
- Steven Swanson, Andrew Putnam, Martha Mercaldi, Ken Michelson, Andrew Petersen, Andrew Schwerin, Mark Oskin, and Susan J. Eggers. 2006. **Area-Performance Trade-offs in Tiled Dataflow Architectures**. In Proceedings of the 33rd annual international symposium on Computer Architecture (ISCA '06). IEEE Computer Society, Washington, DC, USA, 314-326. DOI: <https://doi.org/10.1109/ISCA.2006.10>
- INTEL'S EXASCALE DATAFLOW ENGINE DROPS X86 AND VON NEUMANN. Aug 2018, Timothy Prickett Morgan, <https://www.nextplatform.com/2018/08/30/intels-exascale-dataflow-engine-drops-x86-and-von-neuman/> . See also https://en.wikichip.org/wiki/intel/configurable_spatial_accelerator and US Patent No. US20180189231A1.