Advanced Computer Architecture Chapter 6

Instruction Level Parallelism
- static instruction scheduling, software pipelining, VLIW, EPIC, and multi-threading

February 2018
Paul H J Kelly

These lecture notes are partly based on the course text, Hennessy and Patterson's Computer Architecture, a quantitative approach (3rd and 4th eds), and on the lecture slides of David Patterson and John Kubiatowicz's Berkeley course

Overview

- Up to now: Dynamic scheduling, out-of-order (o-o-o): binary compatible, exploiting ILP in hardware: BTB, ROB, Reservation Stations, ...
- How far can you take it?
- How much of all this complexity can you shift into the compiler?
- What if you can also change instruction set architecture?
- VLIW (Very Long Instruction Word)
- EPIC (Explicitly Parallel Instruction Computer)
 - Intel's (and HP's) multi-billion dollar gamble for the future of computer architecture: Itanium, IA-64
 - Started ca.1994...not dead yet but has it turned a profit?
- Beyond instruction-level parallelism...

Running Example

This code adds a scalar to a vector:

for (i=1000; i>=0; i=i-1)
$$x[i] = x[i] + s;$$

Assume following latency all examples

Instruction producing result	Instruction using result	Execution in cycles	Latency in cycles
FP ALU op	Another FP ALU op	4	3
FP ALU op	Store double	3	2
Load double	FP ALU op	1	1
Load double	Store double	1	0
Integer op	Integer op	1	0

FP Loop: Where are the Hazards?

First translate into MIPS code:

-To simplify, assume 8 is lowest address

```
Loop: L.D F0,0(R1) ;F0=vector element
ADD.D F4,F0,F2 ;add scalar from F2
S.D 0(R1),F4 ;store result
DSUBUI R1,R1,8 ;decrement pointer 8B (DW)
BNEZ R1,Loop ;branch R1!=zero
NOP ;delayed branch slot
```

Where are the stalls?

FP Loop Showing Stalls

```
1 Loop: L.D F0,0(R1) ;F0=vector element
        stall
3
        ADD.D F4,F0,F2 ;add scalar in F2
        stall
4
5
        stall
6
        S.D
               0(R1), F4; store result
        DSUBUI R1,R1,8 ; decrement pointer 8B (DW)
8
               R1,Loop ;branch R1!=zero
        BNEZ
        stall
                         ;delayed branch slot
Instruction
                Instruction
                                    Latency in
producing result
                using result
                                    clock cycles
                Another FP ALU op
FP ALU op
                                    3
FP ALU op
                Store double
Load double
                FP ALU op
```

9 clocks: Rewrite code to minimize stalls?

Revised FP Loop Minimizing Stalls

```
1 Loop: L.D F0,0(R1)
2 stall
3 ADD.D F4,F0,F2
4 DSUBUI R1,R1,8
5 BNEZ R1,Loop ;delayed branch
6 S.D 8(R1),F4 ;altered when moved past DSUBUI
```

Swap BNEZ and S.D by changing address of S.D

```
Instruction Instruction Latency in producing result using result clock cycles

FP ALU op Another FP ALU op 3

FP ALU op Store double 2

Load double FP ALU op 1
```

6 clocks, but just 3 for execution, 3 for loop overhead; How make faster?

Unroll the loop four times

- Four copies of the loop body
- One copy of increment and test
- Adjust register-indirect loads using offsets

```
Loop: L.D
             F0,0(R1)
2
      ADD.D F4,F0,F2
3
      S.D
             0(R1),F4
                         ;drop DSUBUI & BNEZ
4
      L.D F0, -8(R1)
5
      ADD.D F4,F0,F2
6
      S.D -8 (R1), F4; drop DSUBUI & BNEZ
      L.D F0, -16(R1)
8
      ADD.D F4,F0,F2
9
      S.D -16(R1), F4 ; drop DSUBUI & BNEZ
10
      L.D F0, -24(R1)
11
      ADD.D F4,F0,F2
12
      S.D
             -24(R1), F4
13
      DSUBUI R1,R1,#32 ;alter to 4*8
14
      BNEZ
             R1,LOOP
15
      NOP
```

- Re-use of registers creates WAR ("anti-dependences")
 - How can we remove them?

Loop unrolling...

```
Loop: L.D F0, 0 (R1)
2
      ADD.D
             F4,F0,F2
3
      S.D
             0(R1),F4
                          ;drop DSUBUI & BNEZ
4
      L.D F6, -8(R1)
5
      ADD.D F8, F6, F2
6
       S.D
             -8(R1), F8
                          ;drop DSUBUI & BNEZ
      L.D F10, -16(R1)
8
      ADD.D F12,F10,F2
9
       S.D
             -16(R1), F12
                         ;drop DSUBUI & BNEZ
10
      L.D F14, -24(R1)
11
      ADD.D F16,F14,F2
12
       S.D
             -24(R1), F16
13
      DSUBUI R1,R1,#32 ;alter to 4*8
14
      BNEZ
             R1,LOOP
15
      NOP
```

The original "register renaming"

Unrolled Loop That Minimizes Stalls

```
Loop: L.D
              F0,0(R1)
2
       L.D F6, -8(R1)
3
       L.D F10, -16(R1)
4
       L.D
              F14, -24(R1)
5
       ADD.D F4,F0,F2
6
       ADD.D F8, F6, F2
7
       ADD.D F12,F10,F2
8
       ADD.D F16,F14,F2
9
       S.D
              0(R1),F4
10
       S.D -8 (R1), F8
11
              -16(R1),F12
       S.D
12
       DSUBUI R1, R1, #32
13
              R1,LOOP
       BNEZ
              8(R1), F16 ; 8-32 = -24
14
       S.D
```

What assumptions made when moved code?

- OK to move store past DSUBUI even though changes register
- OK to move loads before stores: get right data?
- When is it safe for compiler to do such changes?

14 clock cycles, or 3.5 per iteration

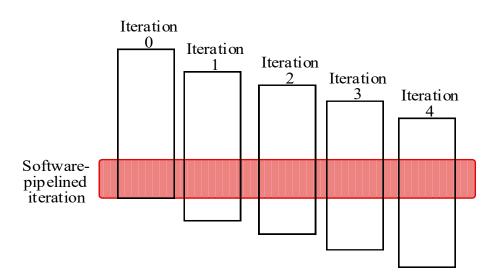
Compare: without scheduling

```
1 cycle stall
              F0,0(R1)
  Loop: L.D
                                  2 cycles stall
2
              F4,F0,F2
       ADD.D
3
       S.D
              0(R1),F4
                             ;drop DSUBUI & BNEZ
4
       L.D F6, -8(R1)
5
       ADD.D F8, F6, F2
6
       S.D
              -8(R1), F8
                             ;drop DSUBUI & BNEZ
7
       L.D F10, -16(R1)
8
       ADD.D F12,F10,F2
                                                         (use a fresh set of
9
       S.D -16(R1), F12
                             ;drop DSUBUI & BNEZ
                                                          registers for each
10
       L.D F14, -24(R1)
                                                          iteration to avoid
11
       ADD.D F16,F14,F2
                                                         WAR hazards/anti-
12
       S.D
               -24 (R1), F16
                                                           dependences)
13
       DSUBUI R1,R1,#32
                             ;alter to 4*8
14
              R1,LOOP
       BNEZ
15
       NOP
```

 $15 + 4 \times (1+2) = 27$ clock cycles, or 6.8 per iteration Assumes R1 is multiple of 4

Static overlapping of loop bodies: "Software Pipelining"

- Observation: if iterations from loops are independent, then can get more ILP by taking instructions from different iterations
- Software pipelining: reorganizes loops so that each iteration is made from instructions chosen from different iterations of the original loop (~ Tomasulo in software)



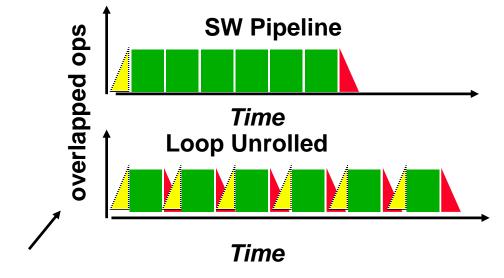
Software Pipelining Example

Before: Unrolled 3 times

```
1 L.D F0,0(R1)
2 ADD.D F4,F0,F2
3 S.D 0(R1),F4
4 L.D F6,-8(R1)
5 ADD.D F8,F6,F2
6 S.D -8(R1),F8
7 L.D F10,-16(R1)
8 ADD.D F12,F10,F2
```

After: Software Pipelined

```
S.D 0(R1),F4; Stores M[i]
ADD.D F4,F0,F2; Adds to M[i-1]
L.D F0,-16(R1); Loads M[i-2]
DSUBUI R1,R1,#8
BNEZ R1,LOOP
```



Symbolic Loop Unrolling

S.D -16(R1), F12

10 DSUBUI R1,R1,#24

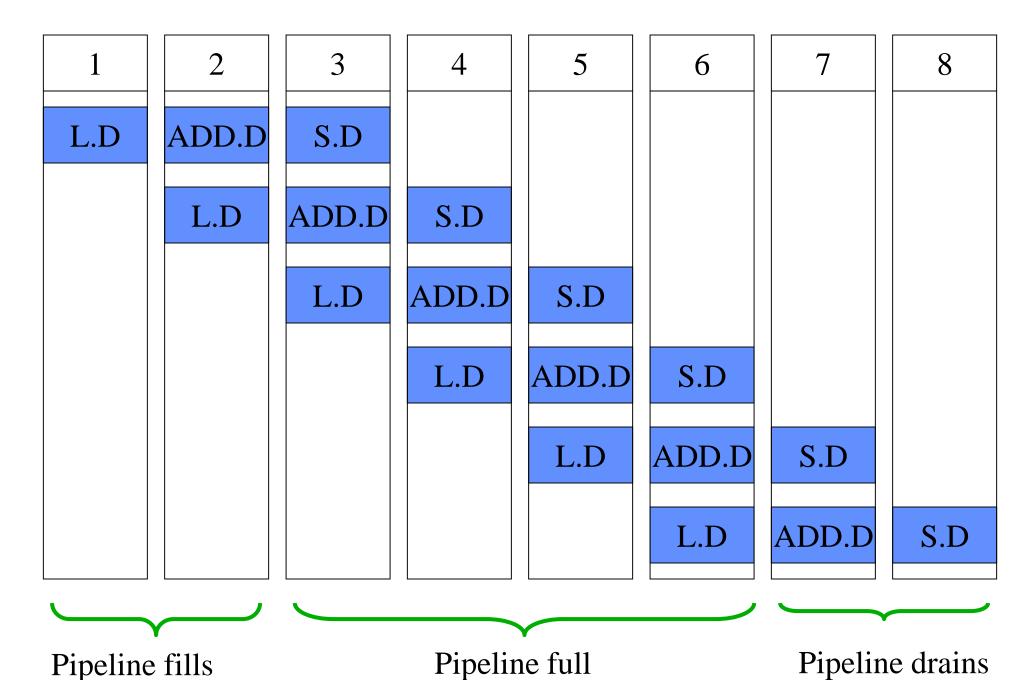
Maximize result-use distance

11 BNEZ R1,LOOP

- Less code space than unrolling
- Fill & drain pipe only once per loop
 vs. once per each unrolled iteration in loop unrolling

5 cycles per iteration

(3 if we can issue DSUBUI and BNEZ in parallel with other instrns)

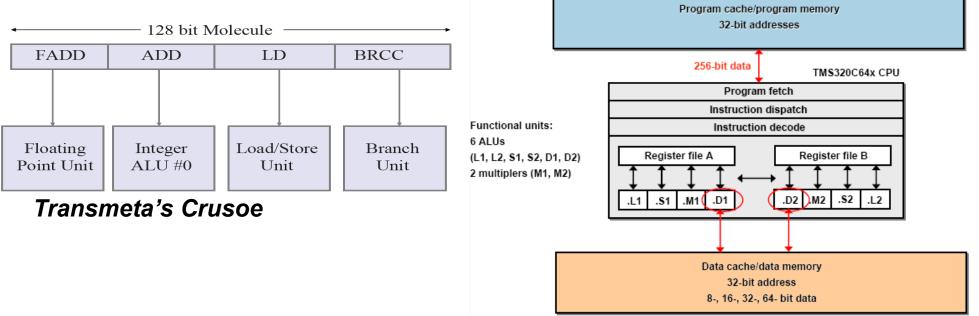


What if We Can Change the Instruction Set?

- Superscalar processors decide on the fly how many instructions to issue in each clock cycle
 - ▶ Have to check for dependences between all n pairs of instructions in a potential parallel issue packet
 - \Rightarrow Hardware complexity of figuring out the number of instructions to issue is $O(n^2)$
 - Entirely doable for smallish n, but tends to lead to multiple pipeline stages between fetch and issue
- Why not allow compiler to schedule instruction level parallelism explicitly?
- Format the instructions into a potential issue packet so that hardware need not check explicitly for dependences

VLIW: Very Large Instruction Word

- Each "instruction" has explicit coding for multiple operations
 - ▶ In IA-64, grouping called a "packet"
 - ▶ In Transmeta, grouping called a "molecule" (with "atoms" as ops)



Texas Instruments TMS320C64x

- All the operations the compiler puts in the long instruction word are independent, so can be issued and can execute in parallel
- E.g., 2 integer operations, 2 FP ops, 2 Memory refs, 1 branch
- **№** 16 to 24 bits per field => 7*16 or 112 bits to 7*24 or 168 bits wide
- Need compiling technique that schedules across several branches itecture Chapter 5.29

VLIW example: Transmeta Crusoe ▶ Transmeta's Shadow Registers Crusoe was a Shadow Registers Debug Reg 32 Floating 64 General Alias Hdw Point Purpose Registers TLB Registers processor T-Bit Buffer Instructions Load/Store FPU ALU0 ALUl Branch dynamically Gated Store translated Buffer from x86 in Local Program Local Data Memory 8KB Memory 8KB Data Flow & Instruction Data Cache Cache Control Data Cache Data Cache Control 64KB 64KB Morphing" 128 bit Molecule Secondary Instruction/Data Cache 256 KB LD**BRCC**

Bus Interface Unit

Note hardware support for speculation

Floating Integer Load/Store Point Unit ALU #0 Unit Instruction encoding

ADD

Branch

Unit

5—issue

VLIW

were

firmware

"Code

FADD

Recall: Unrolled Loop that Minimizes Stalls for Scalar

```
1 Loop: L.D F0,0(R1)
                                      L.D to ADD.D: 1 Cycle
2
       L.D F6, -8(R1)
                                      ADD.D to S.D: 2 Cycles
3
       L.D F10, -16(R1)
4
       L.D F14, -24(R1)
5
       ADD.D F4,F0,F2
6
       ADD.D F8, F6, F2
       ADD.D F12,F10,F2
8
       ADD.D F16,F14,F2
9
       S.D \quad O(R1), F4
10
       S.D -8 (R1), F8
11
       S.D
              -16(R1),F12
12
       DSUBUI R1, R1, #32
13
       BNEZ
              R1,LOOP
14
              8 (R1), F16
                            : 8-32 = -24
       S.D
```

14 clock cycles, or 3.5 per iteration

Loop Unrolling in VLIW

Memory reference 1	Memory reference 2	FP operation 1	FP op. 2	Int. op/ Clobranch	ock
L.D F0,0(R1)	L.D F6,-8(R1)				1
L.D F10,-16(R1)	L.D F14,-24(R1)				2
L.D F18,-32(R1)	L.D F22,-40(R1)	ADD.D F4,F0,F2	ADD.D F8	3,F6,F2	3
L.D F26,-48(R1)		ADD.D F12,F10,F2	ADD.D F1	6,F14,F2	4
		ADD.D F20,F18,F2	ADD.D F2	24,F22,F2	5
S.D 0(R1),F4	3.D -8(R1),F8	ADD.D F28,F26,F2			6
S.D -16(R1),F12	S.D -24(R1),F16				7
S.D -32(R1),F20	S.D -40(R1),F24			DSUBUI R1,R1,#48	8 8
S.D -0(R1),F28				BNEZ R1,LOOP	9

Unrolled 7 times to avoid delays

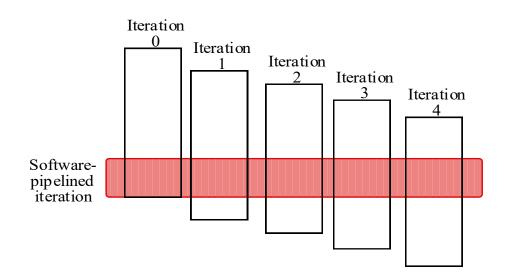
7 results in 9 clocks, or 1.3 clocks per iteration (1.8X)

Average: 2.5 ops per clock, 50% efficiency

Note: Need more registers in VLIW (15 vs. 6 in SS)

Software Pipelining

- Observation: if iterations from loops are independent, then can get more ILP by taking instructions from different iterations
- Software pipelining: reorganizes loops so that each iteration is made from instructions chosen from different iterations of the original loop (~ Tomasulo in SW)



Software Pipelining Example

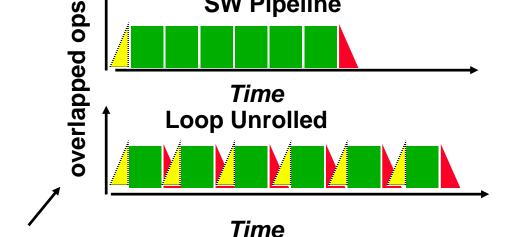
Before: Unrolled 3 times

```
L.D
         F0,0(R1)
  ADD.D F4,F0,F2
  S.D 0(R1), F4
  L.D F6, -8(R1)
5
 ADD.D F8,F6,F2
6 \text{ S.D} -8 (R1), F8
7 L.D F10,-16(R1)
 ADD.D F12,F10,F2
  S.D
         -16(R1),F12
10 DSUBUIR1,R1,#24
```

R1,LOOP

After: Software Pipelined (single-issue)

```
S.D
      0(R1),F4 ; Stores M[i]
ADD.D F4,F0,F2; Adds to M[i-1]
L.D F0,-16(R1); Loads M[i-2]
DSUBUIR1,R1,#8
BNEZ
      R1,LOOP
```



SW Pipeline

Symbolic Loop Unrolling

Maximize result-use distance

11 BNEZ

- Less code space than unrolling
- Fill & drain pipe only once per loop vs. once per each unrolled iteration in loop unrolling

Software Pipelining with Loop Unrolling in VLIW

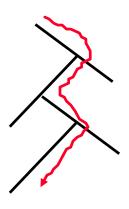
Memory reference 1	Memory reference 2	FP operation 1	FP op. 2	Int. op/ branch	Clock
L.D F0,-48(R1)	ST 0(R1),F4	ADD.D F4,F0,F2			1
L.D F6,-56(R1)	ST -8(R1),F8	ADD.D F8,F6,F2		DSUBUI R1,	R1,#242
L.D F10,-40(R1)	ST 8(R1),F12	ADD.D F12,F10,F2		BNEZ R1,L0	OOP 3

- Software pipelined across 9 iterations of original loop
 - **▶** In each iteration of above loop, we:
 - Store to m,m-8,m-16 (iterations I-3,I-2,I-1)
 - Compute for m-24,m-32,m-40 (iterations I,I+1,I+2)
 - Load from m-48,m-56,m-64 (iterations I+3,I+4,I+5)
- **№** 9 results in 9 cycles, or 1 clock per iteration
- Average: 3.67 instrs per clock, 91.75% efficiency

Note: Need fewer registers for software pipelining (only using 7 registers here, was using 15)

Trace Scheduling

- Parallelism across IF branches vs. LOOP branches?
- Two steps:
 - **→** Trace Selection
 - Find likely sequence of basic blocks (<u>trace</u>)
 of (statically predicted or profile predicted)
 long sequence of straight-line code
 - Trace Compaction
 - Squeeze trace into few VLIW instructions
 - Need bookkeeping code in case prediction is wrong
- This is a form of compiler-generated speculation
 - Compiler must generate "fixup" code to handle cases in which trace is not the taken branch
 - Needs extra registers: undoes bad guess by discarding
- Subtle compiler bugs mean wrong answer vs. poorer performance; no hardware interlocks
- Some scope for support from instruction set to help with fixup/rollback
 - **⇒** Eg Transmeta's shadow registers, store gate



Speculation: HW (Tomasulo) vs. SW (VLIW)

HW advantages:

- HW better at memory disambiguation since actual addresses are known at runtime
- HW better at branch prediction since lower overhead
- HW maintains precise exception model
- HW does not execute bookkeeping instructions
- Same software works across multiple implementations
- Smaller code size (not as many nops filling blank instructions)

SW advantages:

- Window of instructions that is examined for parallelism is unlimited
- Reduced transistor count, and power consumption (?)
- ▶ More elaborate types of speculation can be easier?
- Speculation can be based on large-scale program behavior, not just local information
 Advanced Computer Architecture Chapter 5.37

Problems with "First Generation" VLIW

Increase in code size

- generating enough operations in a straight-line code fragment requires ambitiously unrolling loops
- whenever VLIW instructions are not full, unused functional units translate to wasted bits in instruction encoding

Operated in lock-step; no hazard detection HW

- a stall in any functional unit pipeline caused entire processor to stall, since all functional units must be kept synchronized
- Compiler might know functional unit latencies, but caches harder to predict

Binary code compatibility

Pure VLIW => different numbers of functional units and unit latencies require different versions of the code

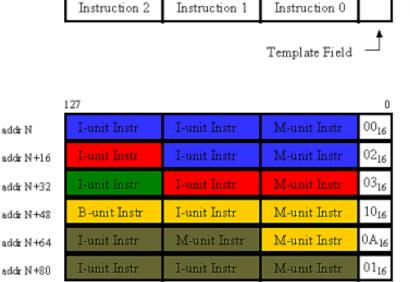
Intel/HP IA-64 "Explicitly Parallel Instruction Computer (EPIC)"

- MA-64: Intel's bid to create a new instruction set architecture
 - **▶**EPIC = "2nd generation VLIW"?
 - ◆ISA exposes parallelism (and many other issues) to the compiler
 - But is binary-compatible across processor implementations
- **Itanium** ™ first implementation (2001)
 - ◆ 6-wide, 10-stage pipeline
- **ltanium 2** (2002-2010)
 - ♦ 6-wide, 8-stage pipeline
 - → http://www.intel.com/products/server/processors/server/itanium2/
- **Itanium 9500 (Poulson) (2012)**
 - **♦ 12-wide, 11-stage pipeline**

(2017: Kittson, but no details available)

Instruction bundling in IA-64

- Instruction group: a sequence of consecutive instructions with no register data dependences
 - All instructions in a group could be executed in parallel, if sufficient hardware resources exist and if any dependences through memory are preserved
 - Instruction group can be arbitrarily long, but compiler must explicitly indicate boundary between one instruction group and another by placing a stop between 2 instructions that belong to different groups



- IA-64 instructions are encoded in bundles, which are 128 bits wide.
- Each bundle consists of a 5-bit template field and 3 instructions, each 41 bits in length
- One purpose of the template field is to mark where instructions in the bundle are dependent or independent, and whether they can be issued in parallel with the *next* bundle
- Eg for Poulson, groups of up to 4 bundles can be issued in parallel
- Smaller code size than old VLIW, larger than x86/RISC

IA-64

Instruction bundle

Instruction group 1
Instruction 2

group 2

Instruction

Instruction bundling in IA-64

Instructions can be explicitly sequential:

```
add r1 = r2, r3 ;;
sub r4 = r1, r2 ;;
shl r2=r4,r8
```

Or not:

```
add r1 = r2, r3
sub r4 = r11, r21
shl r12 = r14, r8 ;;
```

The ";;" syntax sets the "stop" bit that marks the end of a sequence of bundles that can be issued in parallel

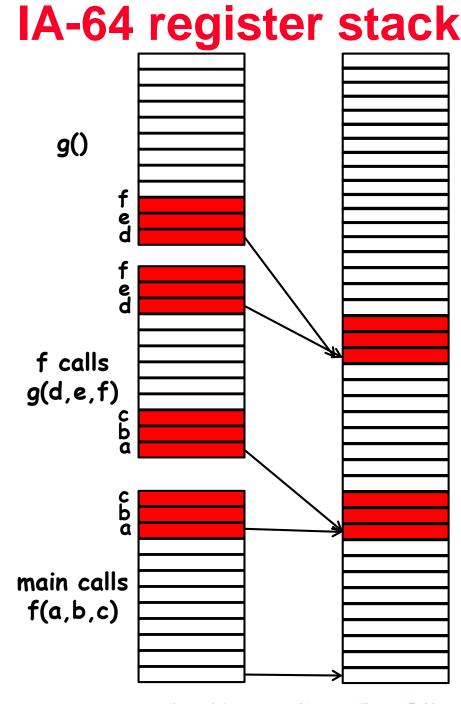
Hardware Support for Exposing More Parallelism at Compile-Time

To help trace scheduling and software pipelining, the Itanium instruction set includes several interesting mechanisms:

- Predicated execution
- Speculative, non-faulting Load instruction
- Software-assisted branch prediction
- Register stack
- Rotating register frame
- Software-assisted memory hierarchy
 - Job creation scheme for compiler engineers
 - ◆We will look at several of these in more detail

▶ General-purpose registers are configured to help accelerate procedure calls using a register stack

- mechanism similar to that developed in the Berkeley RISC-I processor and used in the SPARC architecture.
- Registers 0-31 are always accessible and addressed as 0-31
- Registers 32-128 are used as a register stack and each procedure is allocated a set of registers (from 0 to 96)
- → The new register stack frame is created for a called procedure by renaming the registers in hardware;
- → a special register called the current frame pointer (CFM) points to the set of registers to be used by a given procedure



Predication...

№ 64 1-bit predicate registers

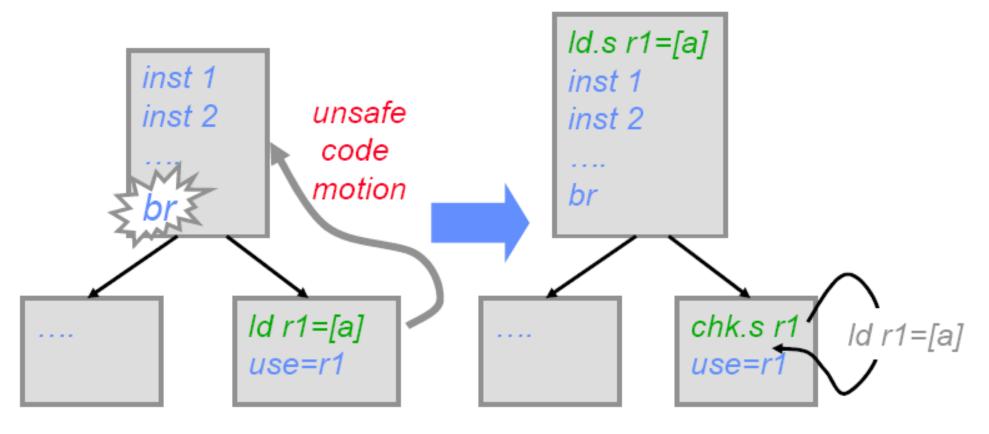
```
(p1) add r1 = r2, r3
(p2) sub r1 = r2, r3 ;;
shl r12 = r1, r8
```

- Predication means
 - **▶**Compiler can move instructions across conditional branches
 - **▶**To pack parallel issue groups
 - ◆May also eliminate some conditional branches completely
 - Avoiding branch prediction and misprediction
- There are also 8 64-bit Branch registers used to hold branch destination addresses for indirect branches – see later

IA64 load instruction variants

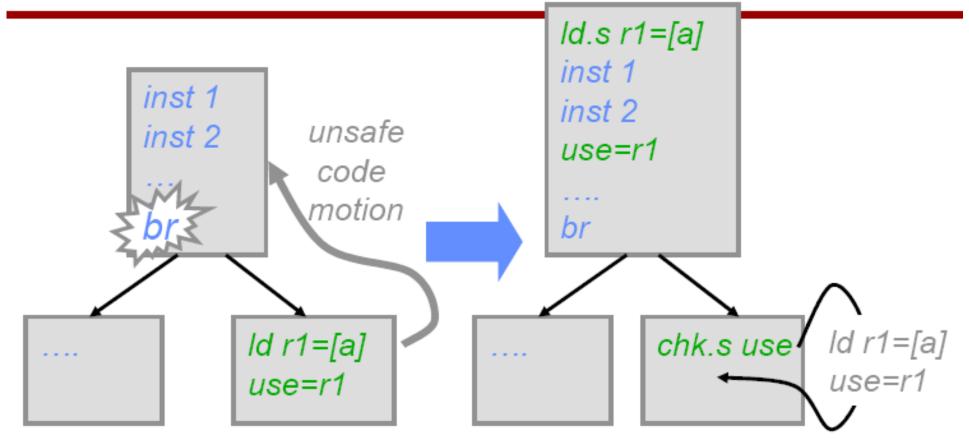
- IA64 has several different mechanisms to enable the compiler to schedule loads
- Id.s speculative, non-faulting
- Id.a speculative, "advanced" checks for aliasing stores
- Register values may be marked "NaT" not a thing
 - If speculation was invalid
- Advanced Load Address Table (ALAT) tracks stores to addresses of "advanced" loads

IA64: Speculative, Non-Faulting Load



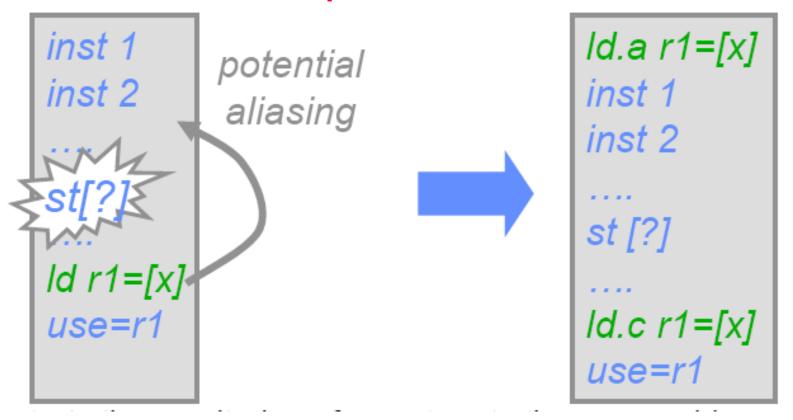
- Id.s fetches speculatively from memory
 - →i.e. any exception due to Id.s is suppressed
- If Id.s r did not cause an exception then chk.s r is an NOP, else a branch is taken to some compensation code

IA64: Speculative, Non-Faulting Load



- Speculatively-loaded data can be consumed prior to check
- "speculation" status is propagated with speculated data
- Any instruction that uses a speculative result also becomes speculative
 (i.e. suppressed exceptions)
- **chk.s** checks the entire dataflow sequence for exceptions

IA64: Speculative "Advanced" Load

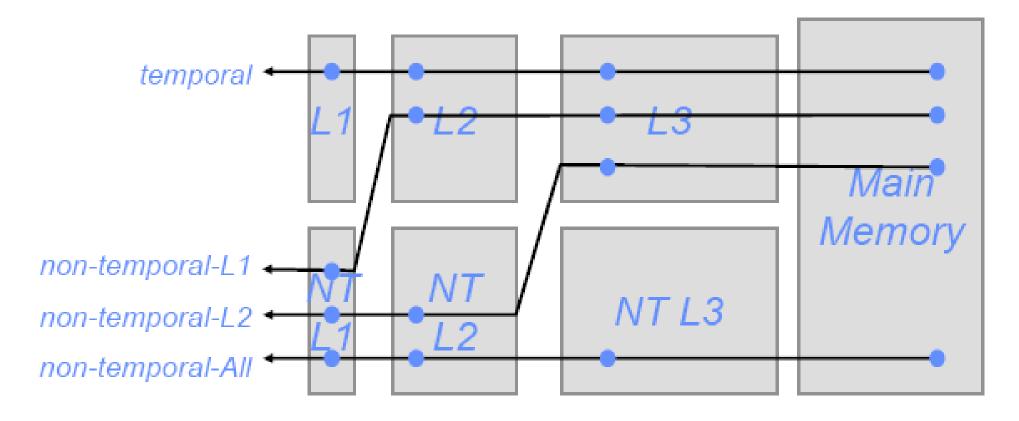


- Id.a starts the monitoring of any store to the same address as the advanced load
- If no aliasing has occurred since Id.a, Id.c is a NOP
- If aliasing has occurred, Id.c re-loads from memory

IA64: Branch prediction

- Static branch hints can be encoded with every branch
 - taken vs. not-taken
 - whether to allocate (or deallocate) an entry in the dynamic BP hardware (Itanium-1 used a 512-entry 2-level BHT and 64-entry BTB)
- TAR (Target Address Register)
 - → a small, fully-associative BTAC-like structure
 - contents are controlled entirely by a "prepare-to-branch" (aka an "importance" hint bit)
 - a hit in TAR overrides all other predictions
- RSB (Return Address Stack)
 - Procedure return addr is pushed (or popped) when a procedure is called (or when it returns)
 - Predicts nPC when executing register-indirect branches

IA64: Software-Assisted Memory Hierarchies



- ISA provides for separate storages for "temporal" vs "non-temporal" data, each with its own multiple level of hierarchies
- Load and Store instructions can give hints about where cached copies should be held after a cache miss

IA-64 Registers

- ▶ Both the integer and floating point registers support register rotation for registers 32-128.
- Register rotation is designed to ease the task of register allocation in software pipelined loops
- When combined with predication, possible to avoid the need for unrolling and for separate prologue and epilogue code for a software pipelined loop
 - makes the SW-pipelining usable for loops with smaller numbers of iterations, where the overheads would traditionally negate many of the advantages

How Register Rotation Helps Software Pipelining sissue packet Il see) 1 each time

The concept of a software pipelining branch:

```
L1:
                                 // post-increment by 4
                r35 = [r4], 4
        ld4
                [r5] = r37, 4
                                 // post-increment by 4
        st4
        br.ctop L1;;
```

One issue packet The br.ctop instruction in the example rotates the general registers (actually br.ctop does more as we shall see)

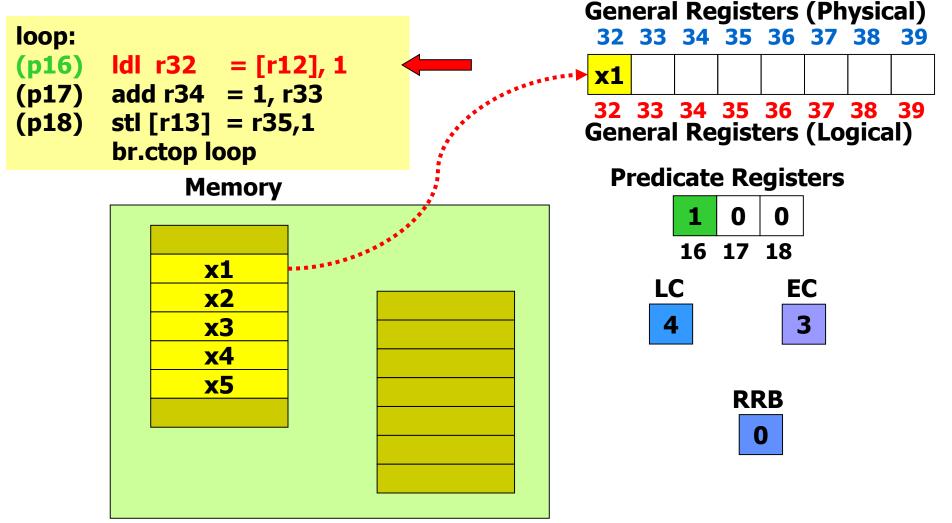
Therefore the value stored into r35 is read in r37 two iterations (and two rotations) later.

The register rotation eliminated a dependence between the load and the store instructions, and allowed the loop to execute in one cycle.

- Register rotation is useful for procedure calls
- It's also useful for software-pipelined loops
- The logical-to-physical register mapping is shifted by 1 each time the branch ("br.ctop") is executed

```
// Clear all rotating predicate registers
       mov pr.rot
                       = 0
       cmp.eq p16,p0 = r0,r0 // Set p16=1
                       = 4 // Set loop counter to n-1
       mov ar.lc
                               // Set epilog counter to 3
                       = 3
       mov ar.ec
                                                       One issue packet
loop:
(p16)
       |d| r32 = [r12], 1
                              // Stage 1: load x
                              // Stage 2: y=x+1
       add r34 = 1, r33
(p17)
       stl [r13] = r35,1
                               // Stage 3: store y
(p18)
       br.ctop loop
                                 Branch back
```

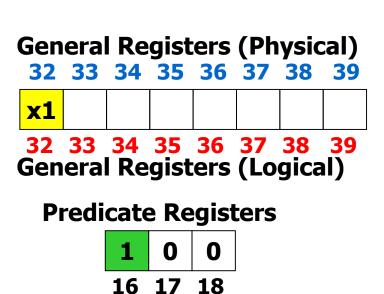
- "Stage" predicate mechanism allows successive stages of the software pipeline to be filled on start-up and drained when the loop terminates
- The software pipeline branch "br.ctop" rotates the predicate registers, and injects a 1 into p16
- Thus enabling one stage at a time, for execution of prologue
- When loop trip count is reached, "br.ctop" injects 0 into p16, disabling one stage at a time, then finally falls-through

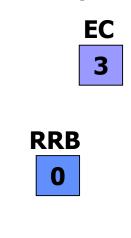


loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory

x1 x2 x3 x4 x5

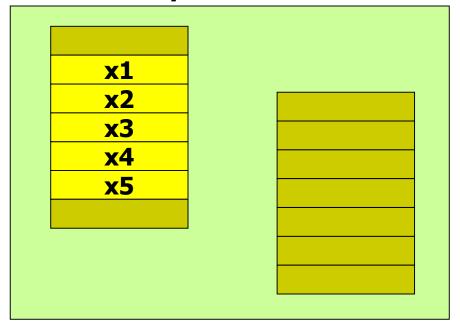


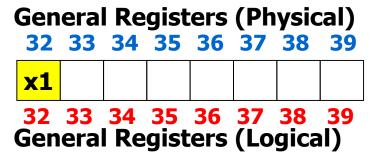


LC

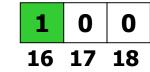
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory





Predicate Registers

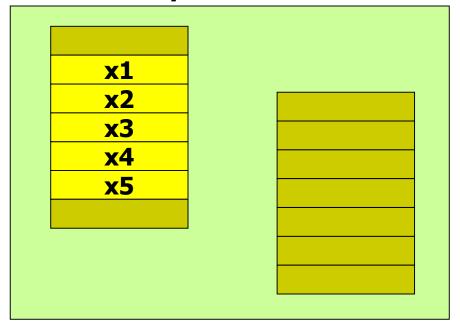


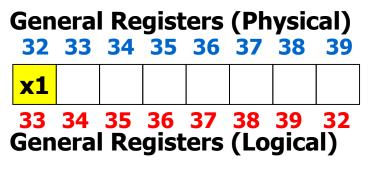
LC 4 **EC**

RRB 0

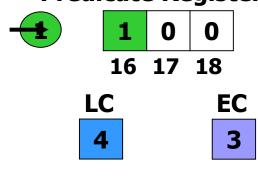
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory





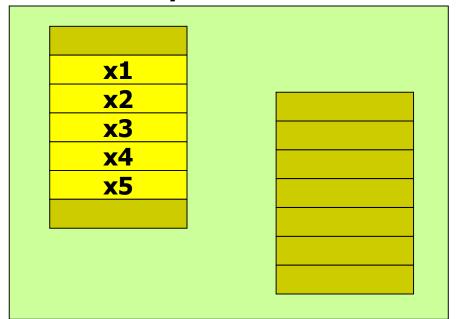
Predicate Registers

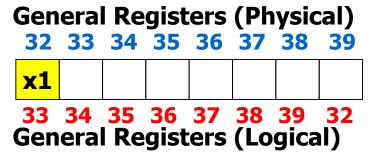


RRB -1

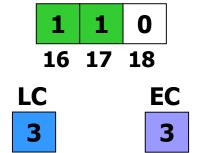
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory

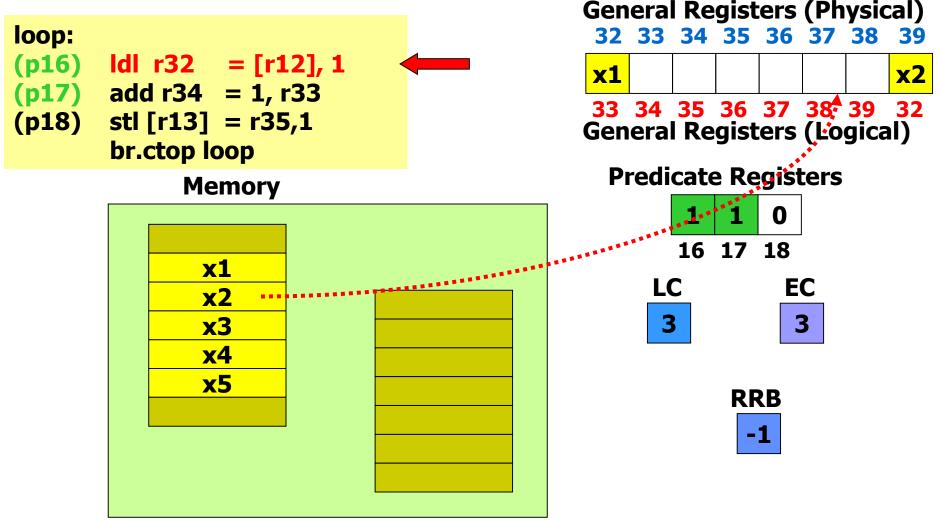




Predicate Registers

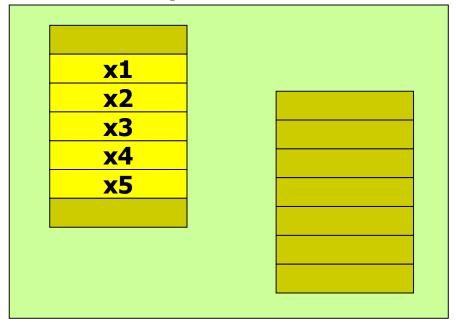


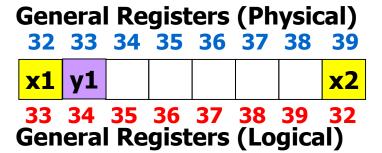
RRB -1



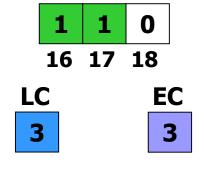
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory





Predicate Registers

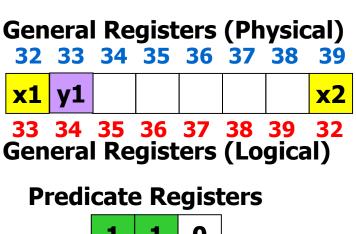


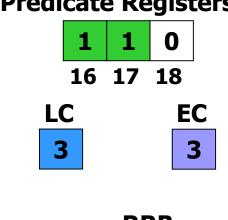
RRB

loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory

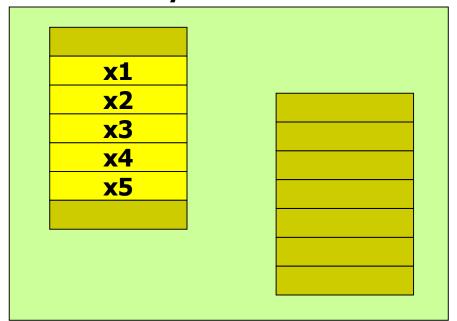
x1 x2 x3 x4 x5

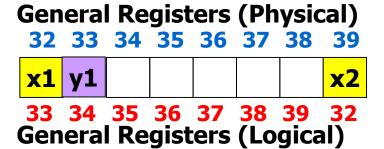




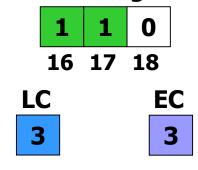
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory





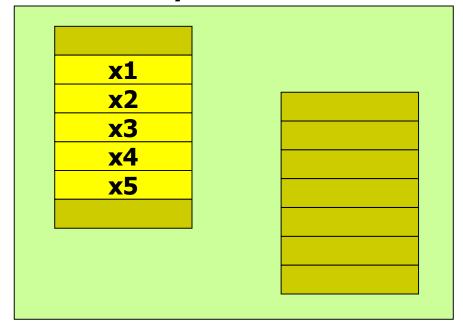
Predicate Registers

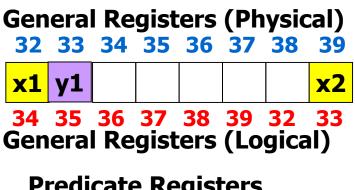




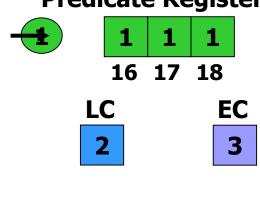
loop: |d| r32 = [r12], 1(p16) (p17) add r34 = 1, r33(p18) stl [r13] = r35,1br.ctop loop

Memory

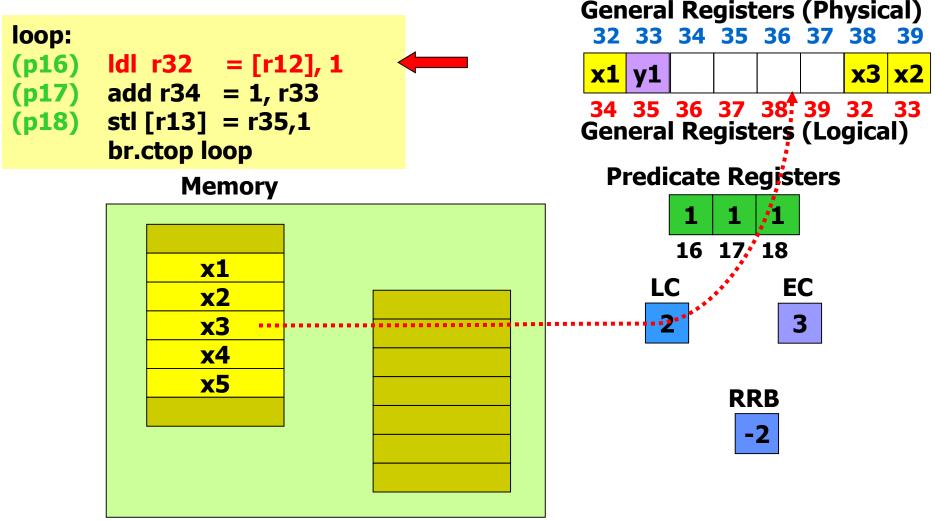




Predicate Registers

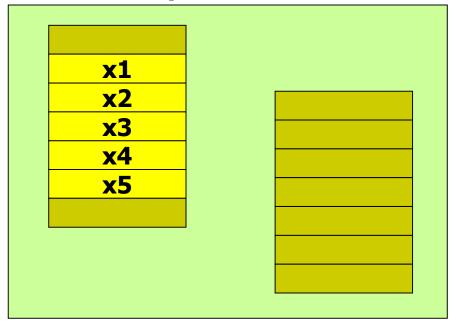


RRB -2



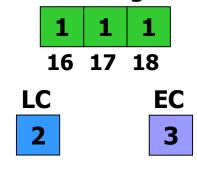
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory

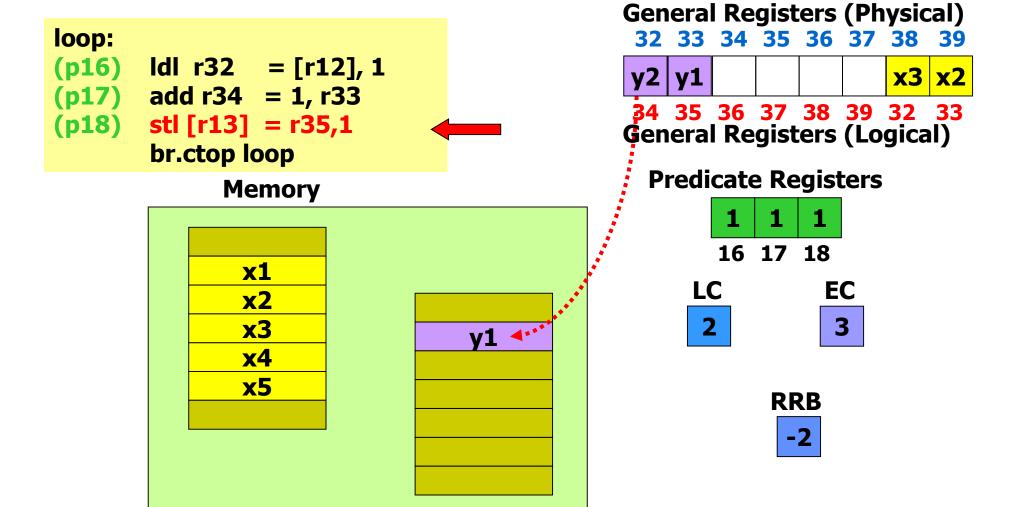


General Registers (Physical) 32 33 34 35 36 37 38 39 y2 y1 x3 x2 34 35 36 37 38 39 32 33 General Registers (Logical)

Predicate Registers

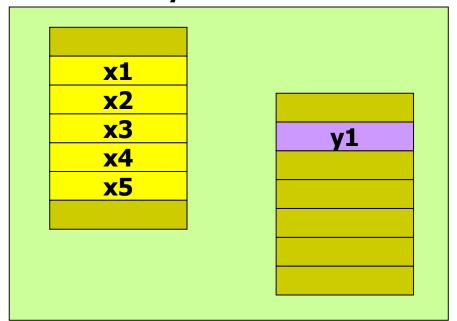


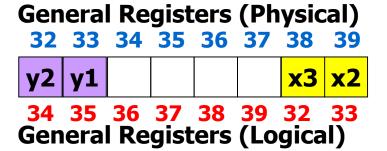
RRB -2



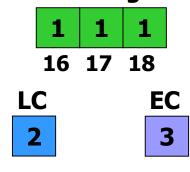
loop: (p16) | Idl r32 = [r12], 1 (p17) | add r34 = 1, r33 (p18) | stl [r13] = r35,1 | br.ctop loop

Memory





Predicate Registers



RRB -2

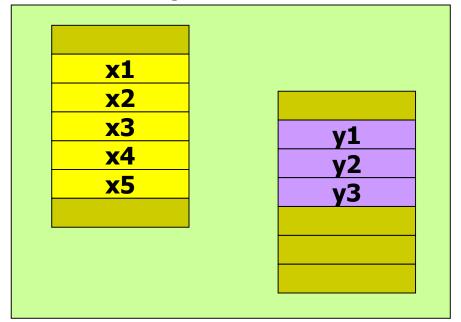
Hidden slides...

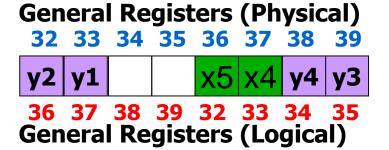
Some hidden slides are not in handout

We continue with start of pipeline drain phase

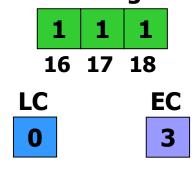
loop: (p16) ldl r32 = [r12], 1 (p17) add r34 = 1, r33 (p18) stl [r13] = r35,1 br.ctop loop

Memory





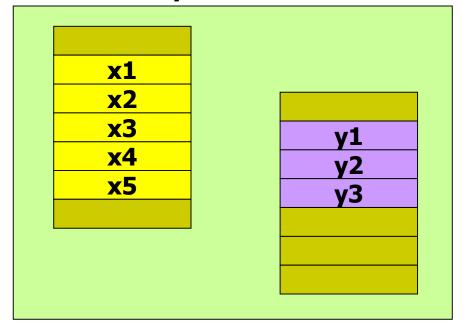
Predicate Registers

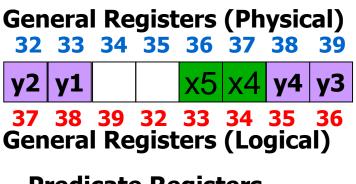


RRB -4

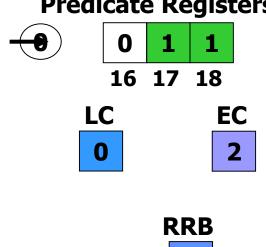
```
loop:
      |d| r32 = [r12], 1
(p16)
(p17) add r34 = 1, r33
(p18) stl [r13] = r35,1
       br.ctop loop
```

Memory

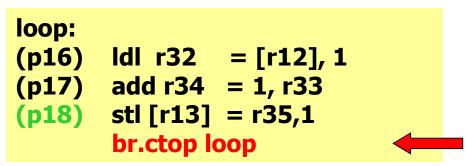




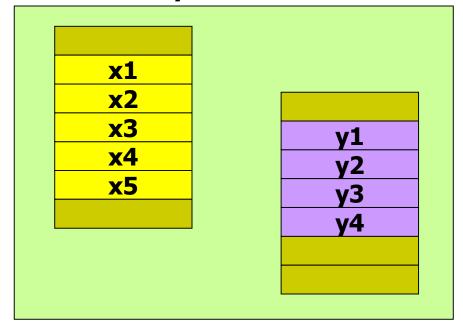
Predicate Registers



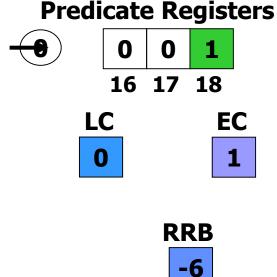
-5



Memory

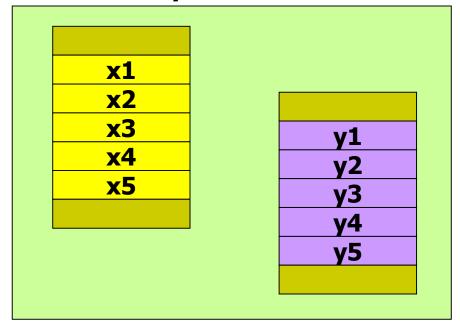


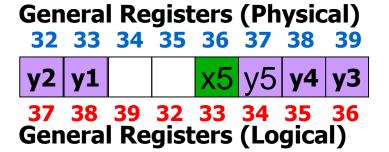
General Registers (Physical) 32 33 34 35 36 37 38 36 37 38 39 32 33 34 35 **General Registers (Logical)**



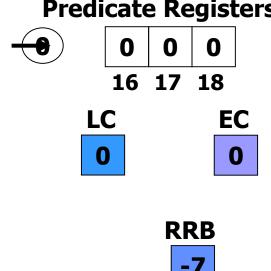
```
loop:
(p16)
      |d| r32 = [r12], 1
(p17) add r34 = 1, r33
(p18) stl [r13] = r35,1
       br.ctop loop
```

Memory



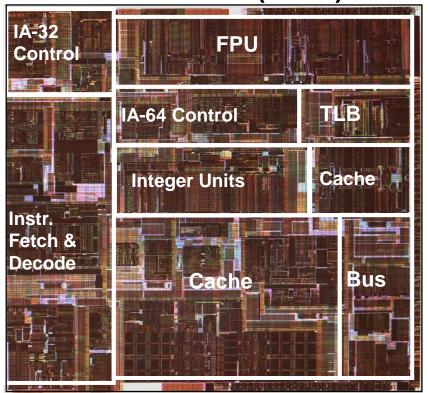


Predicate Registers



Caches

- → 32KB L1 (2 cycle)
- **▶** 96KB L2 (7 cycle)
- **▶** 2 or 4 MB L3 (off chip)
- 133 MHz 64-bit bus
- SpecFP: 711
- SpecInt: 404
- 36-bit addresses (64GB)



Core Processor Die

Itanium™ Processor Silicon

(Copyright: Intel at Hotchips '00)



4 x 1MB L3 cache

Itanium™ EPIC Design Maximizes SW-HW Synergy

(Copyright: Intel at Hotchips '00)

Architecture Features programmed by compiler:

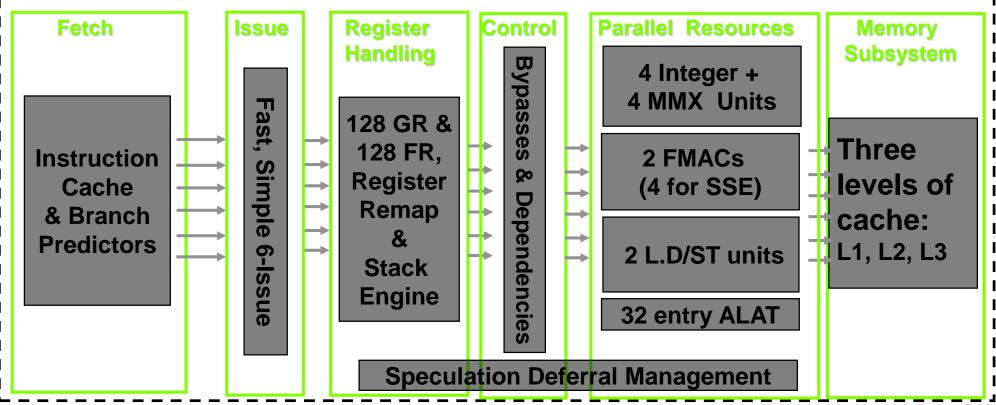
Branch Hints Explicit Parallelism Register Stack & Rotation

Predication

Data & Control Speculation

Memory Hints

Micro-architecture Features in hardware:



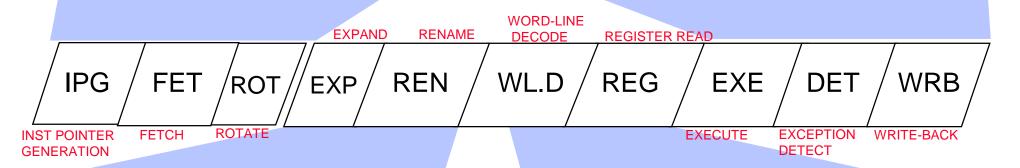
10 Stage In-Order Core Pipeline (Copyright: Intel at Hotchips '00)

Front End

- Pre-fetch/Fetch of up to 6 instructions/cycle
- Hierarchy of branch predictors
- Decoupling buffer

Execution

- 4 single cycle ALUs, 2 ld/str
- Advanced load control
- Predicate delivery & branch
- Nat/Exception//Retirement

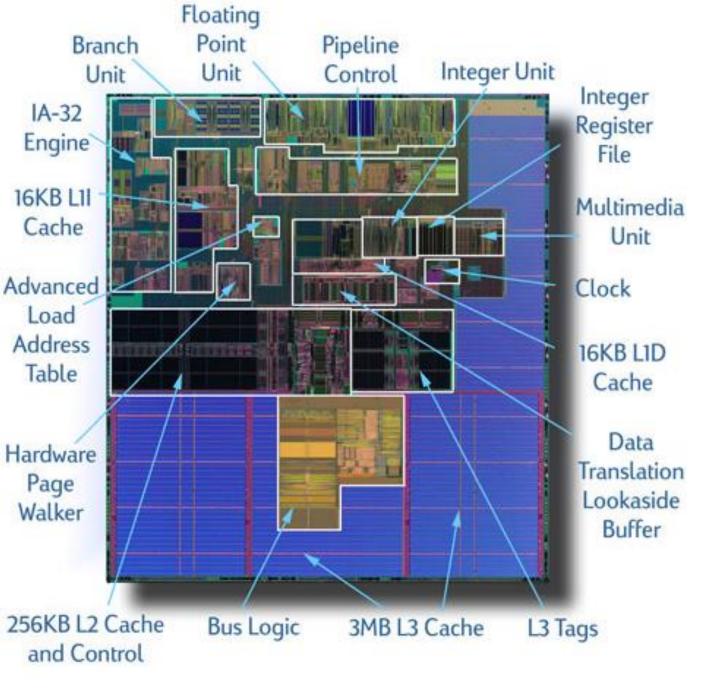


Instruction Delivery

- Dispersal of up to 6 instructions on 9 ports
- Reg. remapping
- Reg. stack engine

Operand Delivery

- Reg read + Bypasses
- Register scoreboard
- Predicated dependencies



Itanium 2

- Caches
 - 32KB L1 (1 cycle)
 - **→** 256KB L2 (5 cycle)
 - → 3 MB L3 (on chip)
- 200 MHz 128-bit Bus
- SpecFP: 1356
- SpecInt: 810
- 44-bit addresses (18TB)
- 221M transistors
- **19.5 x 21.6 mm**
- http://cpus.hp.com/technical references/

Comments on Itanium

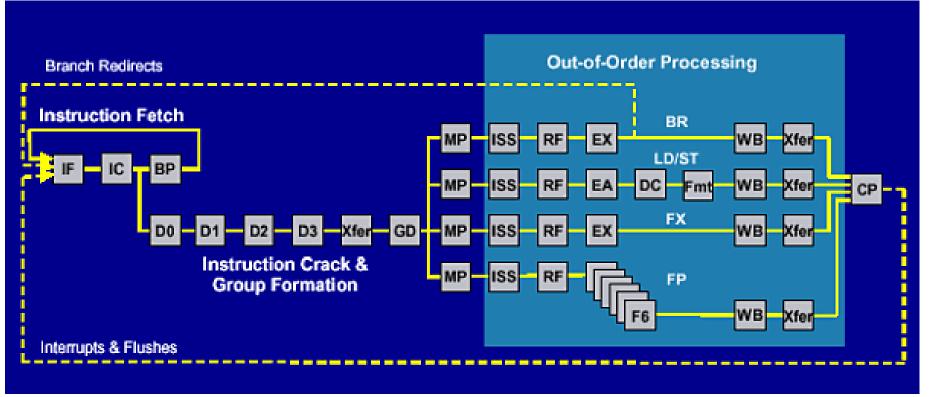
- Remarkably, the Itanium has many of the features more commonly associated with the dynamicallyscheduled pipelines
 - strong emphasis on branch prediction
 - register renaming,
 - scoreboarding,
 - a deep pipeline with many stages before execution (to handle instruction alignment, renaming, etc.),
 - several stages following execution to handle exception detection
- Surprising that an approach whose goal is to rely on compiler technology and simpler HW seems to be at least as complex as dynamically scheduled processors!

Comments on Itanium

Compare Itanium II



With IBM Power 4:



EPIC/IA-64/Itanium principles

Start loads early

- advance loads move above stores when alias analyis is incomplete
- speculative loads move above branches

Predication to eliminate many conditional branches

- 64 predicate registers
- almost every instruction is predicated

register rich

- 128 integer registers (64 bits each)
- 128 floating-point registers

Independence architecture

- VLIW flavor, but fully interlocked (i.e., no delay slots)
- three 41-bit instruction syllables per 128-bit "bundle"
- each bundle contains 5 "template bits" which specify independence of following syllables (within bundle and between bundles)

unbundled branch architecture

- eight branch registers
- multiway branches

Rotating register files

- lower 48 of the predicate registers rotate
- **▶** lower 96 of the integer registers rotate

1983: Josh Fisher describes ELI-512 VLIW design and trace scheduling 1983-1988: Rau at Cydrome works on VLIW design called Cydra-5, but company folds 1988 1984-1990: Fisher at Multiflow works on VLIW design called Trace, but company folds 1990 1988: Dick Lampman at HP hires Bob Rau and Mike Schlansker from Cydrome and also gets IP rights from Cydrome 1989: Rau & Schlansker begin FAST (Fine-grained Architecture & Software Technologies) research

1981: Bob Rau leads Polycyclic Architecture project at TRW/ESL

- 1990-1993: Bill Worley leads PA-WW (Precision Architecture Wide-Word) effort at HP Labs to be successor to PA-RISC architecture; also called SP-PA (Super-Parallel Processor Architecture) & SWS (SuperWorkStation)
- **HP hires Josh Fisher, input to PA-WW** Input to PA-WW from Hitachi team, led by Yasuyuki Okada
- 1991: Hans Mulder joins Intel to start work on a 64-bit architecture
- 1992: Worley recommends HP seek a semiconductor manufacturing partner

project at HP; later develop HP PlayDoh architecture

- 1993: HP starts effort to develop PA-WW as a product
- Dec 1993: HP investigates partnership with Intel
- June 1994: announcement of cooperation between HP & Intel; PA-WW starting point for joint design; John Crawford of Intel leads joint team
- 1997: the term EPIC is coined
- Oct 1997: Microprocessor Forum presentations by Intel and HP
- July 1998: Carole Dulong of Intel, "The IA-64 Architecture at Work," IEEE Computer
- Feb 1999: release of ISA details of IA-64
- 2001: Intel marketing prefers IPF (Itanium Processor Family) to IA-64

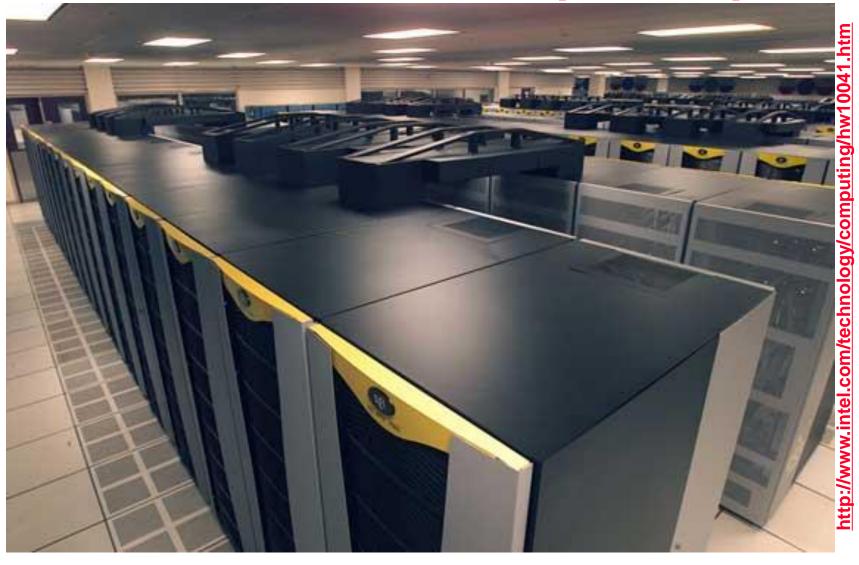
Dec 2004: HP transfers last of Itanium development to Intel

May 2001 - Itanium (Merced)

Itanium Timeline

- July 2002 Itanium 2 (McKinley)
 - Aug 2004: "Itanium sales fall \$13.4bn shy of \$14bn forecast" (The Register)
 - Advanced Computer Architecture Chapter 5.106

Itanium supercomputer



 NASA's 10,240-processor Columbia supercomputer (2004-2013) built from 20 Altix systems, each powered by 512 Intel Itanium 2 processors. Peak performance 42.7 TeraFlops; 20TB RAM. Single shared address space. Ran Linux. (Image courtesy of Silicon Graphics, Inc.)

Гор	20	SPEC	systems	

int peak int base Full results MHz

HTML

HTML

HTML

HTML

HTML

HTML

HTML

HTML

HTML

1854 HTMI

3108

3089

2844

2826

1942

1923

1749

1820

1870

Top 20 SPECfp2000

3642

3098

3056

3050

3044

3017

2850

2796

2497

2462

Processor

2300 POWER5+

3000 Xeon 51xx

2660 Xeon 30xx

2667 Core 2 Duo

2800 Onteron AM2

1600 Itanium 2

1900 POWER5

3000 Opteron

1600 DC Itanium 2

2933 Core 2 Duo EE

fp peak fp base Full results

3369 HTML

3098 HTML

2811 HTML

3048 HTML

2763 HTML

3017 HTML

2847 HTML

2585 HTML

2260 HTML

2230 HTMI

Top 20 SPECint2000

3119

3102

2848

2835

2119

2061

1960

1900

1872

1856

Processor

2933 Core 2 Duo EE

3000 Xeon 51xx

2666 Core 2 Duo

2660 Xeon 30xx

2800 Athlon 64 FX

2800 Opteron AM2

2300 POWER5+

3733 Pentium 4 E

3800 Pentium / Yeon

3000 Opteron

MHz

10	3800 Pentium 4 Xeon	1850	1854	HIM	L	2800 Opteron AM2	2402	2230	HIML	ı
11	2260 Pentium M	1839	1812	HTM	L	3733 Pentium 4 E	2283	2280	HTML	
12	3600 Pentium D	1814	1810	HTM	<u>L</u>	2800 Athlon 64 FX	2261	2086	HTML	
13	2167 Core Duo	1804	1796	HTM	<u>L</u>	2700 PowerPC 970N	IP 2259	2060	HTML	
14	3600 Pentium 4	1774	1772	HTM	<u>L</u>	2160 SPARC64 V	2236	2094	HTML	
15	3466 Pentium 4 EE	1772	1701	HTM	L	3730 Pentium 4 Xeor	n 2150	2063	HTML	
16	2700 PowerPC 970MP	1706	1623	HTM	Ĺ	3600 Pentium D	2077	2073	HTML	1
17	2600 Athlon 64	1706	1612	HTM	L	3600 Pentium 4	2015	2009	HTML	1
18	2000 Pentium 4 Xeon LV	1668	1663	HTM	L	2600 Athlon 64	1829	1700	HTML	١,
19	2160 SPARC64 V	1620	1501	HTM	L	1700 POWER4+	1776	1642	HTML	
20	1600 Itanium 2	1590	1590	HTM	<u>L</u>	3466 Pentium 4 EE	1724	1719	HTML	
With Auto-parallelisation										
Top 20 SPECint2000 Top 20 SPECfp2000								 h		
# :	MHz Processor int peak	int base	Full r	esults	\mathbf{MHz}	Processor	fp peak	fp base	Full results	"
1					2100	POWER5+	4051	3210	HTML	
2					3000	Opteron	3538	2851	HTML	
3					2600	Opteron AM2	3338	2711	HTML	
4					1200	UltraSPARC III Cu	1344	1074	HTML	ute

Aces Hardware analysis of SPEC benchmark data http://www.aces hardware.com/S PECmine/top.jsp

http://www.spec.org/cpu20 06/results/cpu2006.html

uter Architecture Chapter 5.108

Summary#1: Hardware versus Software Speculation Mechanisms

- To speculate extensively, must be able to disambiguate memory references
 - Much easier in HW than in SW for code with pointers
- HW-based speculation works better when control flow is unpredictable, and when HW-based branch prediction is superior to SW-based branch prediction done at compile time
 - Mispredictions mean wasted speculation
- HW-based speculation maintains precise exception model even for speculated instructions
- HW-based speculation does not require compensation or bookkeeping code

Summary#2: Hardware versus Software Speculation Mechanisms cont'd

- Compiler-based approaches may benefit from the ability to see further in the code sequence, resulting in better code scheduling
- HW-based speculation with dynamic scheduling does not require different code sequences to achieve good performance for different implementations of an architecture
 - may be the most important in the long run?
- Example: ARM's "big.LITTLE" architecture
 - Multicore processor with a mixture of large out-of-order cores (A15) and small in-order cores (A7) (eg Exynos 5 Octa in Samsung Galaxy S4)
 - Compiler is configured to schedule for in-order, assuming the out-of-order processor is less sensitive to instruction scheduling

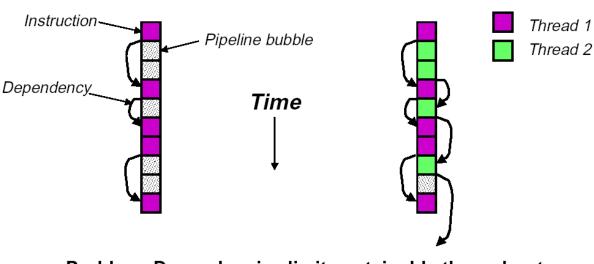
Summary #3: Software Scheduling

- Instruction Level Parallelism (ILP) found either by compiler or hardware.
- Loop level parallelism is easiest to see
 - SW dependencies/compiler sophistication determine if compiler can unroll loops
 - **▶** Memory dependencies hardest to determine => Memory disambiguation
 - Very sophisticated transformations available
- Trace scheduling to parallelize if statements
- Superscalar and VLIW: CPI < 1 (IPC > 1)
 - **▶** Dynamic issue vs. Static issue
 - → More instructions issue at same time => larger hazard penalty
 - Limitation is often number of instructions that you can successfully fetch and decode per cycle

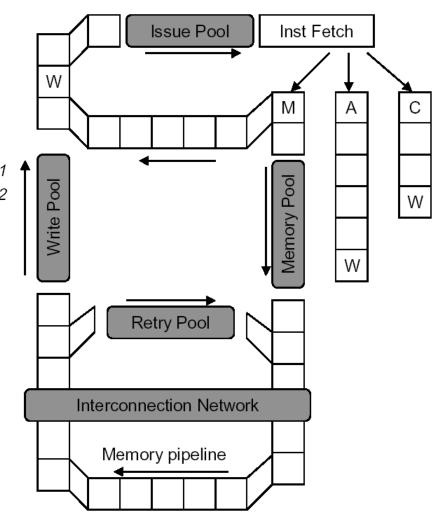
Beyond ILP: Multithreading, Simultaneous Multithreading (SMT)

Cray/Tera MTA

http://www.cray.com/products/systems/ mta/, http://www.utc.edu/~jdumas/cs460/pape rsfa01/craymta/



- Problem: Dependencies limit sustainable throughput of single instruction stream
- Solution: Interleave execution of two or more instruction streams on same hardware to increase utilization

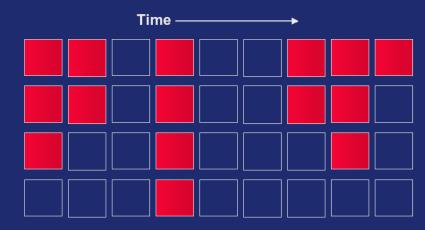


Time ————

Reduced function unit utilization due to dependencies

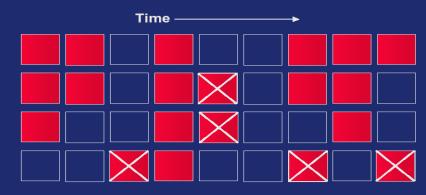
Superscalar Issue

Instruction Issue



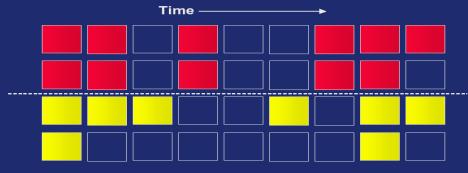
Superscalar leads to more performance, but lower utilization

Predicated Issue



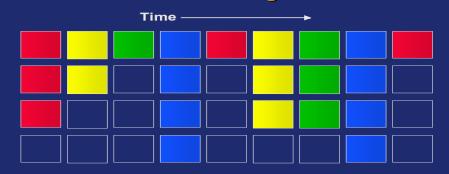
Adds to function unit utilization, but results are thrown away





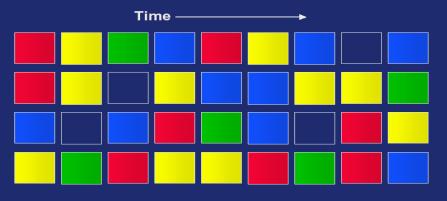
Limited utilization when only running one thread

Fine Grained Multithreading



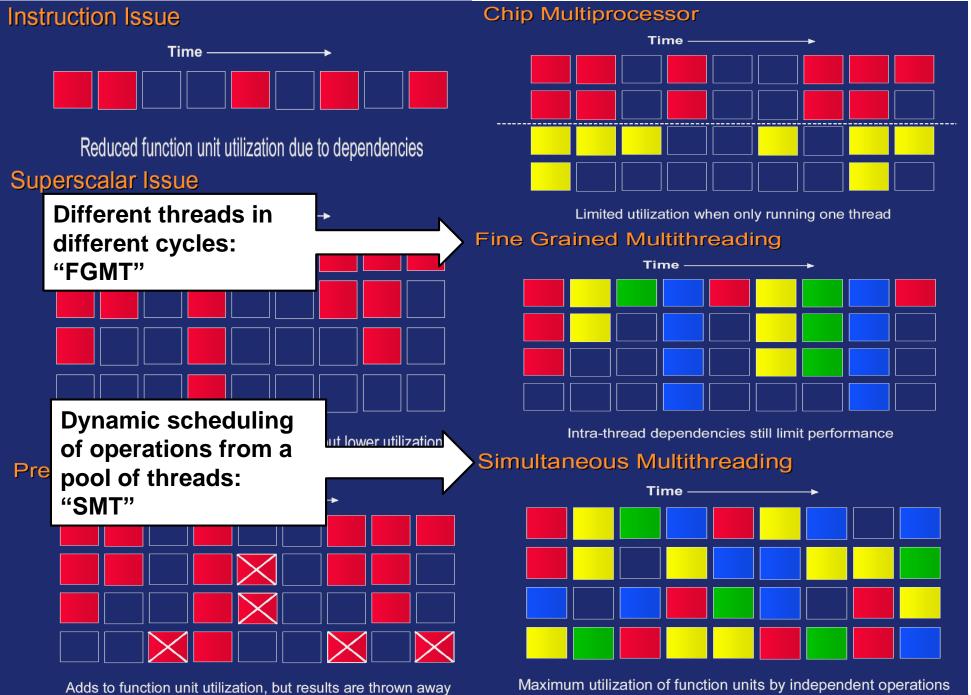
Intra-thread dependencies still limit performance

Simultaneous Multithreading



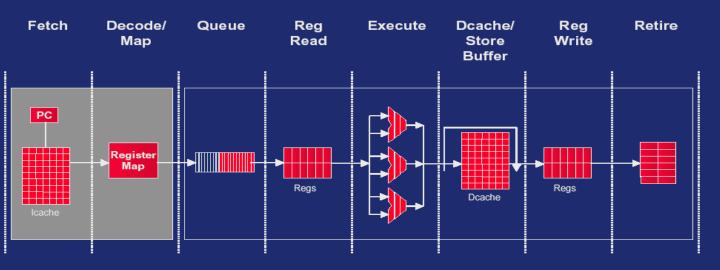
Maximum utilization of function units by independent operations







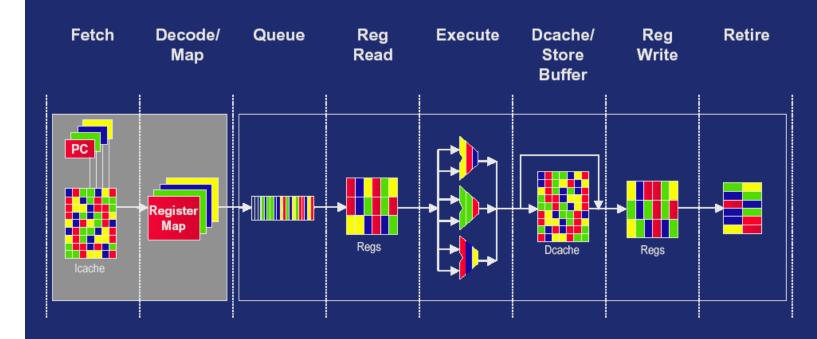
Basic Out-of-order Pipeline



SMT

SMT Pipeline

- Alpha 21464
- One CPU with 4 Thread Processing Units (TPUs)
- "6% area overhead over single-thread 4-issue CPU"

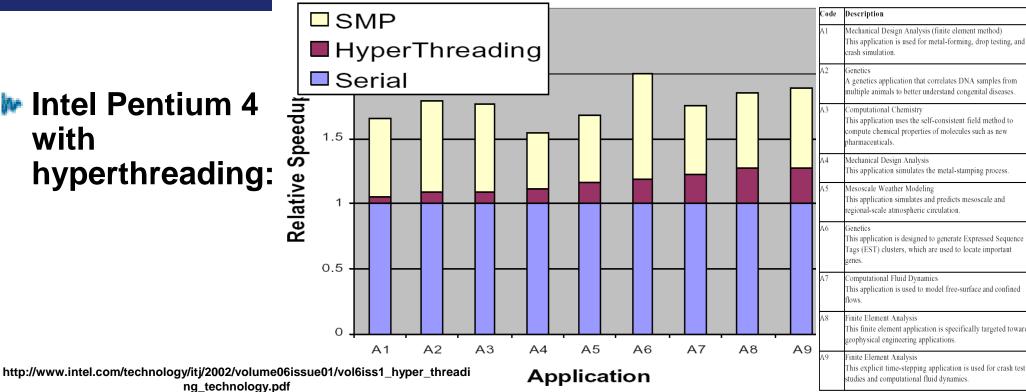




performance

Mar. Alpha 21464

Intel Pentium 4 with hyperthreading:

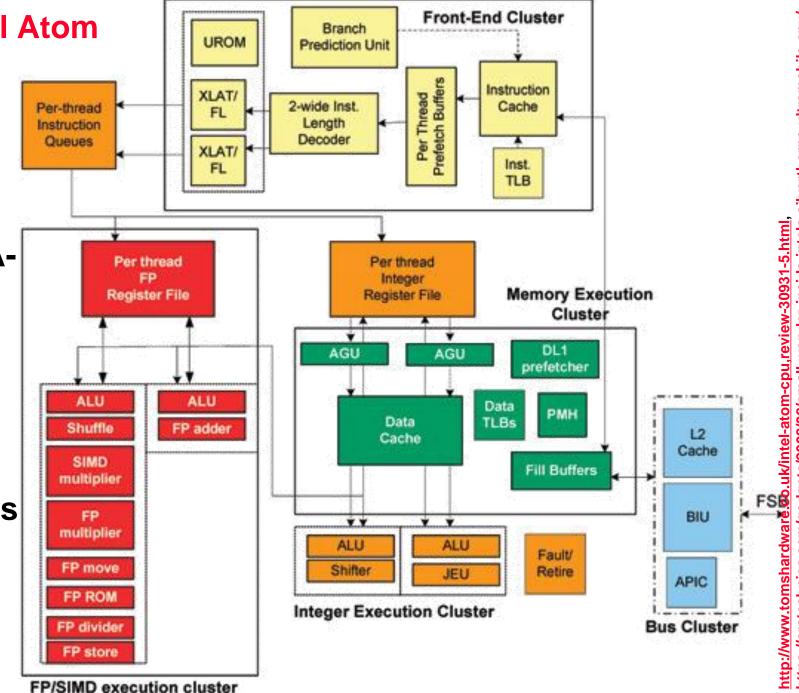


SMT in the Intel Atom

Intel's bid to steal back some of the low-power market for IA-32 and Windows

- **In-order**
- 2-way SMT
- 2 instructions per cycle

(from same or different threads)



Each thread runs slow?

SMT issues

→ The point of Simultaneous Multithreading is that resources are dynamically assigned, so if only one thread can run it can run faster

SMT threads contend for resources

- Possibly symbiotically?
 - One thread is memory-intensive, one arithmetic-intensive?
- Possibly destructively
 - thrashing the cache? Other shared resources.... (TLB?)
- Which resources should be partitioned per-thread, and which should be shared on-demand?
- SMT threads need to be scheduled fairly
 - Can one thread monopolise the whole CPU?
 - Denial of service risk
 - Slow thread that suffers lots of cache misses fills RUU and blocks issue

Side channels:

one thread may be able observe another's traffic and deduce what it's doing

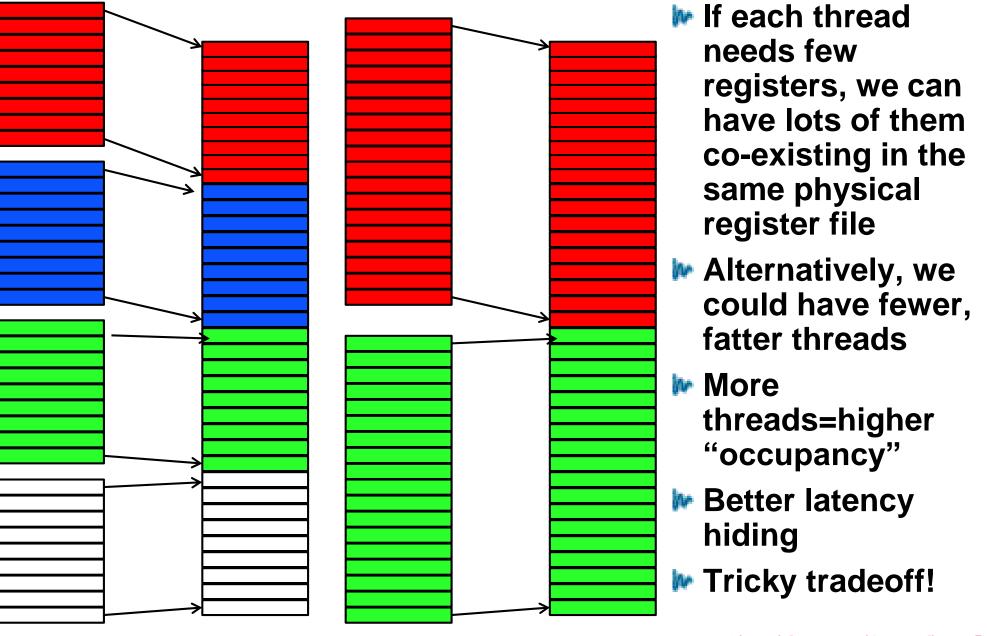
SMT – latency-hiding

- SMT threads exploit memory-system parallelism
 - Easy way to get lots of memory accesses in-flight
 - "Latency hiding" overlapping data access with compute
- What limits the number of threads we can have?
- SMT threads need a *lot* of registers
 - **▶** A lot of logical registers but they share physical registers?

In a machine without register renaming

- What about statically partitioning the register file based on the number of registers each thread actually needs?
- **➡** This is what many GPUs do
- Leads to tradeoff: lots of lightweight threads to maximise latency hiding? Or fewer heavyweight threads that benefit from lots of registers?
- **▶** Nvidia and AMD call this "occupancy"

Mapping threads into the register file



CUDA GPU Occupancy Calculator

Just follow steps 1, 2, and 3 below! (or click here for help)

1.) Select Compute Capability (click):

1.b) Select Shared Memory Size Config (bytes)

2.) Enter your resource usage:
Threads Per Block
Registers Per Thread
Shared Memory Per Block (bytes)

(Help)

(Help)

(Don't edit anything below this line)

3.) GPU Occupancy Data is displayed here and in the grap	hs:
Active Threads per Multiprocessor	1024
Active Warps per Multiprocessor	32
Active Thread Blocks per Multiprocessor	16
Occupancy of each Multiprocessor	50%

Physical Limits for GPU Compute Capability: Threads per Warp Warps per Multiprocessor Threads per Multiprocessor 2048 Thread Blocks per Multiprocessor 65536 Total # of 32-bit registers per Multiprocessor 256 warp 255 Shared Memory per Multiprocessor (bytes) 49152 Shared Memory Allocation unit size 256 Warp allocation granularity Maximum Thread Block Size 1024

				= Allocatable
Allocated Re	sources	Per Block	Limit Per SM	Blocks Per SM
Warps	(Threads Per Block / Threads Per Warp)	2	64	16
Registers	(Warp limit per SM due to per-warp reg count)	2	64	32
Shared Memo	ory (Bytes)	0	49152	16

(Help)

Occupancy = 32 / 64 = 50%

Note: SM is an abbreviation for (Streaming) Multiprocesso

Maximum Thread Blocks Per Multiprocessor	Blocks/SM	* Warps/Block	= Warps/SM
Limited by Max Warps or Max Blocks per Multiprocessor	16	2	32
Limited by Registers per Multiprocessor	32		
Limited by Shared Memory per Multiprocessor	16	2	32
Note: Occupancy limiter is shown in orange		Physical Ma	x Warps/SM = 64

CUDA Occupancy Calculator

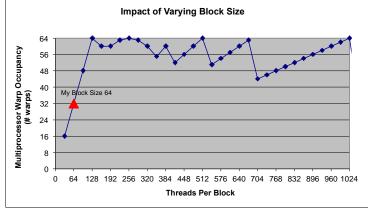
Version: 5.1

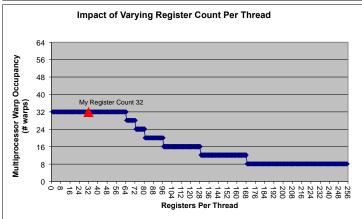
Copyright and License

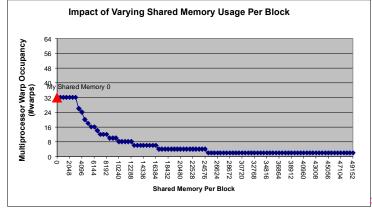
Click Here for detailed instructions on how to use this occupancy calculator.

For more information on NVIDIA CUDA, visit http://developer.nvidia.com/cuda

Your chosen resource usage is indicated by the red triangle on the graphs. The other data points represent the range of possible block sizes, register counts, and shared memory allocation.





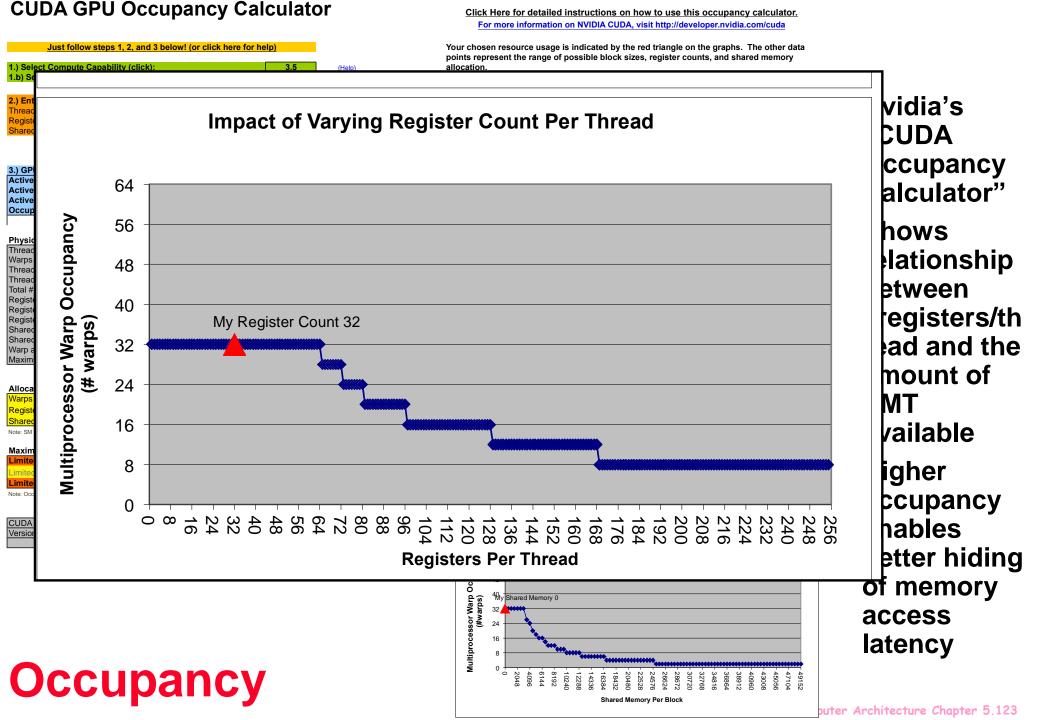


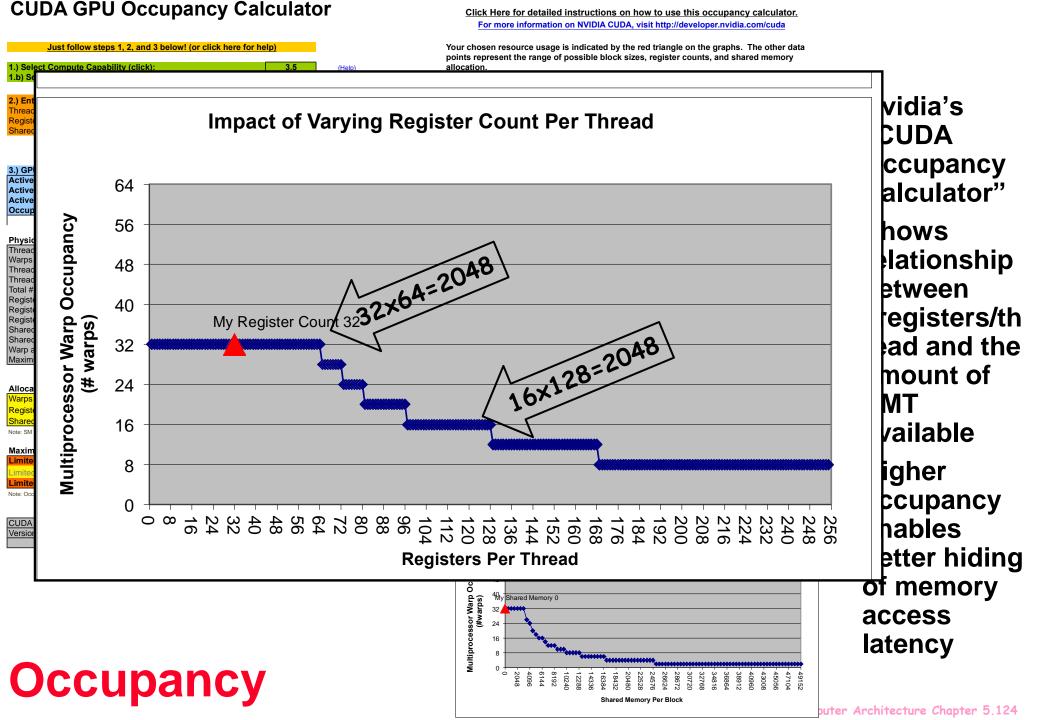
- Nvidia's "CUDA Occupancy Calculator"
- Shows relationship between #registers/th read and the amount of SMT available
- Higher
 occupancy
 enables
 better hiding
 of memory
 access
 latency

Occupancy

CUDA (BU A A decidate			ΤΙ.	
I	1.) Select Compute Capability (click):	3.5	(Help)	a ⁻	
	1.b) Select Shared Memory Size Config (bytes)	49152		L.	
Just foll				ata	
1.) Select Comput 1.b) Select Shared	2.) Enter your resource usage:				
nis) delect charet	Threads Per Block	64	(Help)		
2.) Enter your reso Threads Per Block	Registers Per Thread	32	7		№ Nvidia's
Registers Per Thre	Shared Memory Per Block (bytes)	0		1	
Shared Memory Pe					"CUDA
(Do	(Don't edit anything below this line)				
3.) GPU Occupano	,			Н	Occupancy
Active Threads pe Active Warps per	3.) GPU Occupancy Data is displayed here and in the graphs:				Calculator"
Active Thread Blo	Active Threads per Multiprocessor	1024	(Help)	Н	Calculator
Occupancy of eac	Active Warps per Multiprocessor	32		H	
Dhuniaal Limita fa	Active Thread Blocks per Multiprocessor	16		Н	№ Shows
Physical Limits fo Threads per Warp	Occupancy of each Multiprocessor	50%		004	
Warps per Multipro Threads per Multip				024	relationship
Thread Blocks per					botwoon
Total # of 32-bit reg Register allocation	Physical Limits for GPU Compute Capability:	3.5			between
Register allocation	Threads per Warp	32			#rogistors/th
Registers per Threa Shared Memory pe	Warps per Multiprocessor	64		Ы	#registers/th
Shared Memory All	Threads per Multiprocessor	2048			read and the
Warp allocation gra Maximum Thread E	Thread Blocks per Multiprocessor	16			Teau and the
	Total # of 32-bit registers per Multiprocessor	65536			amount of
Allocated Resource	Register allocation unit size	256			_
Warps (Thi	Register allocation granularity	warp			SMT
Registers (Wa Shared Memory (B	Registers per Thread	255			
Note: SM is an abbreviation	Shared Memory per Multiprocessor (bytes)	49152			available
Maximum Thread	Shared Memory Allocation unit size	256			
Limited by Max W	Warp allocation granularity	4			Higher
Limited by Register Limited by Shared	Maximum Thread Block Size	1024		256	ingnei
Note: Occupancy limiter is s		1021		L	occupancy
			= Allocatable		
CUDA Occupancy (Version:	Allocated Resources	Per Block	Limit Per SM Blocks Per SM		enables
	Warps (Threads Per Block / Threads Per Warp)	2	64 16	7 I	
	Registers (Warp limit per SM due to per-warp reg count)	2	64 32		better hiding
	Shared Memory (Bytes)	0	49152		
		l O	49102	-	of memory
	Note: SM is an abbreviation for (Streaming) Multiprocessor				_
	Maximum Thread Blocks Per Multiprocessor	Blocks/SM	* Warps/Block = Warps/SM		access
	Limited by Max Warps or Max Blocks per Multiprocessor	16	2 32	H	latency
_	Limited by Registers per Multiprocessor	32	02		iatoricy
			2	¶ 49152	
	Limited by Shared Memory per Multiprocessor	16	2 32	152	
	J. 33.2.3.3		Shared Memory Per Block		puter Architecture Chapter 5.122

puter Architecture Chapter 5.122





Chapter summary

We have explored:

- Pipeline parallelism
- Dynamic instruction scheduling
- Static instruction scheduling
- **→** Multiple instructions per cycle
- Very long instruction words (VLIW)
- SIMT single-instruction multiple thread and how it maps onto SIMD
- Multi-threading
 - Coarse-grain
 - Fine-grain
 - Simultaneous multithreading (SMT)
 - Statically-partitioned multithreading

Vector instructions and SIMD - coming soon

Extra slides for interest/fun

Is the "minimum" operator associative?

min(min(X, Y), Z = min(X, min(Y, Z))?

min(X, Y) = if X<Y then X else Y</pre>

$$min(min(10, x), 100) = 100$$

Extra slides for interest/fun

Is the "minimum" operator associative?

- min(min(X, Y), Z = min(X, min(Y, Z))?
- min(X, Y) = if X<Y then X else Y</pre>

All comparisons on NaNs always fail....

min(min(10, NaN), 100) = 100

Extra slides for interest/fun

Is the "minimum" operator associative?

```
min(min(X, Y), Z = min(X, min(Y, Z))?
```

All comparisons on NaNs always fail....

```
min(X , NaN) = NaN
```

$$\mu \min(\min(X, NaN), Y) = \min(NaN, Y) = Y$$

$$min(X,min(NaN,Y)) = min(X,Y)$$

min(min(10, NaN), 100) = 100

Associativity in floating point

$$(a+b)+c = a+(b+c)$$
?

Example: Consider 3-digit base-10 floating-point

Consequence: many compilers use loop unrolling and reassociation to enhance parallelism in summations

And results are different!

But you can tell the compiler not to, eg:

- "-fp-model precise" with Intel's compilers
- (What's the right way to sum an array? See
 - http://en.wikipedia.org/wiki/Kahan_summation_algorithm)

- In the example processor that can only execute one instruction per cycle, unrolling is important because the loop control instructions become the critical factor.
- In machines that can issue multiple instructions per cycle, this is likely not the case there are opportunities to issue some instructions "for free" if you can schedule them into unused issue slots.
- In that case, software pipelining should lead to better performance than unrolling, though the difference might be small with a sufficiently-high unroll factor.
- You might also consider the energy cost: unrolling means we cache and store more instructions. But software pipelining without unrolling means we execute more loop-control instructions.
- Deviously if loop unrolling were to lead to instruction-cache misses, that'd be bad.
- So actually, the optimum strategy is likely to be a hybrid.
- This is actually only the beginning. You can sometimes do better by unrolling an *outer* loop this is called "unroll and jam", because we unroll the outer loop to produce two copies of the inner loop, then we jam them together. Consider matrix-matrix multiply (again!):

```
for(i=O; i<4; i++)
for(j=0; j<4; j++) {
    c[i][j] = 0;
    for(k=0; k<4; k++)
        c[i][j] = a[i][k]*b[k][j]+c[i][j];
}</pre>
```

This has limited parallelism due to the (loop-carried dependence involved in the) summation into C[i][j]. After unroll-and-jam of the j-loop by 1, we have:

```
for(i=O; i<4; i++)
  for(j=0; j<4; j+=2) {
    c[i][j] = 0;
    c[i][j+1] = 0;
    for(k=0; k<4; k++) {
      c[i][j]=a[i][k]*b[k][j]+c[i][j];
      c[i][j+1]=a[i][k]*b[k][j+1]+c[i][j+1];
}}</pre>
```

Unrolling versus software pipelining, and unroll-and-jam

- Now the inner loop has two summations to do, which are independent from one another. So it's more likely that you can fill the schedule more tightly.
- This example is taken from: http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.1.9319&rep=rep1&type=pdf
 Advanced Computer Architecture Chapter 5.130

What does this do?